

CS159 Lecture 1: Markov Decision Processes

Ugo Rosolia

Caltech

Spring 2021

Table of Contents

Markov Decision Processes

- Problem Formulation

- Control Policies and Value Functions

Solution Strategies

- Value Iteration

- Policy Iteration

- Linear Programming

Approximate Dynamic Programming

- Summary Policy and Value Iteration

- Approximate Policy Iteration

Table of Contents

Markov Decision Processes

- Problem Formulation

- Control Policies and Value Functions

Solution Strategies

- Value Iteration

- Policy Iteration

- Linear Programming

Approximate Dynamic Programming

- Summary Policy and Value Iteration

- Approximate Policy Iteration

Markov Decision Process


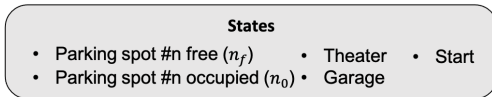
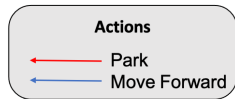
A Markov decision process (MDP) is a tuple $(\mathcal{S}, \mathcal{A}, T_s, c)$, where

- ▶ $\mathcal{S} = \{1, \dots, |\mathcal{S}|\}$ is a set of states;
- ▶ $\mathcal{A} = \{1, \dots, |\mathcal{A}|\}$ is a set of actions;
- ▶ The function $T_s : \mathcal{S} \times \mathcal{A} \times \mathcal{S} \rightarrow [0, 1]$ describes the probability of transitioning to a state s' given the action a and the system's state s ,

$$T_s(s, a, s') := \mathbb{P}(s_{k+1} = s' | s_k = s, a_k = a) = p(s' | s, a);$$

- ▶ The cost function $c : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$ assigns an instantaneous cost to each state-action pairs;

Markov Decision Process



Deterministic And Random Policies

Deterministic Policies

Define the set of deterministic policies Π^d . A deterministic policy $\pi^d \in \Pi^d$ maps states to actions, i.e.,

$$a_k = \pi^d(s_k).$$

Define the set of random policies Π^r . A random policy $\pi^r \in \Pi^r$ maps states to probability distributions, i.e.,

$$a_k \sim \pi^r(s_k).$$

Deterministic And Random Policies

Deterministic Policies

Define the set of deterministic policies Π^d . A deterministic policy $\pi^d \in \Pi^d$ maps states to actions, i.e.,

$$a_k = \pi^d(s_k).$$

Random Policies

Define the set of random policies Π^r . A random policy $\pi^r \in \Pi^r$ maps states to probability distributions, i.e.,

$$a_k \sim \pi^r(s_k).$$

Markov Decision Process

A Markov decision process (MDP) is a tuple $(\mathcal{S}, \mathcal{A}, T_s, R)$, where

- ▶ $\mathcal{S} = \{1, \dots, |\mathcal{S}|\}$ is a set of states;
- ▶ $\mathcal{A} = \{1, \dots, |\mathcal{A}|\}$ is a set of actions;
- ▶ The function $T_s : \mathcal{S} \times \mathcal{A} \times \mathcal{S} \rightarrow [0, 1]$ describes the probability of transitioning to a state s' given the action a and the system's state s ,

$$T_s(s, a, s') := \mathbb{P}(s_{k+1} = s' | s_k = s, a_k = a) = p(s' | s, a);$$

- ▶ The cost function $R : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$ assigns an instantaneous cost to each state-actions pairs;

Goal

Find a policy $\pi^* = [\pi_0^*, \pi_1^*, \dots]$ defined as

$$\pi^* = \arg \min_{\pi} \mathbb{E} \left[\sum_{t=0}^{\infty} \lambda^t c(s_t, a_t) | \pi \right]$$

Markov Decision Process

Goal

Find a policy $\pi^* = [\pi_0^*, \pi_1^*, \dots]$ defined as

$$\pi^* = \arg \min_{\pi} \mathbb{E} \left[\sum_{t=0}^{\infty} \lambda^t c(s_t, a_t) | \pi \right]$$

- ▶ The discount factor $\lambda \in (0, 1)$.
- ▶ The action $a_t = \pi_t(s_t)$ or $a_t \sim \pi_t(s_t)$.
- ▶ $\mathbb{E}[\sum_{t=0}^{\infty} \lambda^t c(s_t, a_t) | \pi]$ denotes the expectation under the policy π .

Markov Decision Process – Assumptions

Assumption 1. (Stationary costs and transition probabilities)

The cost function $c(s, a)$ and the transition probabilities $\mathbb{P}(s'|s, a)$ do not vary.

Assumption 2. (Bounded costs) The cost function

$|c(s, a)| \leq M < \infty$ for all $a \in \mathcal{A}$ and $s \in \mathcal{S}$.

Assumption 3. (Discrete State and Action Spaces) The state space \mathcal{S} and the action space \mathcal{A} are finite and discrete.

Assumption 4. (Discounting) The future costs are discounted by a factor λ and $0 \leq \lambda < 1$.

Table of Contents

Markov Decision Processes

Problem Formulation

Control Policies and Value Functions

Solution Strategies

Value Iteration

Policy Iteration

Linear Programming

Approximate Dynamic Programming

Summary Policy and Value Iteration

Approximate Policy Iteration

Deterministic And Random Policies

Deterministic Policies

Define the set of deterministic policies Π^d . A deterministic policy $\pi^d \in \Pi^d$ maps states to actions, i.e.,

$$a_k = \pi^d(s_k).$$

Random Policies

Define the set of random policies Π^r . A random policy $\pi^r \in \Pi^r$ maps states to probability distributions, i.e.,

$$a_k \sim \pi^r(s_k).$$

Deterministic Vs Random Policies

- For **unconstrained problems** we have that

$$\min_{\pi \in \Pi^d} \mathbb{E} \left[\sum_{t=0}^{\infty} \lambda^t c(s_t, a_t) | \pi \right] = \min_{\pi \in \Pi^r} \mathbb{E} \left[\sum_{t=0}^{\infty} \lambda^t c(s_t, a_t) | \pi \right]$$

There is no performance gain in optimizing over the larger set of random policies.

- For **constrained problems** we have that

$$\begin{array}{ll} \min_{\pi \in \Pi^d} & \mathbb{E} \left[\sum_{t=0}^{\infty} \lambda^t c(s_t, a_t) | \pi \right] \\ \text{s.t.} & \mathbb{E} \left[\sum_{t=0}^{\infty} \lambda^t g(s_t, a_t) | \pi \right] \leq \epsilon. \end{array} \geq \begin{array}{ll} \min_{\pi \in \Pi^r} & \mathbb{E} \left[\sum_{t=0}^{\infty} \lambda^t c(s_t, a_t) | \pi \right] \\ \text{s.t.} & \mathbb{E} \left[\sum_{t=0}^{\infty} \lambda^t g(s_t, a_t) | \pi \right] \leq \epsilon. \end{array}$$

A randomized policy perform better for constrained problems.

Deterministic Vs Random Policies

- For **unconstrained problems** we have that

$$\min_{\pi \in \Pi^d} \mathbb{E} \left[\sum_{t=0}^{\infty} \lambda^t c(s_t, a_t) | \pi \right] = \min_{\pi \in \Pi^r} \mathbb{E} \left[\sum_{t=0}^{\infty} \lambda^t c(s_t, a_t) | \pi \right]$$

There is no performance gain in optimizing over the larger set of random policies.

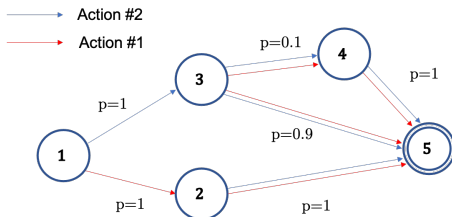
- For **constrained problems** we have that

$$\begin{array}{ll} \min_{\pi \in \Pi^d} \mathbb{E} \left[\sum_{t=0}^{\infty} \lambda^t c(s_t, a_t) | \pi \right] & \geq \min_{\pi \in \Pi^r} \mathbb{E} \left[\sum_{t=0}^{\infty} \lambda^t c(s_t, a_t) | \pi \right] \\ \text{s.t. } \mathbb{E} \left[\sum_{t=0}^{\infty} g(s_t, a_t) | \pi \right] \leq \epsilon. & \text{s.t. } \mathbb{E} \left[\sum_{t=0}^{\infty} g(s_t, a_t) | \pi \right] \leq \epsilon. \end{array}$$

A randomized policy performs better for constrained problems.

Deterministic Vs Random Policies

- ▶ The action space $\mathcal{A} = \{\text{Action 1}, \text{Action 2}\}$, state space $\mathcal{S} = \{1, 2, 3, 4, 5\}$ and the state $s = 5$ is a sink state.
- ▶ The cost function $c(s, a) = 0$ for all $s \in \mathcal{S} \setminus \{3\}$, $a \in \mathcal{A}$ and $c(3, a) = -1$ for all $a \in \mathcal{A}$.
- ▶ The constraint function $g(s, a) = 0$ for all $s \in \mathcal{S} \setminus \{4\}$, $a \in \mathcal{A}$ and $g(4, a) = 1$ for all $a \in \mathcal{A}$.
- ▶ Pick $\epsilon < 0.1$, the a deterministic policy must choose Action 1 from $s = 1$ to meet the constraint $\mathbb{E}[\sum_{t=0}^H g(s_t, a_t) | \pi] \leq \epsilon$.



Value Functions

Value Function

The value function v_π is a vector in $\mathbb{R}^{|\mathcal{S}|}$ where each entry $v_\pi(s)$ represents the cumulative cost of applying the policy $\pi \in \Pi^d$ from the state $s \in \mathcal{S}$, i.e.,

$$v_\pi(s) = \mathbb{E} \left[\sum_{t=0}^{\infty} \lambda^t c(s_t, a_t) | \pi, s \right].$$

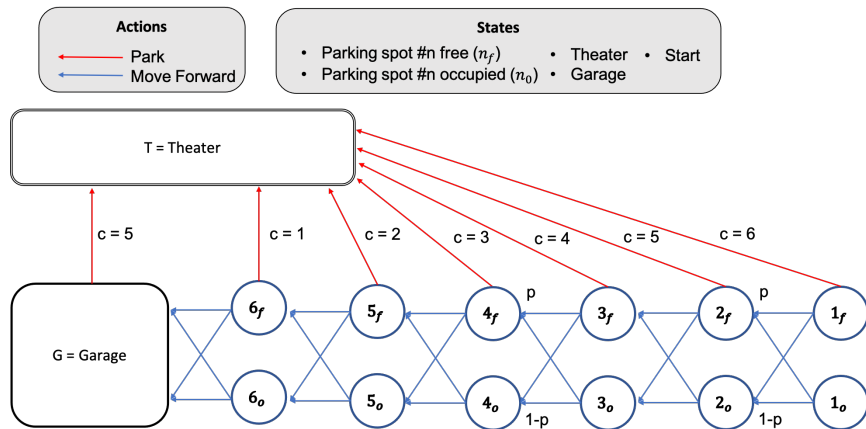
Consider a stationary policy $\pi = [\pi, \pi, \dots]$ with $\pi \in \Pi^d$. Then v_π is the unique solution of

$$v = r_\pi + \lambda P_\pi v$$

where

- ▶ the vector $r_\pi \in \mathbb{R}^{|\mathcal{S}|}$ where $r_\pi(s) = c(s, \pi(s))$
- ▶ the matrix $P_\pi \in \mathbb{R}^{|\mathcal{S}| \times |\mathcal{S}|}$ where $P_\pi(s, s') = p(s' | s, \pi(s))$
- ▶ the value function $v = (I - \lambda P_\pi)^{-1} r_\pi = \sum_{t=0}^{\infty} \lambda^t P_\pi^t r_\pi$

Value Function – The Parking Example



The set of states

$$\mathcal{S} = \{1_f, 1_o, 2_f, 2_o, 3_f, 3_o, 4_f, 4_o, 5_f, 5_o, 6_f, 6_o, G, T\}.$$

Value Function – The Parking Example

Two actions are available: {move forward, park}.

Let π_m be a deterministic policy that selects the action move forward, then P_{π_m} is defined by the following table:

1f	1o	2f	2o	...	G	T	
		p	1-p				1f
		p	1-p				1o
				\ddots			\vdots
					1		6f
					1		6o
						1	G
						1	T

where each entry $P_{\pi}(s, s') = p(s'|s, \pi(s))$ for $s \in \mathcal{S}$, $s' \in \mathcal{S}$ and $\mathcal{S} = \{1_f, 1_o, 2_f, 2_o, 3_f, 3_o, 4_f, 4_o, 5_f, 5_o, 6_f, 6_o, G, T\}$.

Value Function – The Parking Example

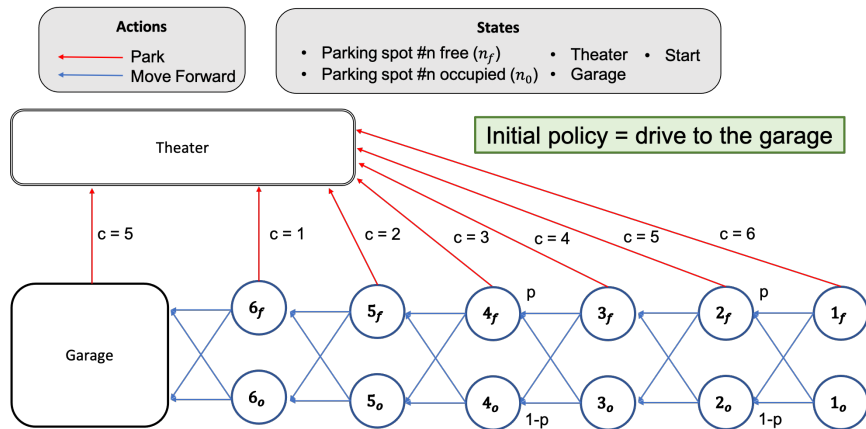
Two actions are available: {move forward, park}.

Let π_p be a deterministic policy that selects the action park, then P_{π_p} is defined by the following table:

	1f	1o	2f	2o	...	G	T	
							1	1f
			p	1-p				1o
					...			:
						1		6f
						1		6o
							1	G
							1	T

where each entry $P_{\pi}(s, s') = p(s'|s, \pi(s))$ for $s \in \mathcal{S}$, $s' \in \mathcal{S}$ and $\mathcal{S} = \{1_f, 1_o, 2_f, 2_o, 3_f, 3_o, 4_f, 4_o, 5_f, 5_o, 6_f, 6_o, G, T\}$.

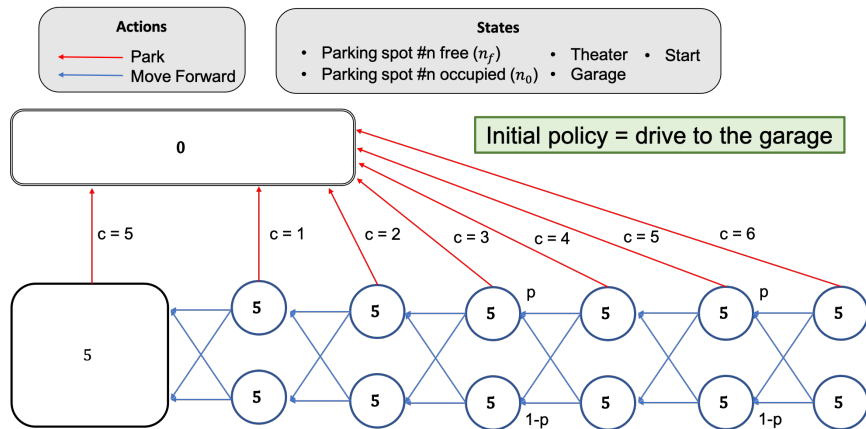
Value Function – The Parking Example



State Vector = $[1_f, 1_o, 2_f, 2_o, 3_f, 3_o, 4_f, 4_o, 5_f, 5_o, 6_f, 6_o, G, T]$.

Value Function = $[5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 0]$.

Value Function – The Parking Example



State Vector = $[1_f, 1_o, 2_f, 2_o, 3_f, 3_o, 4_f, 4_o, 5_f, 5_o, 6_f, 6_o, G, T]$.

Value Function = $[5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 0]$.

Markov Decision Process

Goal

Find a stationary policy $\pi^* = [\pi^*, \pi^*, \dots]$ defined as

$$\pi^* = \arg \min_{\pi} \mathbb{E} \left[\sum_{t=0}^{\infty} \lambda^t c(s_t, a_t) | \pi \right]$$

Given a value function which satisfies:

$$v^*(s) = \arg \min_{a \in \mathcal{A}} c(s, a) + \sum_{s' \in \mathcal{S}} \lambda v^*(s') p(s' | s, a).$$

Then, the optimal policy is:

$$\pi(s) = \min_{a \in \mathcal{A}} c(s, a) + \sum_{s' \in \mathcal{S}} \lambda v^*(s') p(s' | s, a)$$

Markov Decision Process

Goal

Find a stationary policy $\pi^* = [\pi^*, \pi^*, \dots]$ defined as

$$\pi^* = \arg \min_{\pi} \mathbb{E} \left[\sum_{t=0}^{\infty} \lambda^t c(s_t, a_t) | \pi \right]$$

Optimality Conditions

Given the optimal value function v^* that satisfies the Bellman recursion $v^* = Bv^*$ defined as follows:

$$v^*(s) = \min_{a \in \mathcal{A}} [c(s, a) + \sum_{s' \in \mathcal{S}} \lambda v^*(s') p(s' | s, a)], \quad \forall s \in \mathcal{S}.$$

Then, the optimal policy is:

$$\pi^*(s) = \arg \min_{a \in \mathcal{A}} [c(s, a) + \sum_{s' \in \mathcal{S}} \lambda v^*(s') p(s' | s, a)]$$

Table of Contents

Markov Decision Processes

- Problem Formulation

- Control Policies and Value Functions

Solution Strategies

- Value Iteration

- Policy Iteration

- Linear Programming

Approximate Dynamic Programming

- Summary Policy and Value Iteration

- Approximate Policy Iteration

Value Iteration

Algorithm Steps:

1. Select $v^0 \in \mathbb{R}^{|\mathcal{S}|}$, set $k = 0$ and pick a tolerance $\epsilon \geq 0$

Value Iteration

Algorithm Steps:

1. Select $v^0 \in \mathbb{R}^{|\mathcal{S}|}$, set $k = 0$ and pick a tolerance $\epsilon \geq 0$
2. For each $s \in \mathcal{S}$ compute $v^{k+1} \in \mathbb{R}^{|\mathcal{S}|}$ where

$$v^{k+1}(s) = \min_{a \in \mathcal{A}} [c(s, a) + \sum_{s' \in \mathcal{S}} \lambda p(s'|s, a) v^k(s')]$$

Value Iteration

Algorithm Steps:

1. Select $v^0 \in \mathbb{R}^{|\mathcal{S}|}$, set $k = 0$ and pick a tolerance $\epsilon \geq 0$
2. For each $s \in \mathcal{S}$ compute $v^{k+1} \in \mathbb{R}^{|\mathcal{S}|}$ where

$$v^{k+1}(s) = \min_{a \in \mathcal{A}} [c(s, a) + \sum_{s' \in \mathcal{S}} \lambda p(s'|s, a) v^k(s')]$$

3. If

$$\|v^{k+1} - v^k\| \geq \epsilon \frac{(1 - \lambda)}{2\lambda}$$

set $k = k + 1$ and go to step 2.

Value Iteration

Algorithm Steps:

1. Select $v^0 \in \mathbb{R}^{|\mathcal{S}|}$, set $k = 0$ and pick a tolerance $\epsilon \geq 0$
2. For each $s \in \mathcal{S}$ compute $v^{k+1} \in \mathbb{R}^{|\mathcal{S}|}$ where

$$v^{k+1}(s) = \min_{a \in \mathcal{A}} [c(s, a) + \sum_{s' \in \mathcal{S}} \lambda p(s'|s, a) v^k(s')]$$

3. If

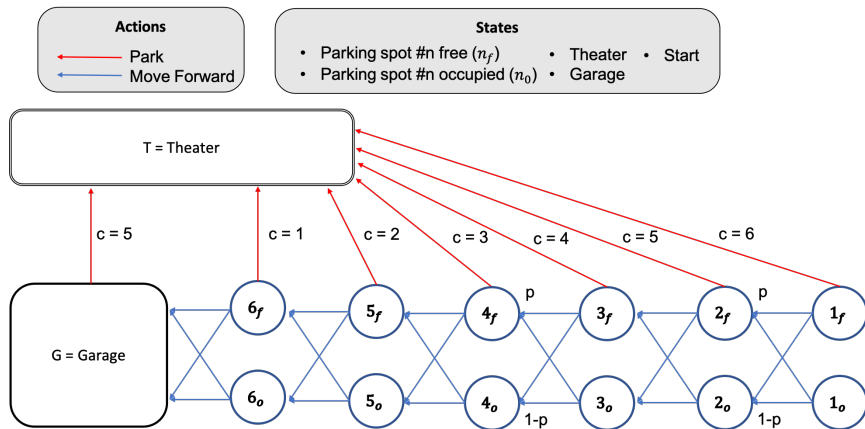
$$\|v^{k+1} - v^k\| \geq \epsilon \frac{(1 - \lambda)}{2\lambda}$$

set $k = k + 1$ and go to step 2.

4. Define the control policy

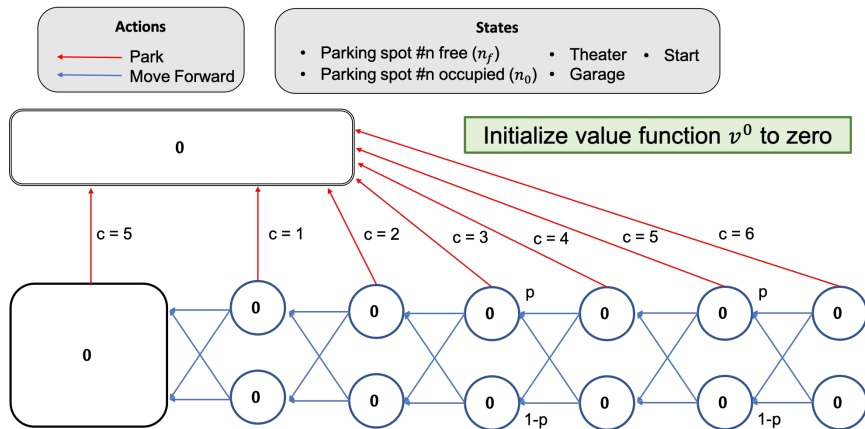
$$\pi^{\text{vi}}(s) = \arg \min_{a \in \mathcal{A}} [c(s, a) + \sum_{s' \in \mathcal{S}} \lambda p(s'|s, a) v^{k+1}(s')]$$

Value Iteration



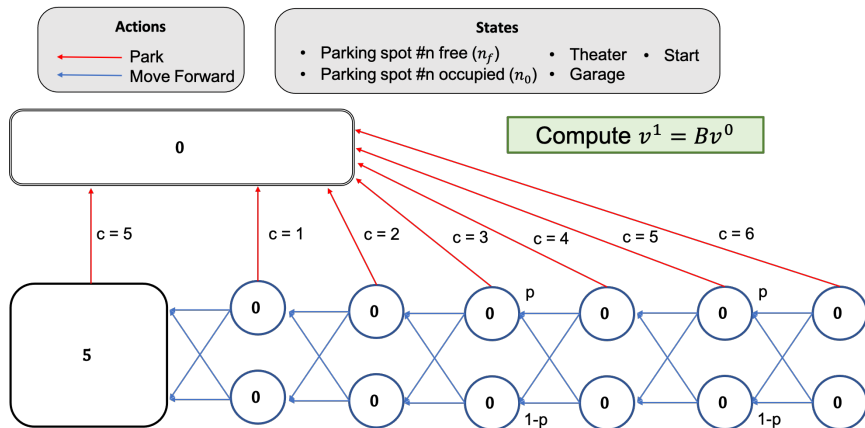
State Vector = $[1_f, 1_o, 2_f, 2_o, 3_f, 3_o, 4_f, 4_o, 5_f, 5_o, 6_f, 6_o, G, T]$.

Value Iteration



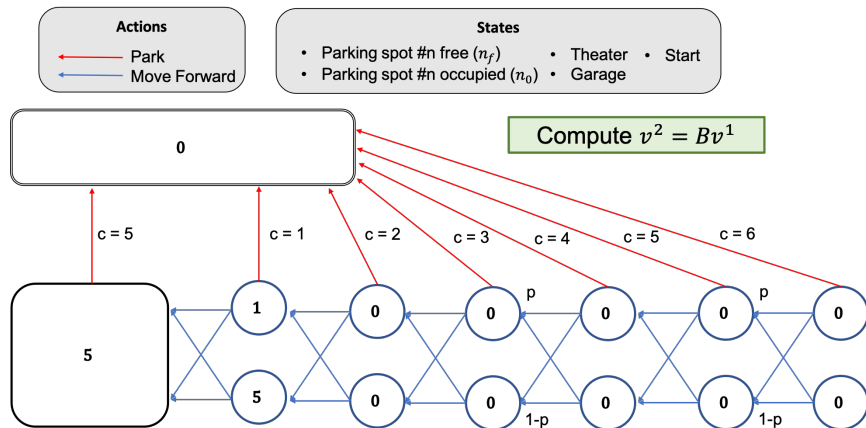
State Vector = $[1_f, 1_o, 2_f, 2_o, 3_f, 3_o, 4_f, 4_o, 5_f, 5_o, 6_f, 6_o, G, T]$.

Value Iteration



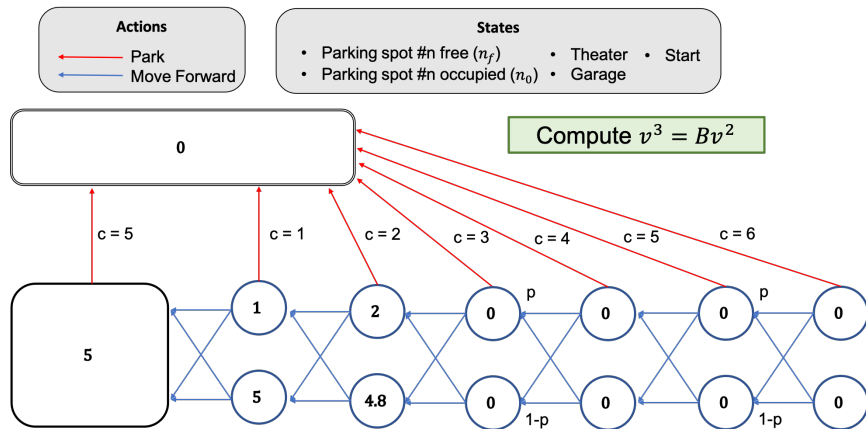
State Vector = $[1_f, 1_o, 2_f, 2_o, 3_f, 3_o, 4_f, 4_o, 5_f, 5_o, 6_f, 6_o, G, T]$.

Value Iteration



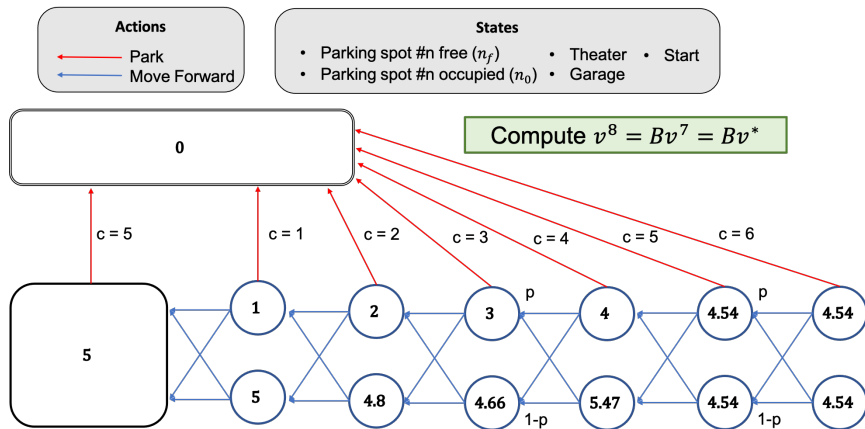
State Vector = $[1_f, 1_o, 2_f, 2_o, 3_f, 3_o, 4_f, 4_o, 5_f, 5_o, 6_f, 6_o, G, T]$.

Value Iteration



State Vector = $[1_f, 1_o, 2_f, 2_o, 3_f, 3_o, 4_f, 4_o, 5_f, 5_o, 6_f, 6_o, G, T]$.

Value Iteration



State Vector = $[1_f, 1_o, 2_f, 2_o, 3_f, 3_o, 4_f, 4_o, 5_f, 5_o, 6_f, 6_o, G, T]$.

Value Iteration: Properties

Theorem

Let $\{v^k\}$ be a sequence defined by the Bellman recursion and consider the stopping rule

$$\|v^{k+1} - v^k\|_{\infty} < \epsilon \frac{(1 - \lambda)}{2\lambda} \quad (1)$$

Then we have that

- ▶ v^k converges in norm to v^* and the convergence is linear with rate λ .
- ▶ If (1) holds for a finite N , then (1) holds for $k \geq N$.
- ▶ If (1) holds for a finite N , then $\|v^{N+1} - v^*\|_{\infty} < \epsilon/2$ and π^{vi} is ϵ -optimal.

Variants to the Value Iteration with better convergence rate in Chapter 6 of “Markov decision processes: discrete stochastic dynamic programming” by M. Puterman. John Wiley & Sons, 2014.

Value Iteration: Convergence Proof

Define the Bellman backup operator $B : \mathbb{R}^{|S|} \rightarrow \mathbb{R}^{|S|}$

$$Bv(s) = \min_{a \in \mathcal{A}} [c(s, a) + \sum_{s' \in \mathcal{S}} \lambda p(s'|s, a) v(s')]$$

which is a contraction as

$$\begin{aligned} |Bv_0(s) - Bv_1(s)| &= \left| \min_{a \in \mathcal{A}} [c(s, a) + \sum_{s' \in \mathcal{S}} \lambda p(s'|s, a) v_0(s')] \right. \\ &\quad \left. - \min_{a \in \mathcal{A}} [c(s, a) + \sum_{s' \in \mathcal{S}} \lambda p(s'|s, a) v_1(s')] \right| \\ &\leq \max_{a \in \mathcal{A}} \lambda \left| \sum_{s' \in \mathcal{S}} p(s'|s, a) v_0(s') - \sum_{s' \in \mathcal{S}} p(s'|s, a) v_1(s') \right| \\ &= \max_{a \in \mathcal{A}} \lambda \sum_{s' \in \mathcal{S}} p(s'|s, a) |v_0(s') - v_1(s')| \\ &\leq \lambda \max_{s' \in \mathcal{S}} |v_0(s') - v_1(s')|. \end{aligned}$$

Then, by the fixed-point theorem, we have that $Bv^* = v^*$ and the sequence $v^{k+1} = Bv^k = B^{k+1}v^0$ converges to v^* .

Value Iteration: Suboptimality Proof

Now define the Bellman backup for a policy π as $B_\pi : \mathbb{R}^{|S|} \rightarrow \mathbb{R}^{|S|}$

$$B_\pi v(s) = c(s, \pi(a)) + \sum_{s' \in S} \lambda p(s'|s, \pi(a)) v(s').$$

We notice that

$$\begin{aligned} \|v_{\pi^{vi}} - v^*\|_\infty &= \|B_{\pi^{vi}} v_{\pi^{vi}} - v^*\|_\infty \\ &\leq \|B_{\pi^{vi}} v_{\pi^{vi}} - Bv^{k+1}\|_\infty + \|Bv^{k+1} - v^*\|_\infty \\ &= \|B_{\pi^{vi}} v_{\pi^{vi}} - B_{\pi^{vi}} v^{k+1}\|_\infty + \|Bv^{k+1} - Bv^*\|_\infty \\ &\leq \lambda \|v_{\pi^{vi}} - v^{k+1}\|_\infty + \lambda \|v^{k+1} - v^*\|_\infty. \end{aligned}$$

Rearranging terms and leveraging the stopping rule yields to

$$\|v^{k+1} - v^*\|_\infty \leq \frac{\lambda}{1 - \lambda} \|v^{k+1} - v^k\|_\infty \leq \frac{\epsilon}{2}$$

Table of Contents

Markov Decision Processes

Problem Formulation

Control Policies and Value Functions

Solution Strategies

Value Iteration

Policy Iteration

Linear Programming

Approximate Dynamic Programming

Summary Policy and Value Iteration

Approximate Policy Iteration

Policy Iteration

Algorithm Steps:

1. Set $k = 0$ and select a policy $\pi^k \in \Pi^d$.

Policy Iteration

Algorithm Steps:

1. Set $k = 0$ and select a policy $\pi^k \in \Pi^d$.
2. **(Policy Evaluation)**. Compute the value function $v_{\pi^k}^k \in \mathbb{R}^{|\mathcal{S}|}$ that is the solution to the following equation:

$$v_{\pi^k}^k(s) = c(s, \pi^k(s)) + \sum_{s' \in \mathcal{S}} \lambda p(s'|s, \pi^k(s)) v_{\pi^k}^k(s').$$

Recall that $v_{\pi^k}^k = (I - P_{\pi^k})^{-1} r_{\pi^k}$.

Policy Iteration

Algorithm Steps:

1. Set $k = 0$ and select a policy $\pi^k \in \Pi^d$.
2. **(Policy Evaluation)**. Compute the value function $v_{\pi^k}^k \in \mathbb{R}^{|\mathcal{S}|}$ that is the solution to the following equation:

$$v_{\pi^k}^k(s) = c(s, \pi^k(s)) + \sum_{s' \in \mathcal{S}} \lambda p(s'|s, \pi^k(s)) v_{\pi^k}^k(s').$$

Recall that $v_{\pi^k}^k = (I - P_{\pi^k})^{-1} r_{\pi^k}$.

3. **(Policy Improvement)**. Set

$$\pi^{k+1}(s) = \min_{a \in \mathcal{S}} \left[c(s, a) + \sum_{s' \in \mathcal{S}} \lambda p(s'|s, a) v_{\pi^k}^k(s') \right].$$

Policy Iteration

Algorithm Steps:

1. Set $k = 0$ and select a policy $\pi^k \in \Pi^d$.
2. **(Policy Evaluation)**. Compute the value function $v_{\pi^k}^k \in \mathbb{R}^{|\mathcal{S}|}$ that is the solution to the following equation:

$$v_{\pi^k}^k(s) = c(s, \pi^k(s)) + \sum_{s' \in \mathcal{S}} \lambda p(s'|s, \pi^k(s)) v_{\pi^k}^k(s').$$

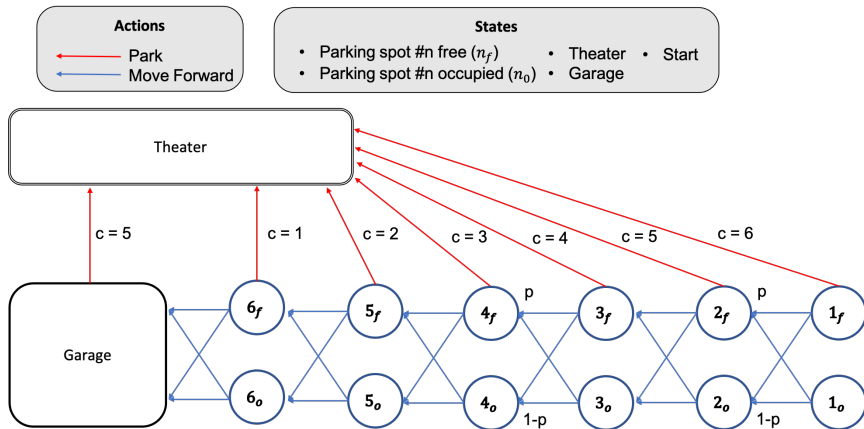
Recall that $v_{\pi^k}^k = (I - P_{\pi^k})^{-1} r_{\pi^k}$.

3. **(Policy Improvement)**. Set

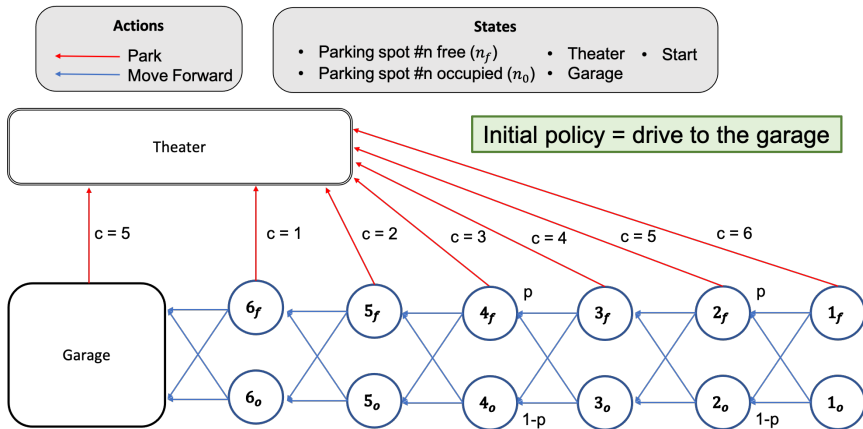
$$\pi^{k+1}(s) = \min_{a \in \mathcal{S}} \left[c(s, a) + \sum_{s' \in \mathcal{S}} \lambda p(s'|s, a) v_{\pi^k}^k(s') \right].$$

4. If $\pi^k = \pi^{k+1}$ stop, $\pi^* = \pi^k$. Otherwise, set $k = k + 1$ and go to Step 2.

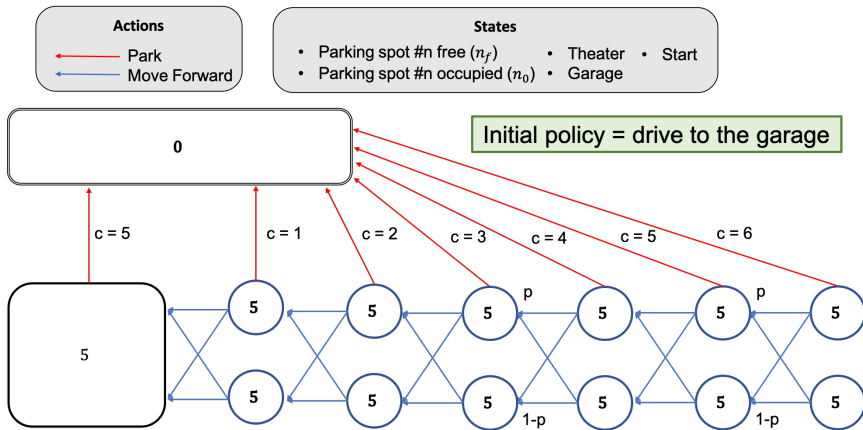
Policy Iteration



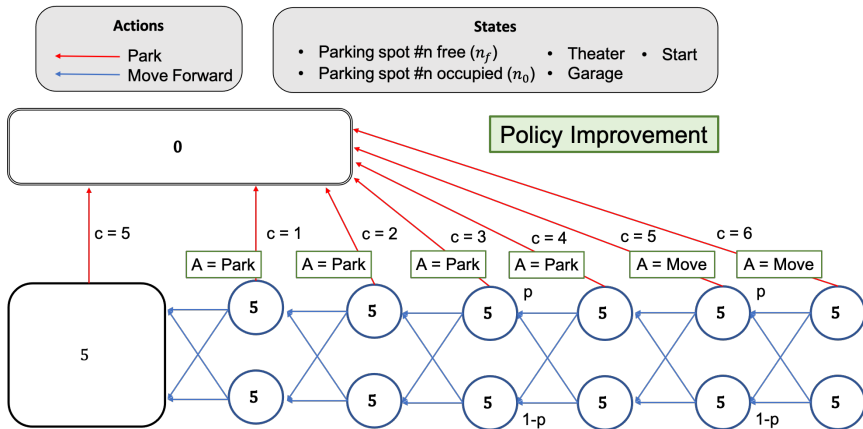
Policy Iteration



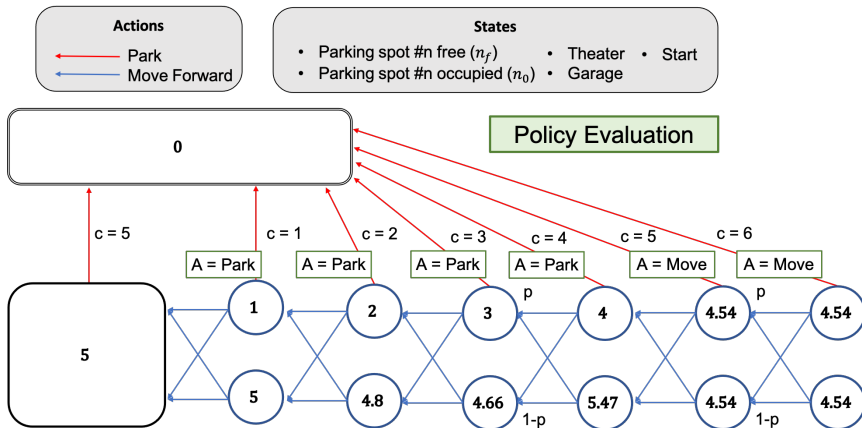
Policy Iteration



Policy Iteration



Policy Iteration



Policy Evaluation Step

Direct Strategy

Solve the linear system of equations

$$v_{\pi^k}^k = (I - \lambda P_{\pi^k})^{-1} r_{\pi^k}$$

Set $v_{\pi^k}^{k,0}(s) = 0$

Iterate $v_{\pi^k}^{k,i+1}(s) = c(s, \pi^k(s)) + \sum_{s' \in \mathcal{S}} \lambda p(s'|s, \pi^k(s)) v_{\pi^k}^{k,i}(s')$

Stop when $v_{\pi^k}^{k,i+1}(s) = v_{\pi^k}^{k,i}(s)$ for all $s \in \mathcal{S}$ and set $v_{\pi^k}^{k,i} = v_{\pi^k}^k$

Policy Evaluation Step

Direct Strategy

Solve the linear system of equations

$$v_{\pi^k}^k = (I - \lambda P_{\pi^k})^{-1} r_{\pi^k}$$

Iterative Strategy

Set $v_{\pi^k}^{k,0}(s) = 0$

Iterate $v_{\pi^k}^{k,i+1}(s) = c(s, \pi^k(s)) + \sum_{s' \in \mathcal{S}} \lambda p(s'|s, \pi^k(s)) v_{\pi^k}^{k,i}(s')$

Stop when $v_{\pi^k}^{k,i+1}(s) = v_{\pi^k}^{k,i}(s)$ for all $s \in \mathcal{S}$ and set $v_{\pi^k}^{k,i} = v_{\pi^k}^k$

Policy Evaluation: Properties

Theorem

For the policy iteration algorithm we have that

- ▶ The value function is non-increasing, i.e., $v_{\pi^{k+1}}^{k+1} \leq v_{\pi^k}^k$
- ▶ The algorithm converges in a finite number of iterations
- ▶ Let π^∞ be the policy at convergence, then $\pi^\infty = \pi^*$

Policy Evaluation: Properties

Proof sketch:

- ▶ The value function is non-increasing and there is a finite number of policies (as the number of action is finite). Therefore, the policy iteration algorithm converges in a finite number of iterations
- ▶ At convergence we have that $\pi^{k+1} = \pi^k$ and therefore

$$v^{k+1}(s) = \min_{a \in \mathcal{A}} \left[c(s, a) + \sum_{s' \in \mathcal{S}} \lambda p(s'|s, a) v^{k+1}(s') \right], \forall s \in \mathcal{S}.$$

Hence, v^{k+1} satisfies the Bellman equation and $\pi^{k+1} = \pi^*$.

Table of Contents

Markov Decision Processes

Problem Formulation

Control Policies and Value Functions

Solution Strategies

Value Iteration

Policy Iteration

Linear Programming

Approximate Dynamic Programming

Summary Policy and Value Iteration

Approximate Policy Iteration

Linear Programming

Linear Programming

Let $\alpha(s) > 0$ for all $s \in \mathcal{S}$ and

$$\bar{v} = \arg \max_{v \in \mathbb{R}^{|\mathcal{S}|}} \sum_{s \in \mathcal{S}} \alpha(s) v(s)$$

$$\text{subject to } v(s) \leq c(s, a) + \sum_{s' \in \mathcal{S}} \lambda p(s'|s, a) v(s'),$$

$$\forall a \in \mathcal{A}, \forall s \in \mathcal{S}.$$

then, we have that $\bar{v} = v^*$.

Linear Programming

Proof Sketch. By feasibility of \bar{v} we have

$$\bar{v}(s) \leq c(s, a) + \sum_{s' \in \mathcal{S}} \lambda p(s'|s, a) \bar{v}(s'), \quad \forall a \in \mathcal{A}, \quad \forall s \in \mathcal{S}.$$

which is equivalent to

$$\bar{v}(s) \leq \min_{a \in \mathcal{A}} \left[c(s, a) + \sum_{s' \in \mathcal{S}} \lambda p(s'|s, a) \bar{v}(s') \right] = Bv(\bar{s}), \quad \forall s \in \mathcal{S}.$$

Now recall that B is monotone and therefore

$v(s) \leq Bv(s) \leq B^2v(s) \leq \dots \leq B^\infty v(s) = v^*(s), \quad \forall s \in \mathcal{S}$. Hence, any feasible solution $v(s) \leq Bv(s) \leq v^*(s) = Bv^*(s)$. Concluding as $\alpha(s) > 0$, the feasible solution $v^*(s)$ is optimal.

Table of Contents

Markov Decision Processes

- Problem Formulation

- Control Policies and Value Functions

Solution Strategies

- Value Iteration

- Policy Iteration

- Linear Programming

Approximate Dynamic Programming

- Summary Policy and Value Iteration

- Approximate Policy Iteration

Summary Policy and Value Iteration

Policy Iteration

Policy Evaluation: Find V_{π^k} by solving

$$V_{\pi^k}(s) = c(s, \pi^k(s)) + \sum_{s' \in \mathcal{S}} \lambda p(s'|s, \pi^k(s)) V_{\pi^k}(s'), \quad \forall s \in \mathcal{S}.$$

Policy Improvement: Compute π^{k+1} as

$$\pi^{k+1}(s) = \min_{a \in \mathcal{A}} \left[c(s, a) + \sum_{s' \in \mathcal{S}} \lambda p(s'|s, a) V_{\pi^k}(s') \right], \quad \forall s \in \mathcal{S}.$$

Value Iteration

For any $V \in \mathbb{R}^{|\mathcal{S}|}$ compute

$$V^*(s) = \lim_{k \rightarrow \infty} B^k V(s), \quad \forall s \in \mathcal{S}.$$

Summary Policy and Value Iteration

Policy Iteration

Policy Evaluation: Find V_{π^k} by solving

$$V_{\pi^k}(s) = c(s, \pi^k(s)) + \sum_{s' \in \mathcal{S}} \lambda p(s'|s, \pi^k(s)) V_{\pi^k}(s'), \quad \forall s \in \mathcal{S}.$$

Policy Improvement: Compute π^{k+1} as

$$\pi^{k+1}(s) = \min_{a \in \mathcal{A}} [c(s, a) + \sum_{s' \in \mathcal{S}} \lambda p(s'|s, a) V_{\pi^k}(s')], \quad \forall s \in \mathcal{S}.$$

Value Iteration

For any $V \in \mathbb{R}^{|\mathcal{S}|}$ compute

$$V^*(s) = \lim_{k \rightarrow \infty} B^k V(s), \quad \forall s \in \mathcal{S}.$$

Table of Contents

Markov Decision Processes

- Problem Formulation

- Control Policies and Value Functions

Solution Strategies

- Value Iteration

- Policy Iteration

- Linear Programming

Approximate Dynamic Programming

- Summary Policy and Value Iteration

- Approximate Policy Iteration

Approximate Policy Iteration

Policy Iteration

Policy Evaluation: Find V_{π^k} by solving

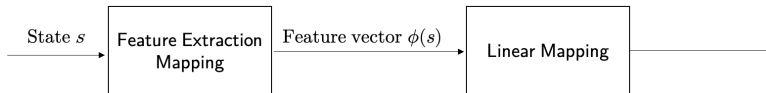
$$V_{\pi^k}(s) = c(s, \pi^k(s)) + \sum_{s' \in \mathcal{S}} \lambda p(s'|s, \pi^k(a)) V_{\pi^k}(s'), \quad \forall s \in \mathcal{S}.$$

Policy Improvement: Compute π^{k+1} as

$$\pi^{k+1}(s) = \min_{a \in \mathcal{A}} \left[c(s, a) + \sum_{s' \in \mathcal{S}} \lambda p(s'|s, a) V_{\pi^k}(s') \right], \quad \forall s \in \mathcal{S}.$$

- ▶ Perform the policy evaluation step for all $s \in \bar{\mathcal{S}} \subset \mathcal{S}$
- ▶ Similar strategies for Value Iteration and Linear Programming

Approximation in the Value Space



Value function approximation

$$\hat{V}_{\theta}(s) = \sum_i \theta_i \phi_i(s) = \theta^{\top} \phi(s)$$

Chess example

- ▶ $\phi_1(s)$ = *material score* computed summing the points with the pieces on the board (pawn = 1, rook = 5, Knight and Bishops = 3, queen = 10)
- ▶ $\phi_2(s)$ = *mobility* given by the legal moves available,
- ▶ $\phi_3(s)$ = *center control* given by the number of pawns in the center
- ▶ $\phi_4(s)$ = *bishop's mobility* given by the amount of squared reachable by the bishop,



Policy Iteration w/ Value Function Approximation

We focus on a variant of approximate policy iteration based on Monte Carlo simulations and function approximation.

Approximate Policy Iteration

Policy Evaluation: For a set of representative states $\bar{\mathcal{S}} \subset \mathcal{S}$ run M simulations using the policy π^k . Then, compute the cost of each i th simulation from the state $s \in \bar{\mathcal{S}}$ denoted as $\bar{c}(i, s)$ and approximate the value function $\hat{V}_\theta(s) = \sum_{s \in \mathcal{S}} \theta^\top \phi(s)$ solving the following problem

$$\theta^k = \arg \min_{\theta} \sum_{s \in \bar{\mathcal{S}}} \sum_{i=1}^M \|\hat{V}_\theta(s) - \bar{c}(i, s)\|.$$

Policy Improvement: Compute π^{k+1} as

$$\pi^{k+1} = \min_{a \in \mathcal{A}} \left[c(s, a) + \sum_{s' \in \mathcal{S}} \lambda p(s'|s, a) \hat{V}_{\theta^k}(s') \right].$$

Theoretical Basis for Approximate Policy Iteration

Theorem

If policies are approximately evaluated using an approximated value function such that

$$\max_s |V_{\theta^k}(s) - V_{\pi^k}(s)| \leq \delta, \quad \forall k = 0, 1, \dots$$

and the policy improvement is approximate

$$\max_s |B_{\pi^{k+1}} V_{\theta^k}(s) - B V_{\theta^k}(s)| \leq \epsilon, \quad \forall k = 0, 1, \dots$$

Then, we have that

$$\limsup_{k \rightarrow \infty} \max_s |V_{\pi^k}(s) - V^*(s)| \leq \frac{\epsilon + 2\lambda\delta}{(1 - \lambda)^2}$$

Readings

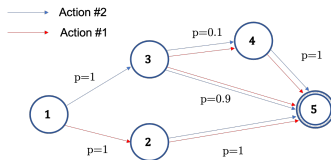
- ▶ Chapter 2 and Chapter 6.2 “Neuro-Dynamic Programming”
Dimitri P. Bertsekas and John Tsitsiklis
- ▶ Chapter 6 “Markov decision processes: discrete stochastic dynamic programming.” M. Puterman
- ▶ D. Bertsekas, “Feature-based aggregation and deep reinforcement learning: A survey and some new implementations.” IEEE/CAA Journal of Automatica Sinica 6.1 (2018): 1-31.
- ▶ D. Bertsekas, “Biased aggregation, rollout, and enhanced policy improvement for reinforcement learning.” arXiv preprint arXiv:1910.02426 (2019).
- ▶ D. P. De Farias, and B. Van Roy. “The linear programming approach to approximate dynamic programming.” Operations research 51.6 (2003): 850-865.

Summary

- ▶ We discussed how to solve optimal control problem with discrete state and action spaces of the form

$$\pi^* = \arg \min_{\pi} \mathbb{E} \left[\sum_{t=0}^{\infty} \lambda^t r(s_t, a_t) | \pi \right].$$

- ▶ The solution can be computed exactly given a known model and state-action spaces of moderate size.
- ▶ Approximate dynamic programming can be used to reduce the computational complexity of synthesis strategies.



What is next?

Optimal Control Problem with Continuous State Spaces:

In the next lectures we will

- ▶ Compute a control policy mapping continuous state to continuous control action

$$\pi : \mathbb{R}^n \rightarrow \mathbb{R}^d$$

- ▶ Leverage the same ideas to synthesize optimal policies, but computing/approximating the value function is harder for problem with constraints.
- ▶ Present learning-based strategies to approximate the value function in continuous state-action spaces.