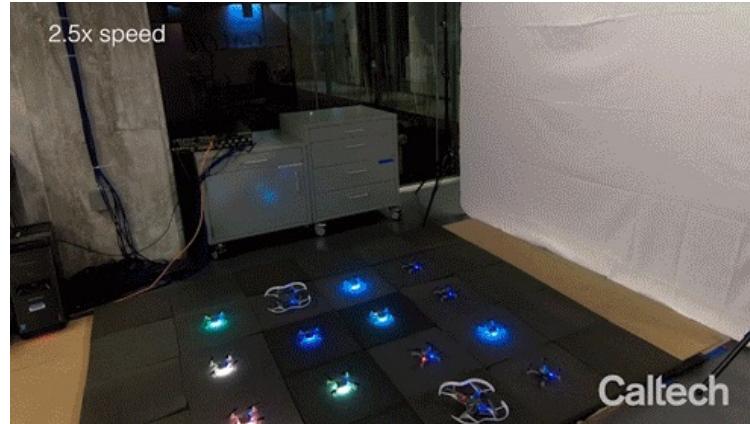


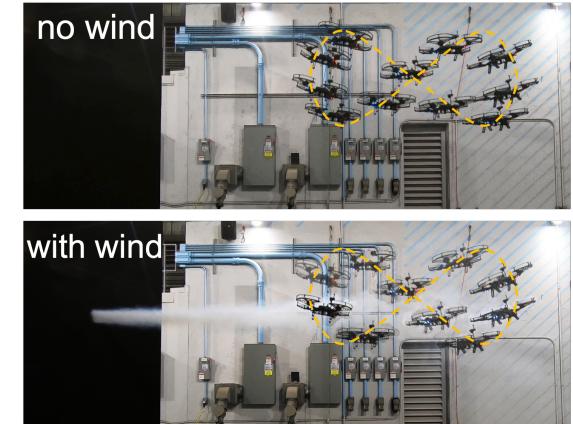
# Deep Residual Learning for Control



*Shi & Shi & O'Connell et al., ICRA' 19*



*Shi et al., ICRA' 20, T-RO*



*O'Connell & Shi et al., ongoing work*



**Autonomous Robotic Camera**

*Le et al., ICML' 16*



**“1+1>2”**

Guanya Shi

[www.gshi.me](http://www.gshi.me)

Caltech CS 159 Guest Lecture

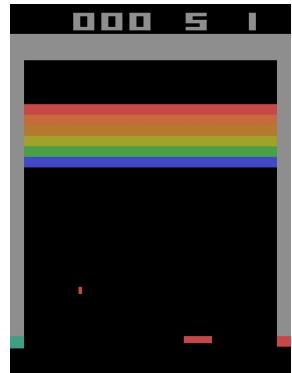


*Zeng et al., RSS' 19, T-RO*

**Caltech**

# Motivations: Real-World Decision Making is Hard

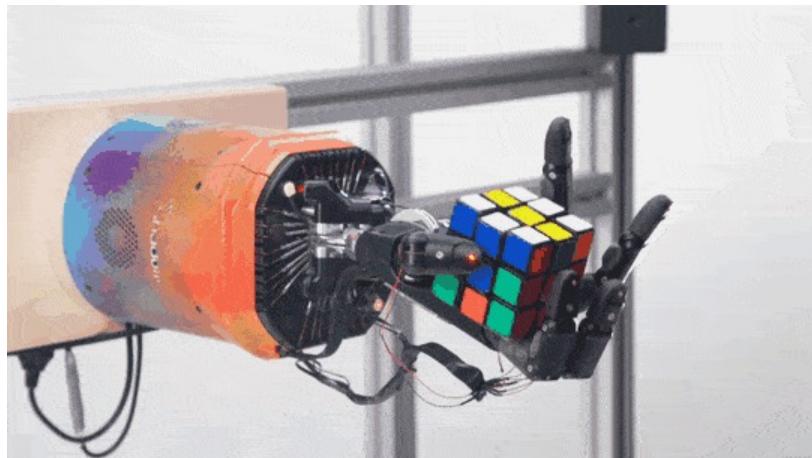
- Q: Can we broaden the successful horizon of deep learning?



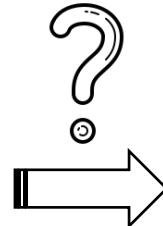
Atari



AlphaGo



Manipulator



"I want to use deep learning to optimize the design, manufacturing and operation of our aircrafts. But I need some guarantees." -- Aerospace Director



**Question: How safe do autonomous vehicles need to be?**

- As safe as human-driven cars (7 deaths every  $10^9$  miles)
- As safe as buses and trains (0.1-0.4 deaths every  $10^9$  miles)
- As safe as airplanes (0.07 deaths every  $10^9$  miles)

I. Savage, "Comparing the fatality risks in United States transportation across modes and over time", *Research in Transportation Economics*, 43:9-22, 2013.

"Can we really use ML in safety critical system?"

Richard Murray, IPAM ICLO workshop 2020

# Where are the Challenges from?

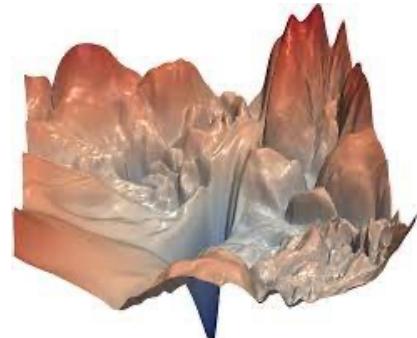
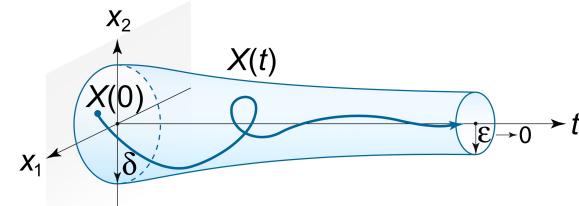
- Uncertainties
  - Modelling mismatch, sim2real gap
  - Delay, motor failures, saturations, ...
- Limited computational power
  - Data collection is also not cheap
- We need some “formal” interpretability
  - E.g., stability, safety
  - Neural networks are “hard” to control



R.I.P., propellers



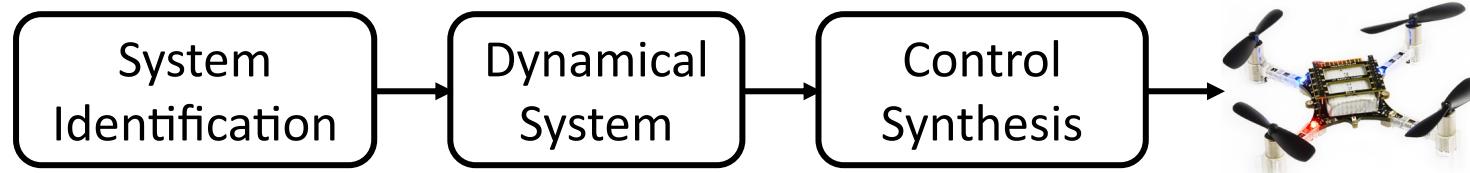
Crazyflie, weight 34g



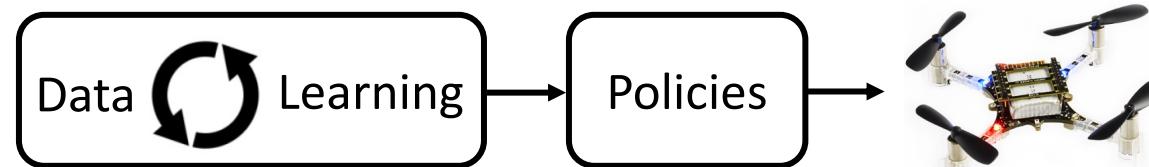
NN landscape, Li et al., NeurIPS 2018

# Idea: Control Theory Meets Machine Learning

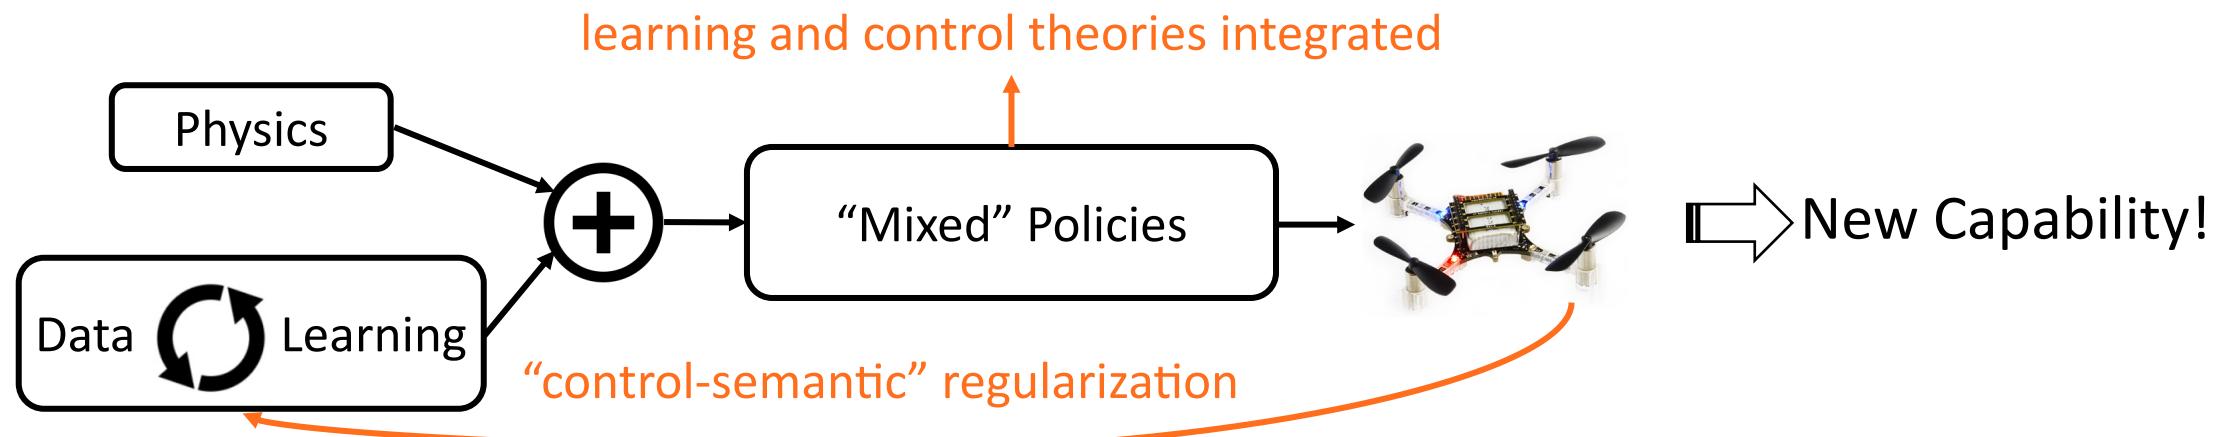
- A typical pipeline in control



- + Formal guarantees
- System Id is often expensive, inaccurate and hard to generalize
- Control synthesis might be too complex or conservative
- A typical pipeline in learning



# This Lecture



- Topic I: dynamics-level residual learning
  - Topic II: action-level residual learning
  - Topic III: program-level residual learning
- \* Not formal classifications. The boundaries are NOT clear!

# Many People Work on this Topic!

## Control Meets Learning

Virtual Seminar Series on the Intersection of Control and Learning



**Control  
Meets  
Learning**

[Edit profile](#)

**Control Meets Learning Virtual Seminars**  
[@ControlMeetsML](#)

Virtual seminar series broadly focused on the intersection of control and learning.  
Organized by [@NavidAzizan](#), [@guannanqu](#), [@GuanyaShi](#), and [@YuanyuanShi2](#).

Control Meets Learning seminar series

Advanced Topics in Machine Learning  
CS 159 · Caltech · Spring 2021

Control Learning

**AA 203: Optimal and Learning-Based Control**  
Spring 2021

**ESE 680, Fall 2019 – Learning and Control**

courses: Caltech CS159, Stanford AA 203, UPenn ESE 680, ...

## Learning for Dynamics and Control (L4DC)

May 30 & 31, 2019 at the Ray and Maria Stata Center  
Massachusetts Institute of Technology, Cambridge, MA

## LEARNING FOR DYNAMICS & CONTROL (L4DC)

Online June 11-12th, 2020

## 3rd Annual Learning for Dynamics & Control Conference

June 7 – 8, 2021 | ETH Zurich, Switzerland Virtual

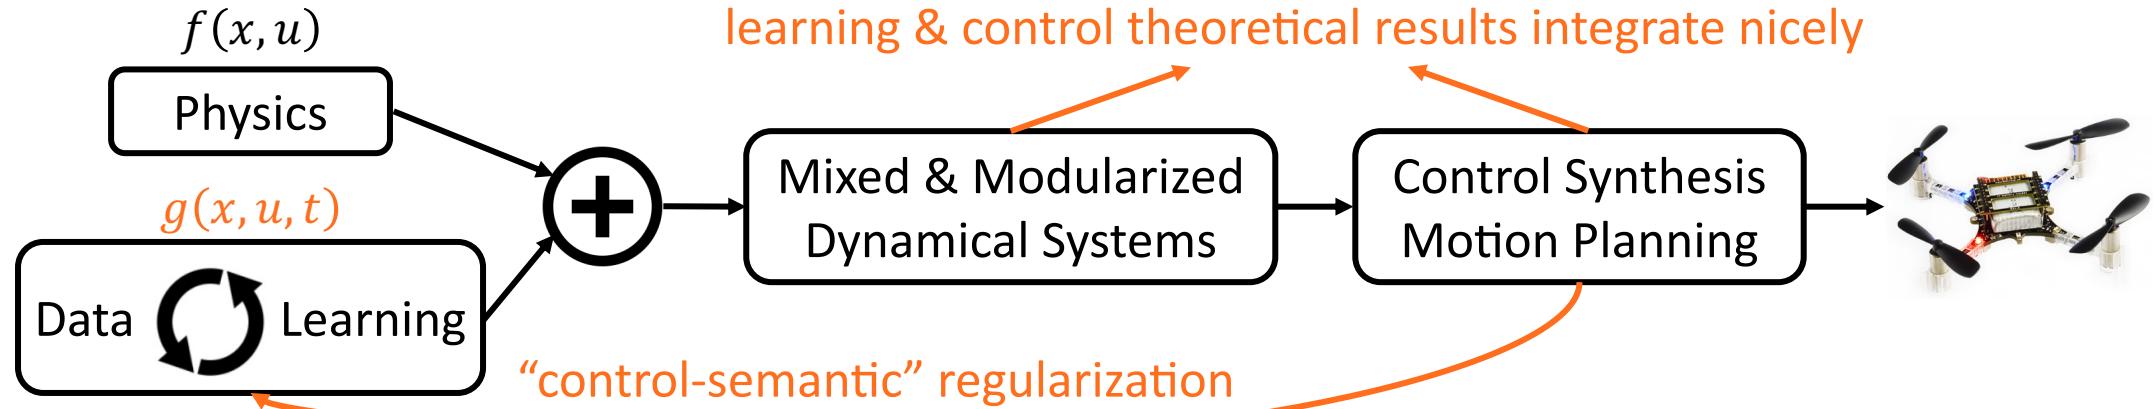
L4DC conference

**Workshop on Learning for Control**  
**57th IEEE Conference on Decision and Control**  
Miami Beach, Florida, December 16, 2018

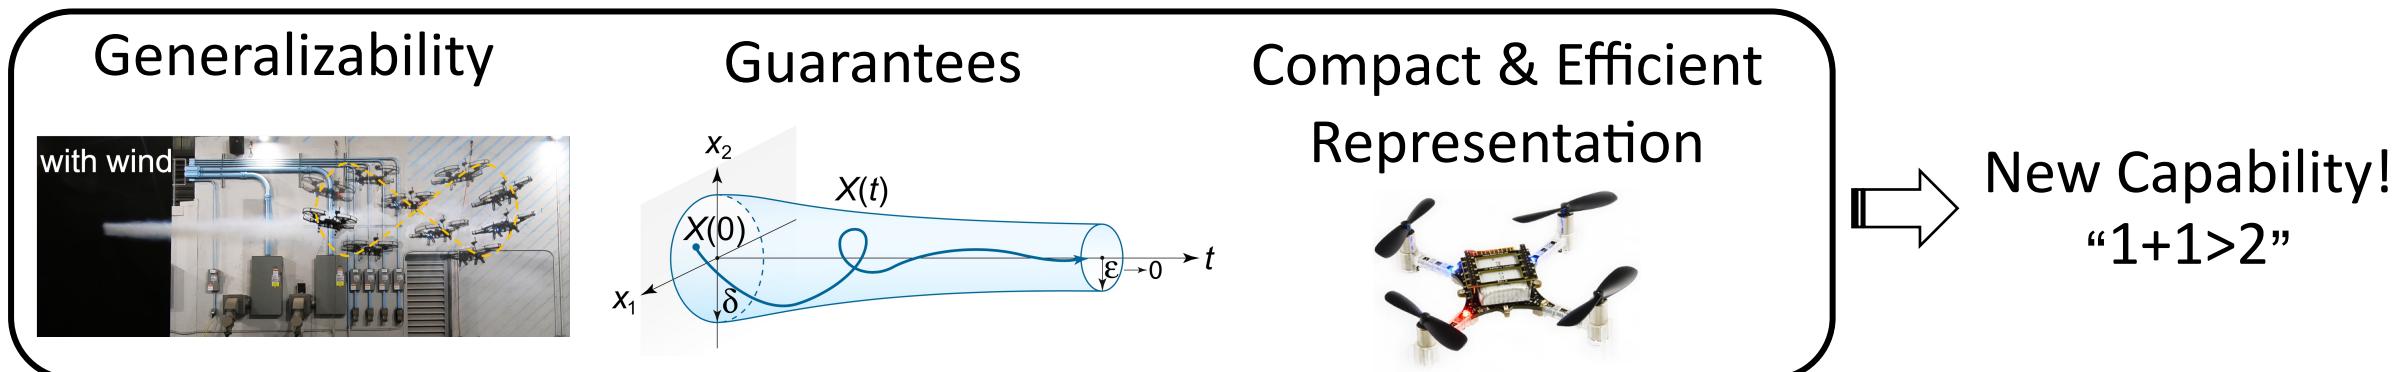
workshops at CDC, NeurIPS, ICML, ICRA, RSS...

# Topic I: Dynamics-level Residual Learning

$$\dot{x} = f(x, u) + \boxed{g(x, u, t)} \text{ unknown}$$



- Key features we pursue:

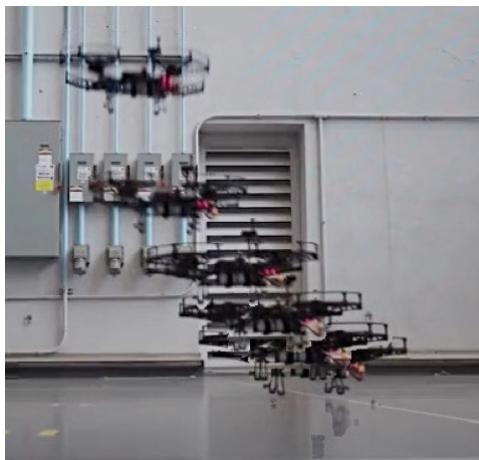


# Neural-Lander Family

$$\dot{x} = f(x, u) + \boxed{g(x, u, t)}$$

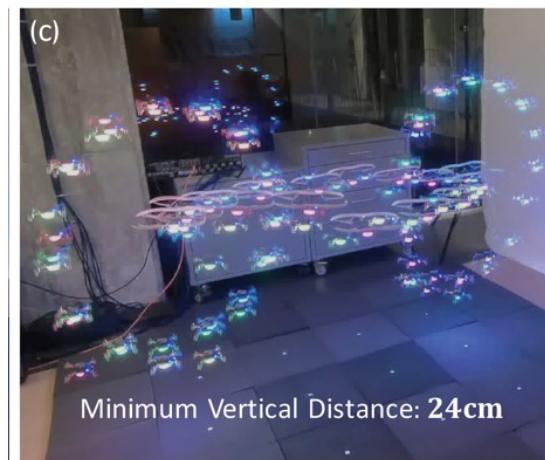
unknown

- (Neural-Lander) learn  $g$  using normalized DNNs
- (Neural-Swarm) learn  $g$  in heterogeneous swarms
- (Neural-Fly) meta-learn  $g$  and online fast adapt
- (Safe Exploration) safe learn  $g$  without experts



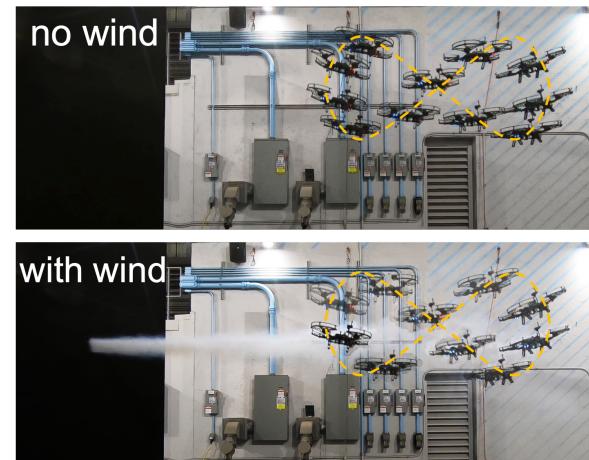
Neural-Lander

*Shi & Shi & O'Connell et al., ICRA 2019*

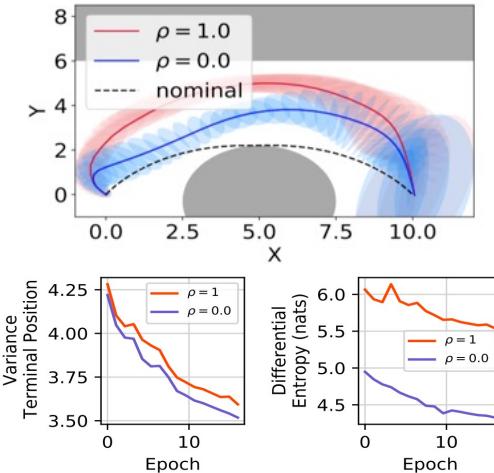


Neural-Swarm and N-Swarm2

*Shi et al., ICRA 2020 and T-RO*



Neural-Fly, *ongoing work*  
*O'Connell & Shi et al., arXiv preprint*



Safe Exploration

*Liu et al., L4DC 2020*

*Nakka et al., RA-L 2020*

# Neural-Lander

$$\dot{x} = f(x, u) + \boxed{g(x, u, t)}$$

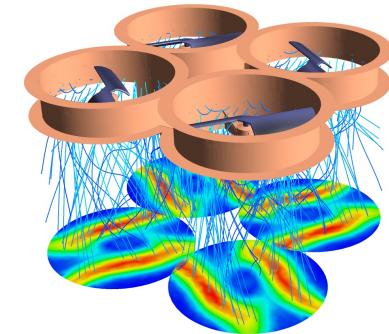
unknown

(Neural-Lander) learn  $g$  using normalized DNNs

(Neural-Swarm) learn  $g$  in heterogeneous swarms

(Neural-Fly) meta-learn  $g$  and online fast adapt

(Safe Exploration) safe learn  $g$  without experts



ground effect

# Quadrotor Dynamics

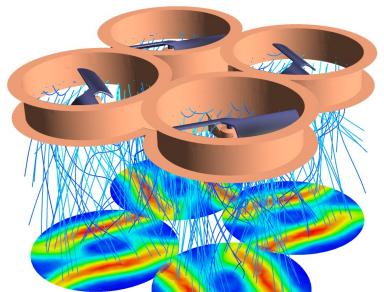
$$\dot{x} = f(x, u) + g(x, u)$$

- $x = [\mathbf{p}, \mathbf{v}, R, \boldsymbol{\omega}], u = [n_1^2, n_2^2, n_3^2, n_4^2]$

- Dynamics

$$\left\{ \begin{array}{l} \dot{\mathbf{p}} = \mathbf{v}, \quad m\dot{\mathbf{v}} = m\mathbf{g} + R\mathbf{f}_u + \mathbf{f}_a \\ \dot{R} = RS(\boldsymbol{\omega}), \quad J\dot{\boldsymbol{\omega}} = J\boldsymbol{\omega} \times \boldsymbol{\omega} + \boldsymbol{\tau}_u + \boldsymbol{\tau}_a \end{array} \right.$$

$$\left\{ \begin{array}{l} \mathbf{f}_u = [0, 0, T]^\top \\ \boldsymbol{\tau}_u = [\tau_x, \tau_y, \tau_z]^\top \\ \begin{bmatrix} T \\ \tau_x \\ \tau_y \\ \tau_z \end{bmatrix} = \begin{bmatrix} c_T & c_T & c_T & -c_T \\ 0 & c_T l_{\text{arm}} & 0 & -c_T l_{\text{arm}} \\ -c_T l_{\text{arm}} & 0 & c_T l_{\text{arm}} & 0 \\ -c_Q & c_Q & -c_Q & c_Q \end{bmatrix} \begin{bmatrix} n_1^2 \\ n_2^2 \\ n_3^2 \\ n_4^2 \end{bmatrix} \end{array} \right.$$



much harder to model!

- $f(x, u)$  is the nominal dynamics
- $g(x, u)$  is the **residual** dynamics (e.g., modelling mismatch, ground effect)

# Controller Design (Sketch)

A feedback-linearization-style nonlinear controller

$$u = \pi(x, x^d, \hat{f}_a(x, u))$$



$\hat{f}_a$  is the learned dynamics (a neural network)

- Challenge I: sample efficiency
  - Learning  $f(x, u) + g(x, u)$  needs **1 hour** flight data
  - Learning the residual  $g(x, u)$  only needs **~5 mins**
- Challenge II: control allocation
  - Non-affine control synthesis problem ( $\hat{f}_a$  explicitly depends on  $u$ )
- Challenge III: stability & robustness in control / generalization in learning

spectrally  
normalized  
DNNs

# Spectral Normalization and Stability Guarantees

- Key idea: use  $u_{t-1}$  in the RHS:

$$u_t = \pi(x_t, x_t^d, \hat{\mathbf{f}}_a(x_t, u_{t-1}))$$

- If the Lipschitz constant of the DNN  $\hat{\mathbf{f}}_q$  is upper bounded by  $\gamma$ , we have

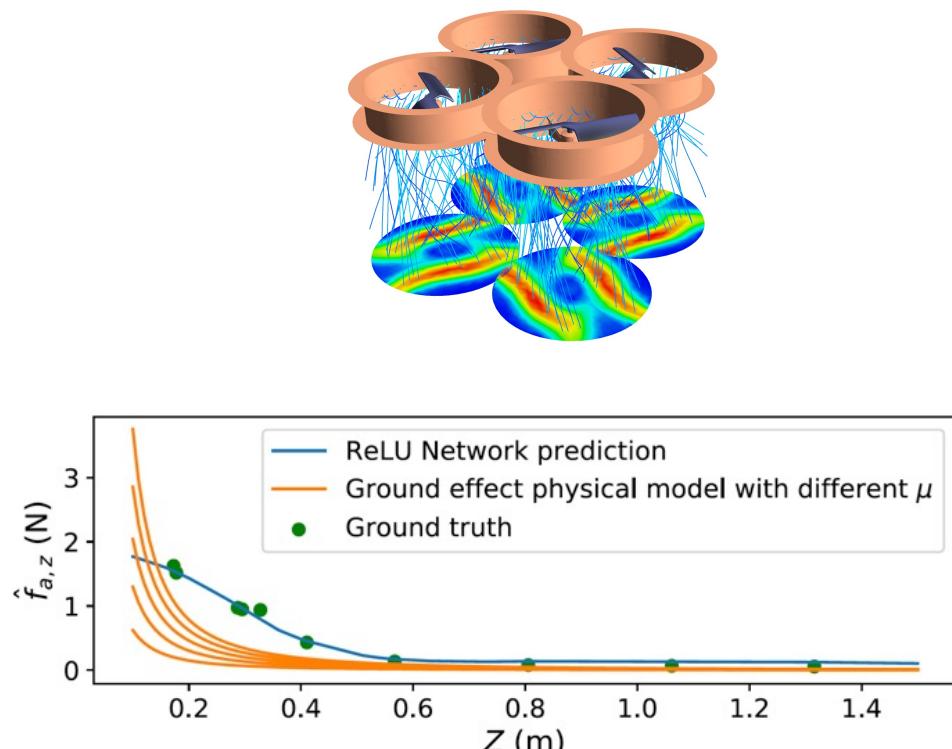
$$\|s(t)\| \leq \|s(t_0)\| \exp\left(-\frac{\lambda - L_a \rho}{m}(t - t_0)\right) + \frac{\epsilon_m}{\lambda - L_a \rho} \rightarrow \text{smoothness of the desired trajectory } x_t^d$$

tracking error    control gain    Lip constant    approximation error of DNN ( $\|\hat{f}_a - f_a\|$ )

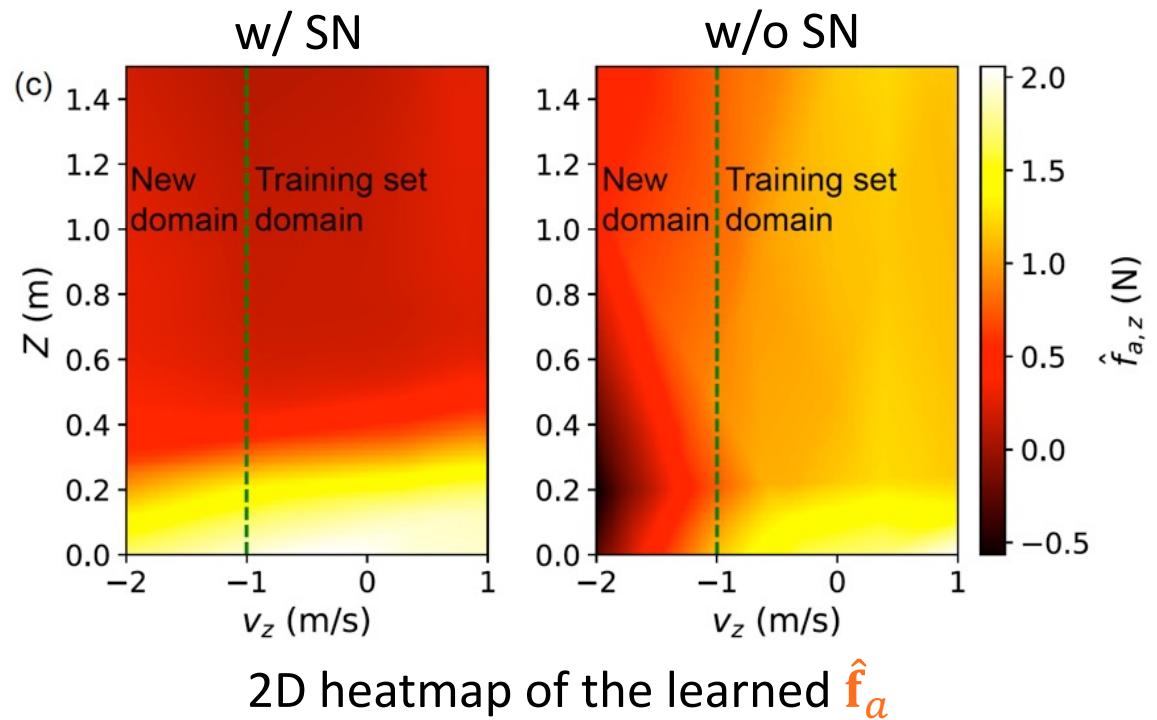
- Thus, we trained the  $\hat{\mathbf{f}}_a$  model using Spectral Normalization (SN)
    - \*SN of DNNs also leads to *good generalization* (stability in a learning-theoretic sense) [Bartlett & Foster & Telgarsky, NeurIPS 2017]

# Training Result Visualization

- Training data: manually fly the drone for 5 minutes

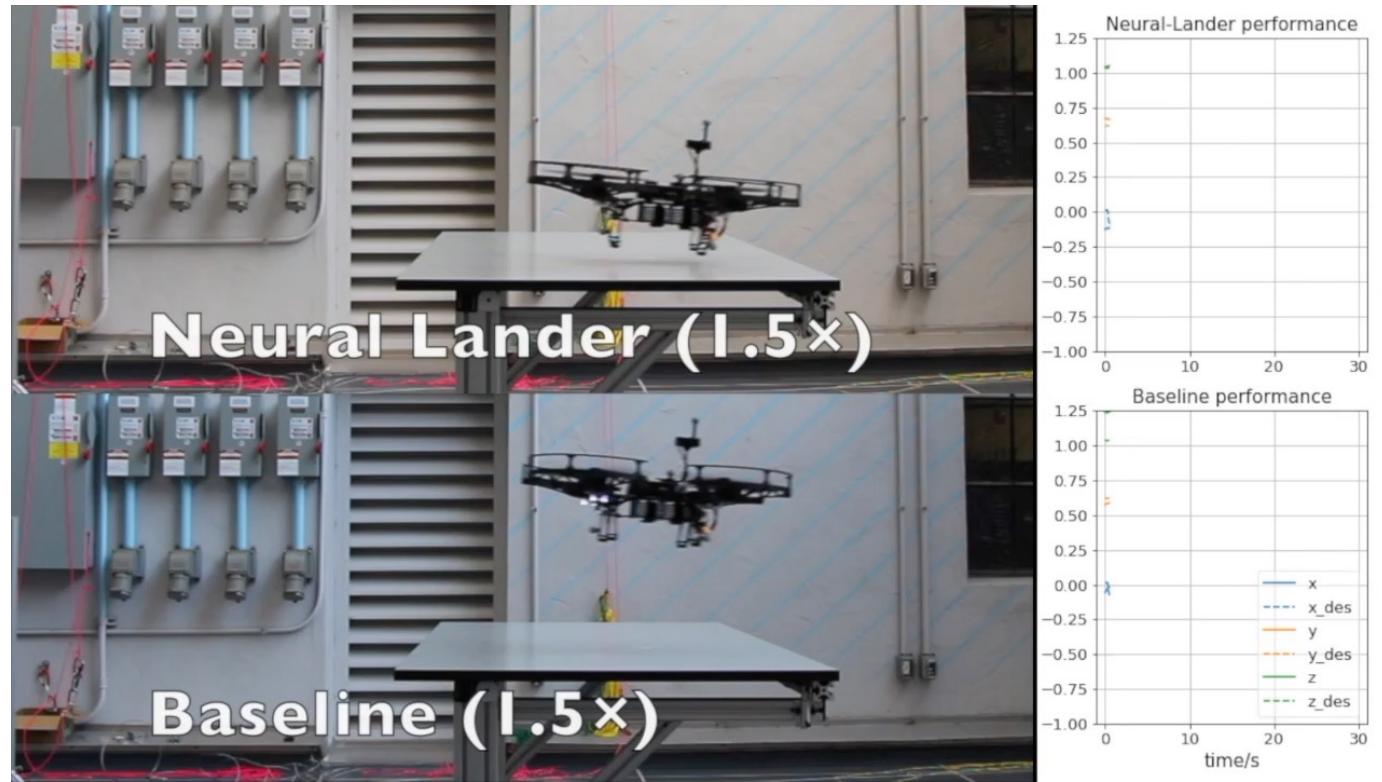
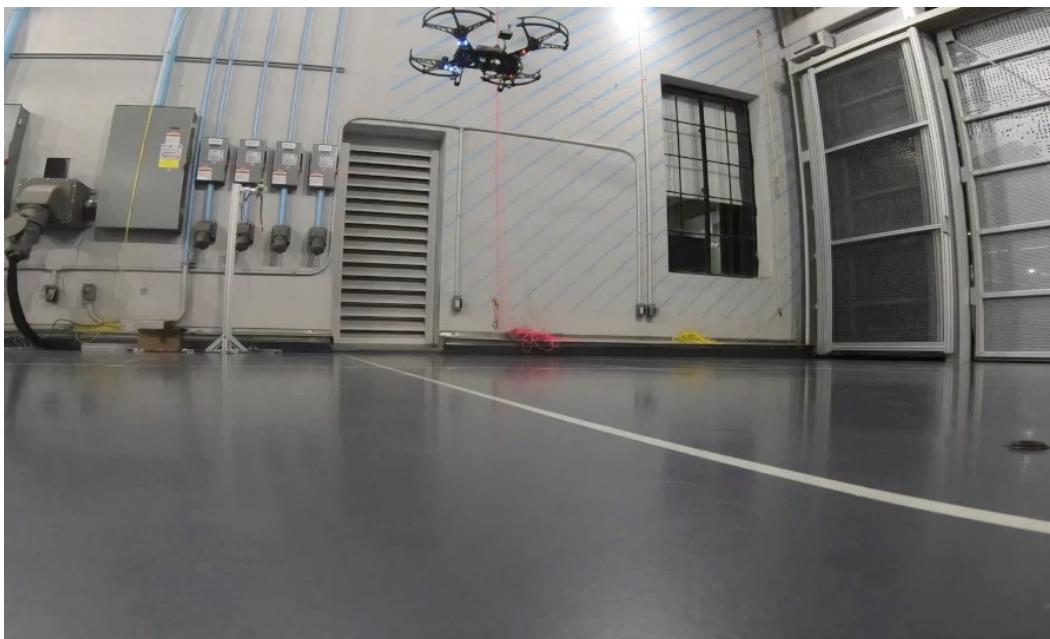


compare with an existing ground effect model



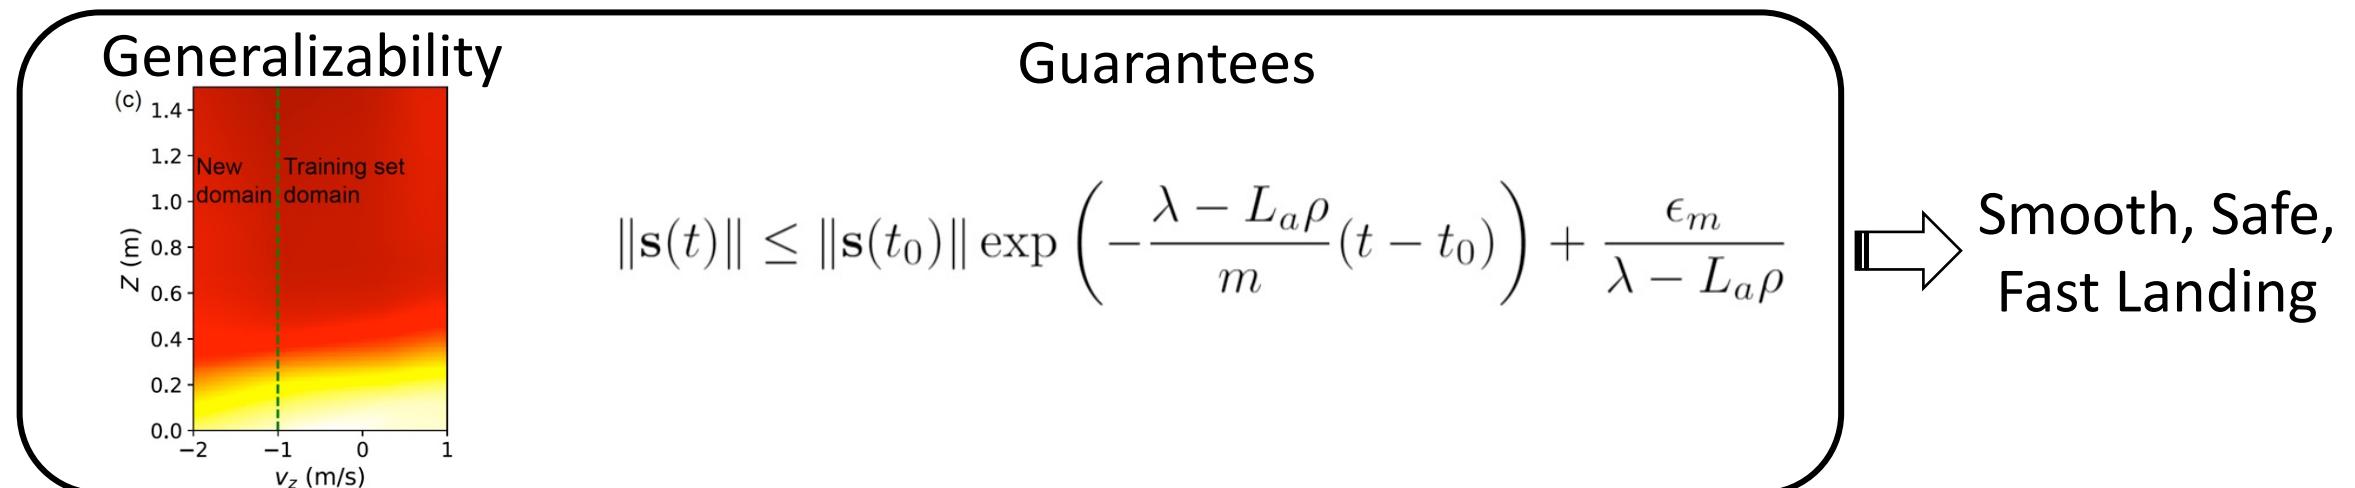
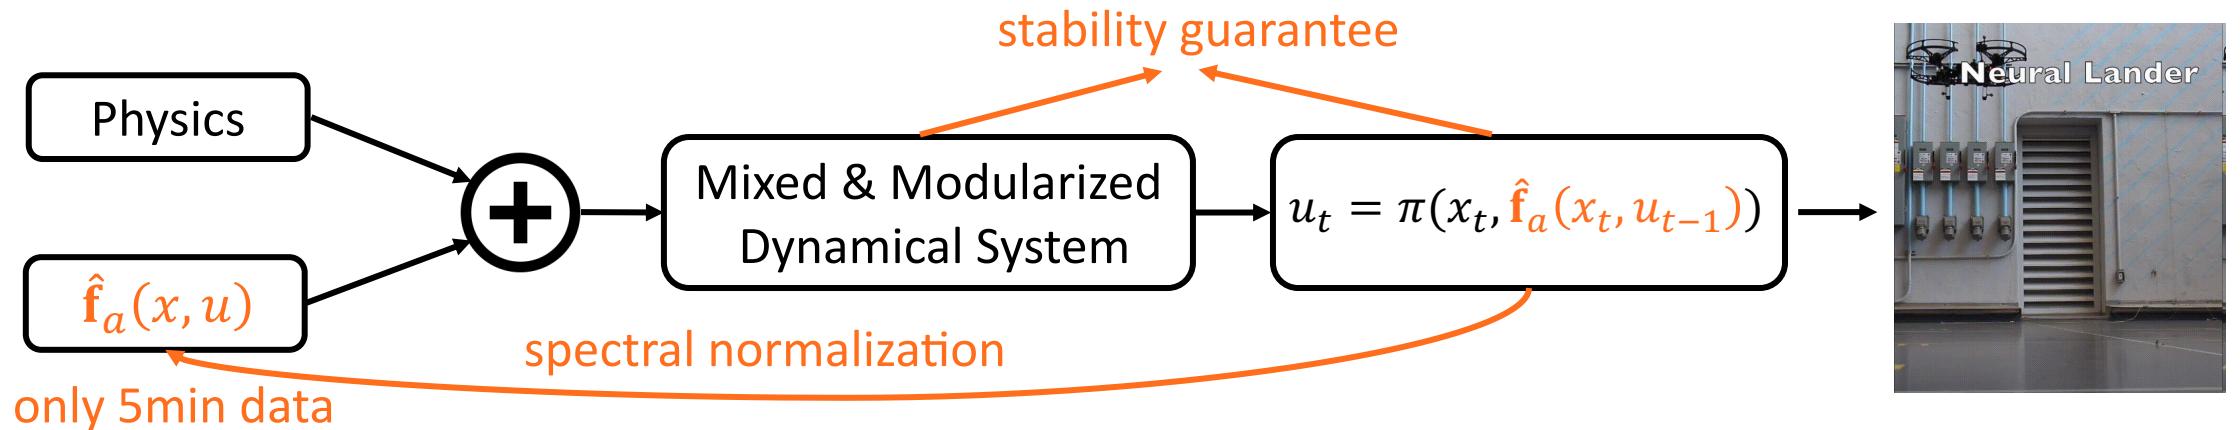
\*Spectrally normalized DNNs can guarantee generalization (*Bartlett et al., NeurIPS 2020; Liu & Shi et al., L4DC 2020*)

# Trajectory Tracking Performance



“table effect”: need to recollect data and  
retrain  $\hat{\mathbf{f}}_a$

# Neural-Lander Takeaways



# Neural-Swarm

$$\dot{x} = f(x, u) + \boxed{g(x, u, t)}$$

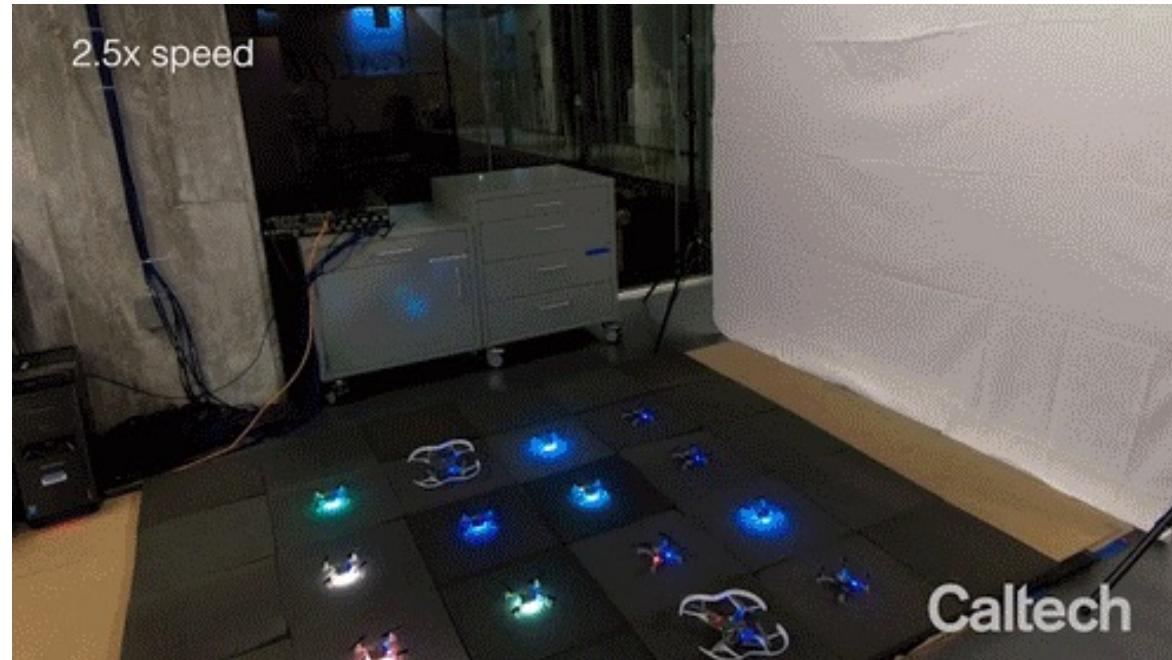
unknown

(Neural-Lander) learn  $g$  using normalized DNNs

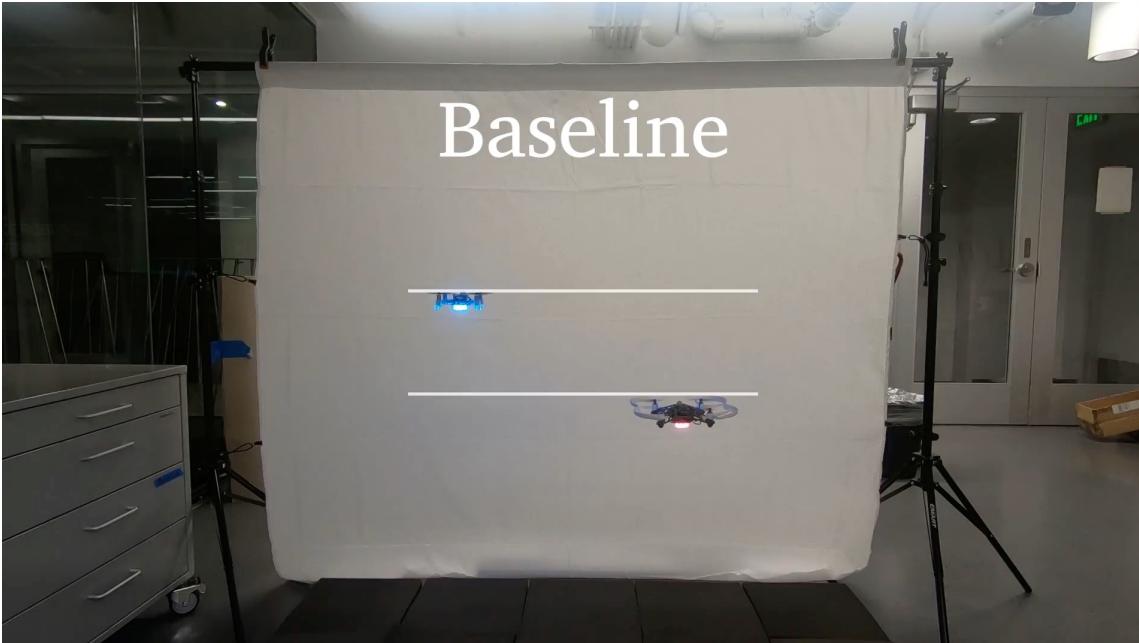
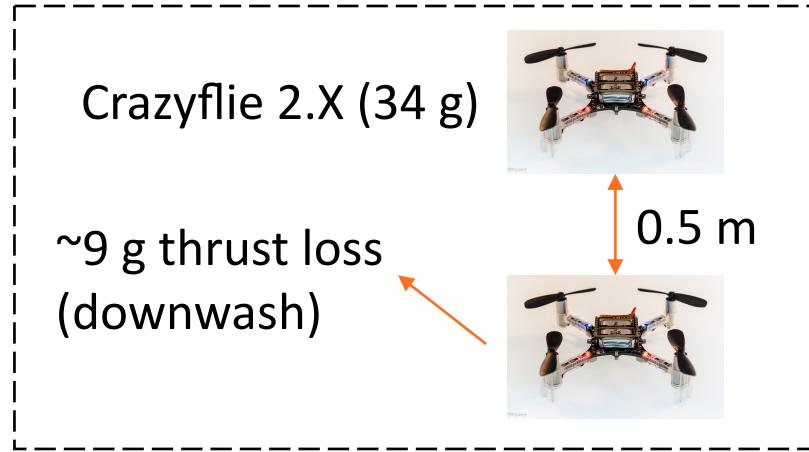
**(Neural-Swarm)** learn  $g$  in heterogeneous swarms

(Neural-Fly) meta-learn  $g$  and online fast adapt

(Safe Exploration) safe learning  $g$  without experts

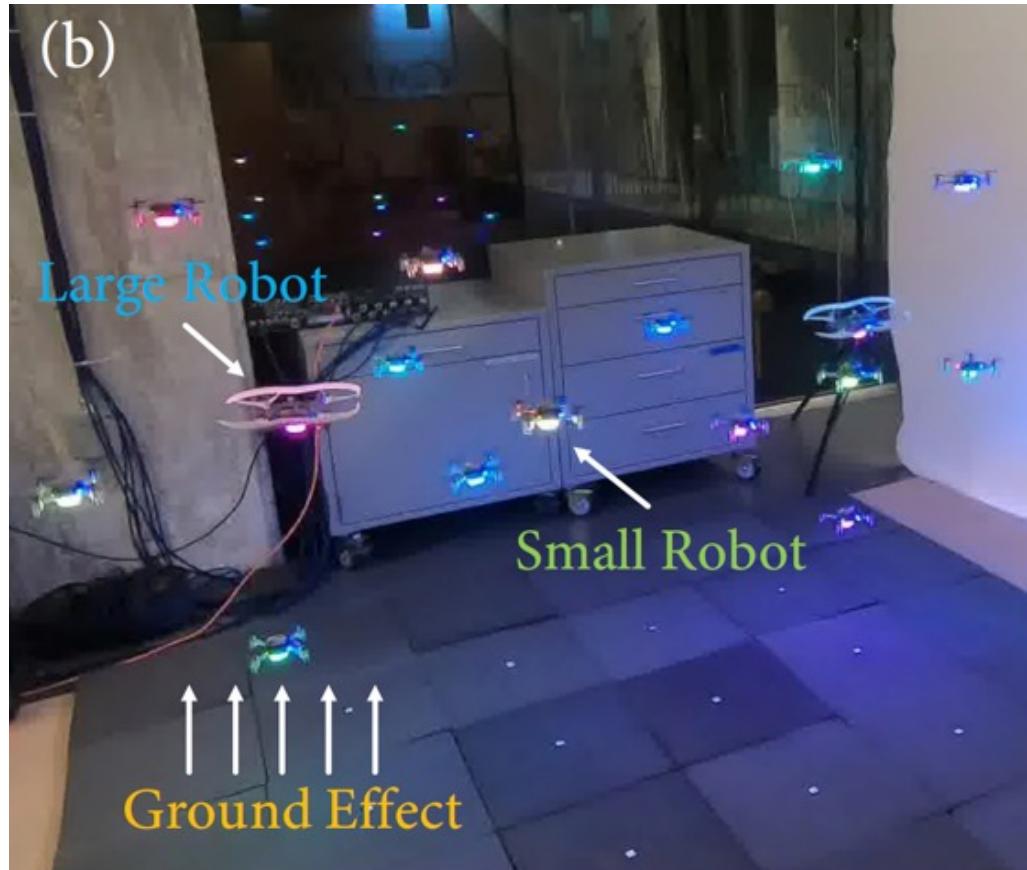


# Motivations



- Interaction matters
  - Prior works require **~60cm** safe vertical distance
- How can we model the interaction in:
  - two and more drones?
  - heterogeneous teams?

# Models and Challenges



14 small robots and 2 large robots

nominal dynamics of type  $I(i)$

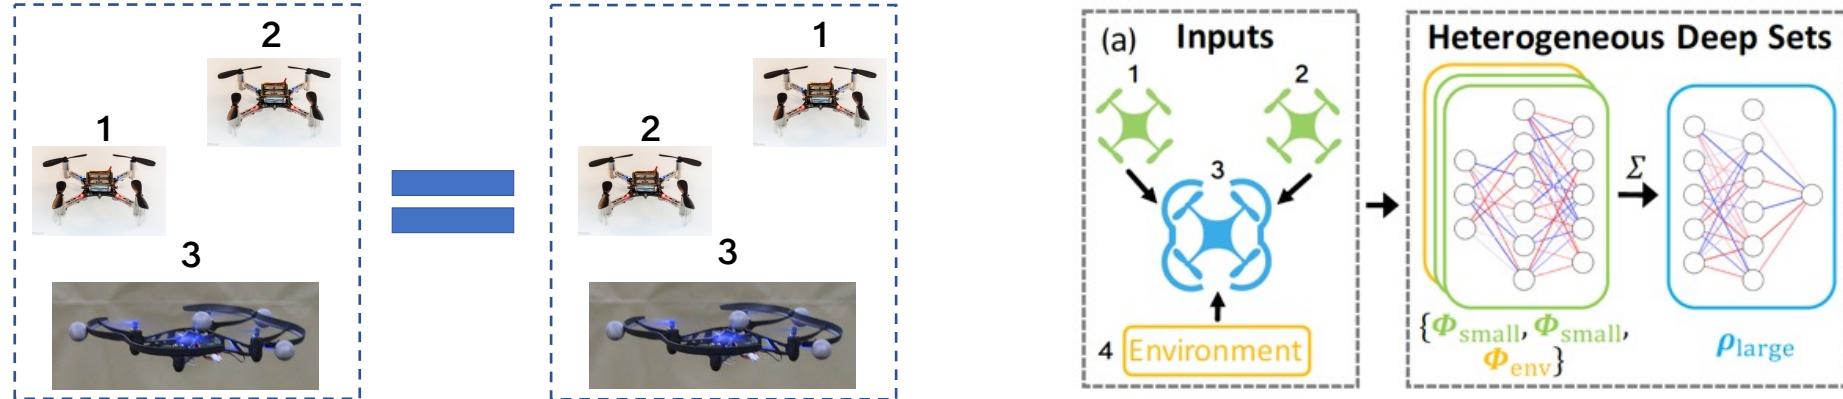
$$\dot{\mathbf{x}}^{(i)} = f_{\mathcal{I}(i)}(\mathbf{x}^{(i)}, \mathbf{u}^{(i)}) + \begin{bmatrix} \mathbf{0} \\ \mathbf{f}_a^{(i)} \\ \mathbf{0} \\ \tau_a^{(i)} \end{bmatrix}, \quad i = 1, \dots, n$$

$$\mathbf{f}_a^{(i)} = \mathbf{f}_a^{\mathcal{I}(i)}(\mathcal{N}_{\text{type}-1}^{(i)}, \dots, \mathcal{N}_{\text{type}-K}^{(i)})$$

$$\tau_a^{(i)} = \tau_a^{\mathcal{I}(i)}(\mathcal{N}_{\text{type}-1}^{(i)}, \dots, \mathcal{N}_{\text{type}-K}^{(i)})$$

- Heterogeneity
- Generalization in the number of robots
  - What if we only have data from 1-3 robots?
- Limited computation power
  - Requires decentralization

# Heterogeneous Permutation Invariance



- Heterogeneous permutation-invariant function:  $h(x, y) = \max\{2x_1, 2x_2\} + \max\{y_1, y_2\}$
- We generalize the “Deep Sets” architecture:

**Theorem** (*Shi et al., accepted by T-RO; generalized Zaheer et al., NeurIPS 2017*)

Any continuous, heterogeneous permutation-invariant function  $[\hat{\mathbf{f}}_a; \hat{\mathbf{t}}_a]$  can be approximated by

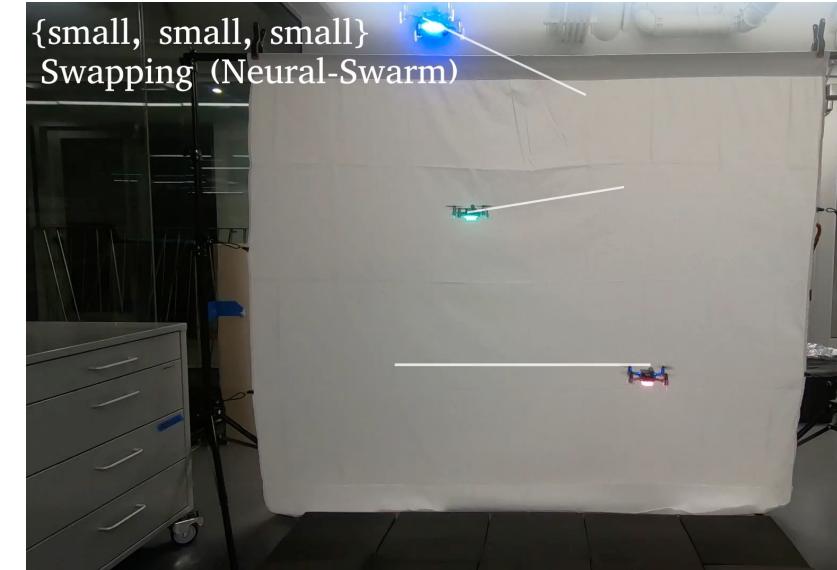
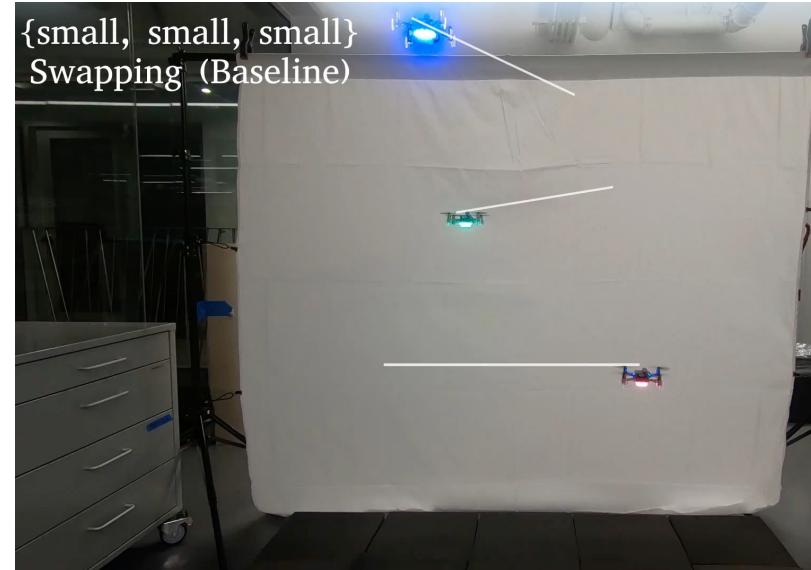
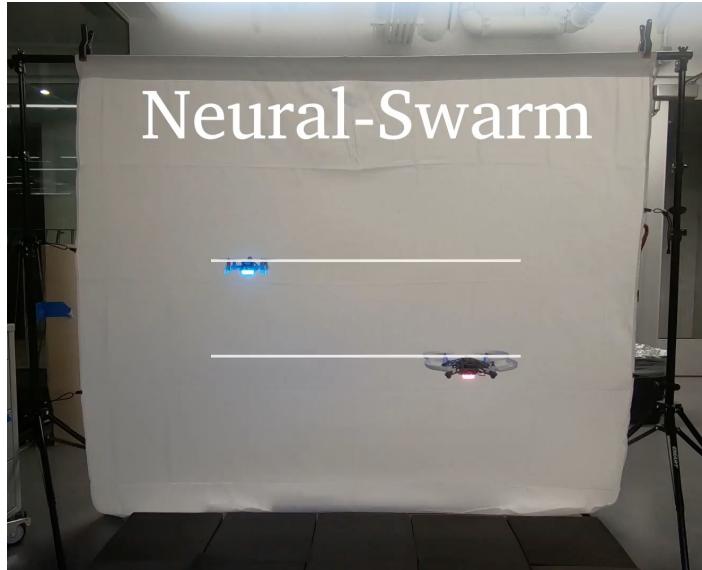
$$\begin{bmatrix} \mathbf{f}_a^{\mathcal{I}(i)}(\mathcal{N}_{\text{type}-1}^{(i)}, \dots, \mathcal{N}_{\text{type}-K}^{(i)}) \\ \mathbf{\tau}_a^{\mathcal{I}(i)}(\mathcal{N}_{\text{type}-1}^{(i)}, \dots, \mathcal{N}_{\text{type}-K}^{(i)}) \end{bmatrix} \approx \rho_{\mathcal{I}(i)} \left( \sum_k \sum_{\mathbf{x}^{(ij)} \in \mathcal{N}_{\text{type}-k}^{(i)}} \phi_{\mathcal{I}(j)}(\mathbf{x}^{(ij)}) \right)$$

superposition in the latent space

Interaction force on type  $I(i)$       decode for type  $I(i)$       encode type  $I(j)$  neighbors

- We only need  $2K$  neural networks for  $K$  types

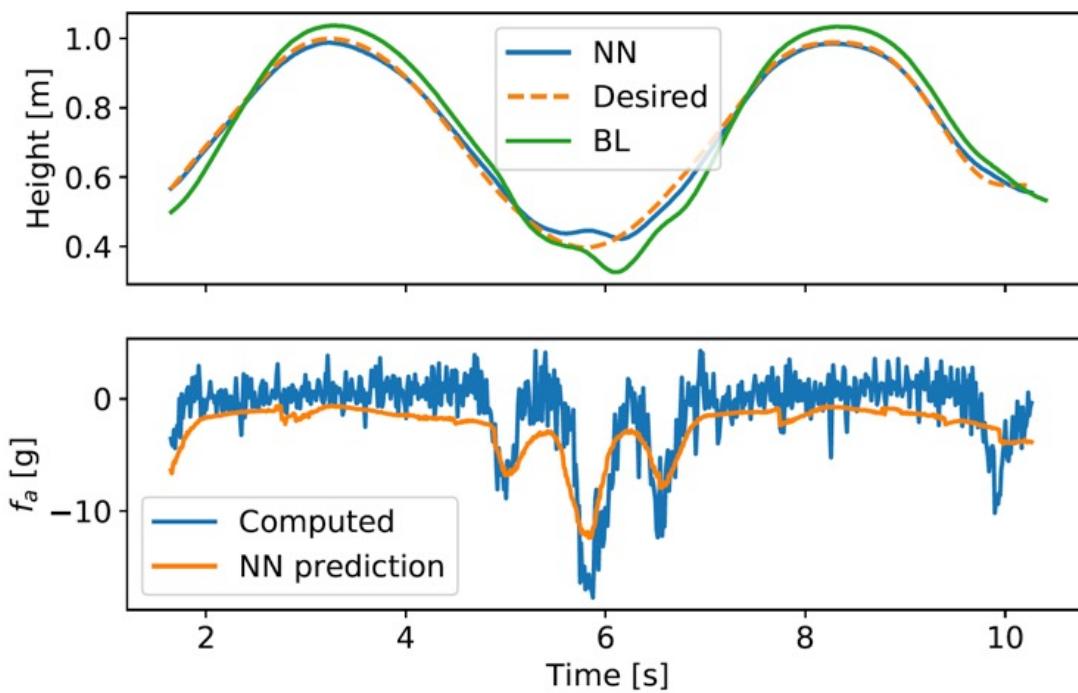
# Use-Case I: Control



- Baseline (BL) is a SOTA nonlinear controller with delay compensation

# Generalization

- We only collected 1-3 robots' data in training
- Can it generalize?



5-robot swapping task (2 larges, 3 smalls)



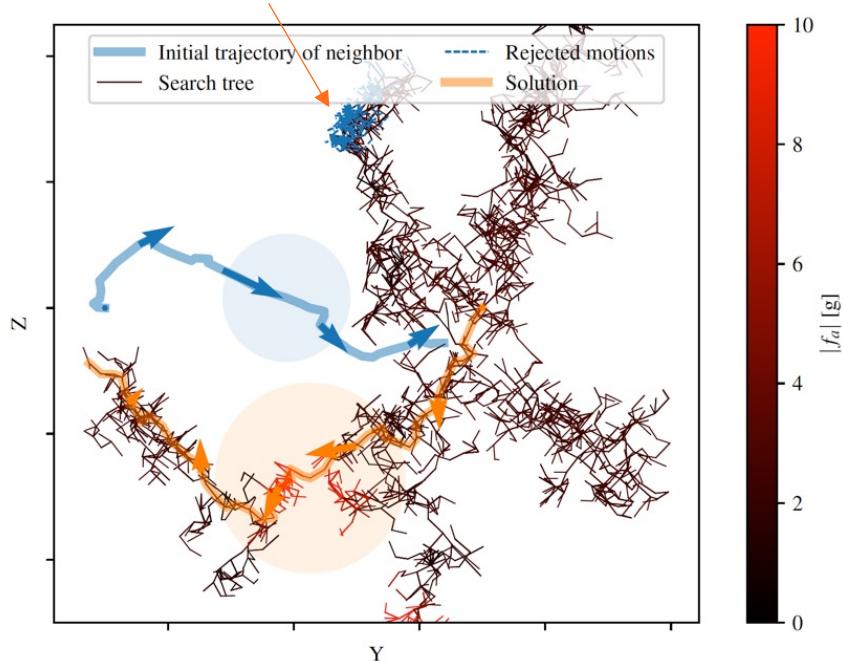
16-robot 3-ring tracking task (2 larges, 14 smalls)

- **24cm** minimum vertical distance
- Prior works: **~60cm** in a homogeneous team

# Use-Case II: Interaction-Aware Motion Planning

- Integrate learned interactions with a **two-stage** planner
- Stage I: AO-RRT type sampling-based planning with learned interactions
- Stage II: Optimal-control-based planning using Sequential Convex Programming (SCP)

Rejected because introducing too large force



**Stage I:** AO-RRT type planning

- Fix the blue robot and plan for the orange

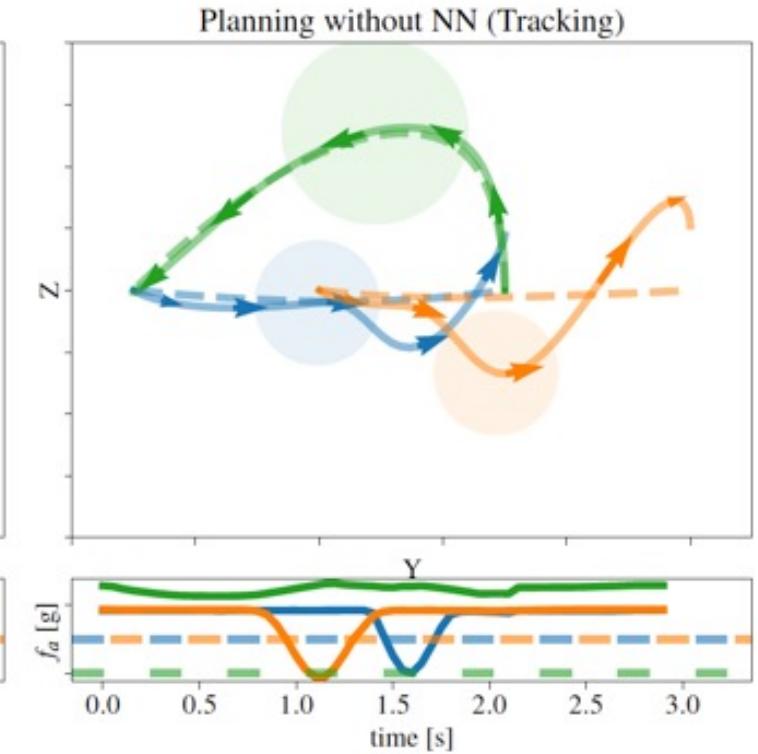
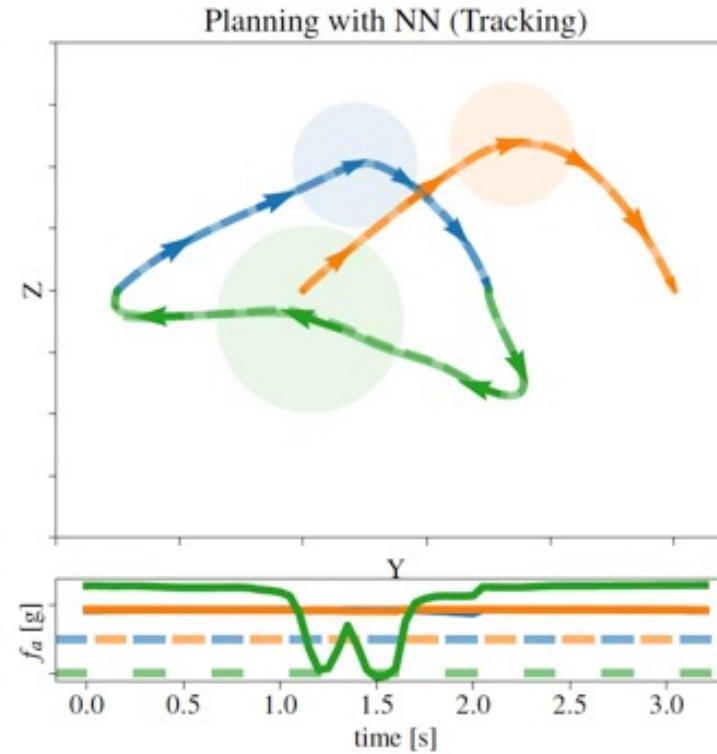
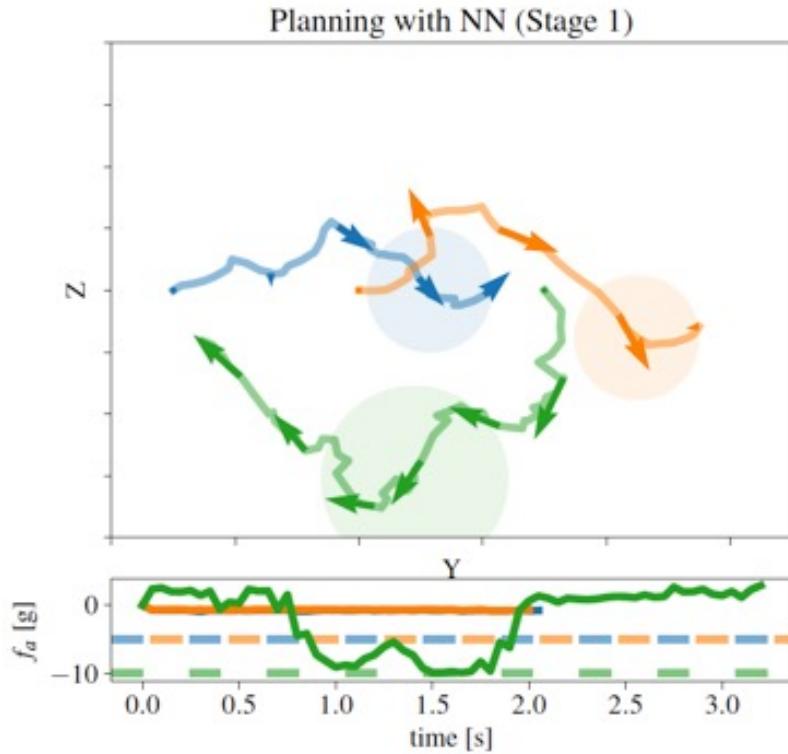
$$\min_{\mathbf{u}^{(i)}, \mathbf{x}^{(i)}, t_f} \sum_{i=1}^N \int_0^{t_f} \|\mathbf{u}^{(i)}(t)\| dt \quad (6)$$

$$\text{s.t. } \begin{cases} \text{robot dynamics (4)} & i \in [1, N] \\ \mathbf{u}^{(i)}(t) \in \mathcal{U}^{\mathcal{I}(i)}; \mathbf{x}^{(i)}(t) \in \mathcal{X}^{\mathcal{I}(i)} & i \in [1, N] \\ \|\mathbf{p}^{(ij)}\| \geq r^{(\mathcal{I}(i)\mathcal{I}(j))} & i < j, j \in [2, N] \\ \|\mathbf{f}_a^{(i)}\| \leq f_{a,\max}^{\mathcal{I}(i)}; \|\boldsymbol{\tau}_a^{(i)}\| \leq \tau_{a,\max}^{\mathcal{I}(i)} & i \in [1, N] \\ \mathbf{x}^{(i)}(0) = \mathbf{x}_s^{(i)}; \mathbf{x}^{(i)}(t_f) = \mathbf{x}_f^{(i)} & i \in [1, N] \end{cases}$$

**Stage II:** optimal control

- Nonconvex so we use SCP
- We can explicitly control the interaction magnitude!

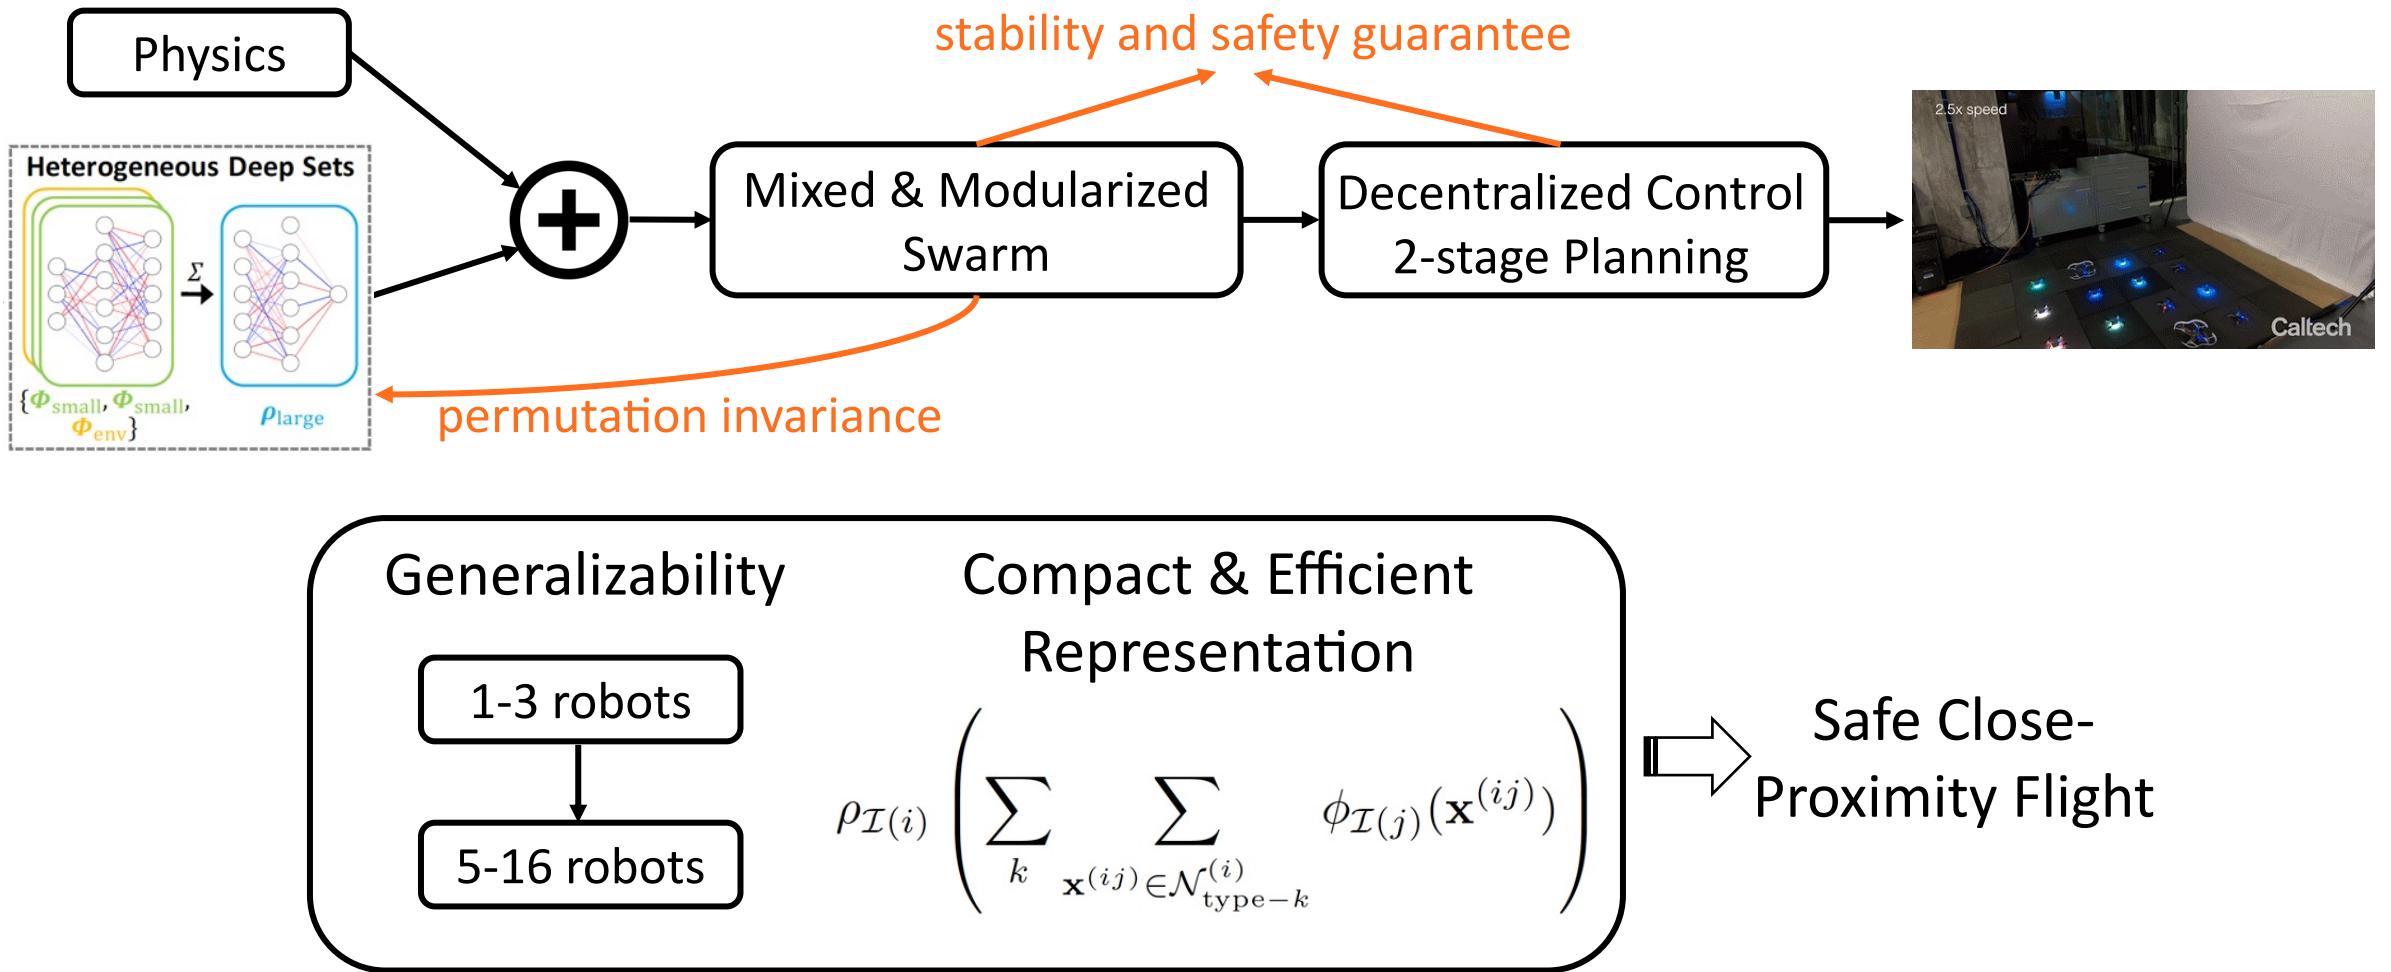
# Use-Case II: Interaction-Aware Motion Planning



The large robot is at the bottom

The large robot is at the top;  
constraint violations

# Neural-Swarm Takeaways



# Neural-Fly (Ongoing Work)

$$\dot{x} = f(x, u) + \boxed{g(x, u, t)}$$

unknown

(Neural-Lander) learn  $g$  using normalized DNNs

(Neural-Swarm) learn  $g$  in heterogeneous swarms

**(Neural-Fly)** meta-learn  $g$  and online fast adapt

(Safe Exploration) safe learning  $g$  without experts

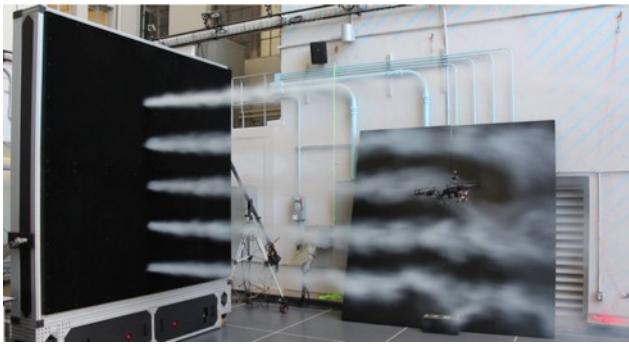


# Meta-Learning Meets Adaptive Control

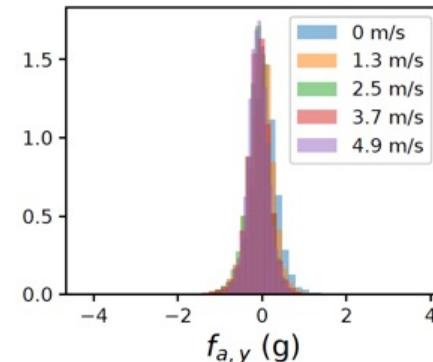
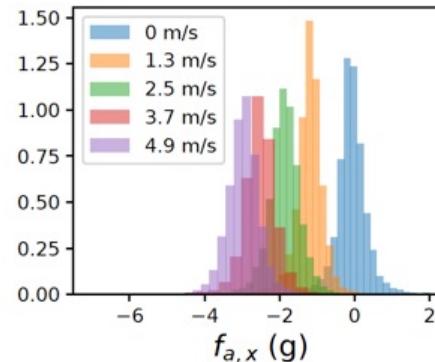
- The unknown residual term is governed by the environment  $c(t)$ :

$$\dot{x} = f(x, u) + g(x, c(t))$$

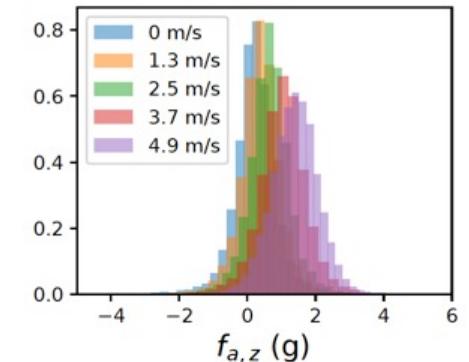
- Example: drone flying in different winds



Caltech CAST fan wall



histogram of  $g$  with different  $c$



- Key idea:

$$g(x, c) \approx \phi(x; \theta)a(c)$$

representations to be meta-learned

a latent state to be online adapted  
(using adaptive control theory)

# Meta-Learning Meets Adaptive Control

- With the meta-learned representation  $\phi(x; \theta)$ , adaptive control comes into play

$$a_{t+1} = \text{adapt}(a_t, \phi, x_t, \dots)$$

- Control stability and robustness can be guaranteed

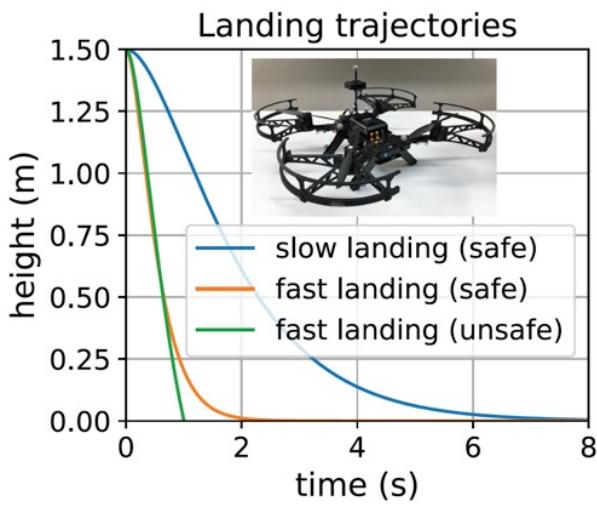


# Safe Exploration

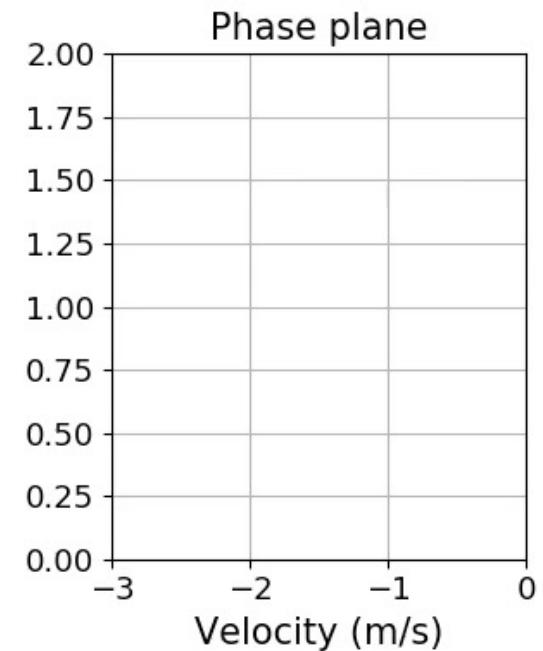
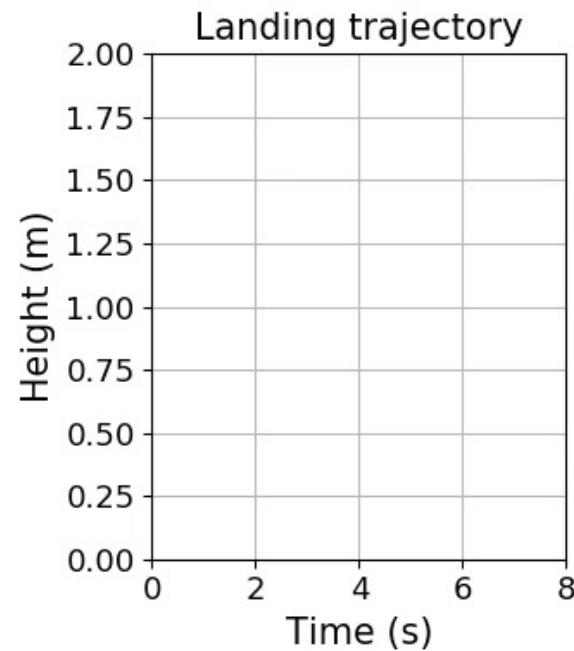
$$\dot{x} = f(x, u) + \boxed{g(x, u, t)}$$

unknown

- (Neural-Lander) learn  $g$  using normalized DNNs
- (Neural-Swarm) learn  $g$  in heterogeneous swarms
- (Neural-Fly) meta-learn  $g$  and online fast adapt
- (Safe Exploration)** safe learning  $g$  without experts



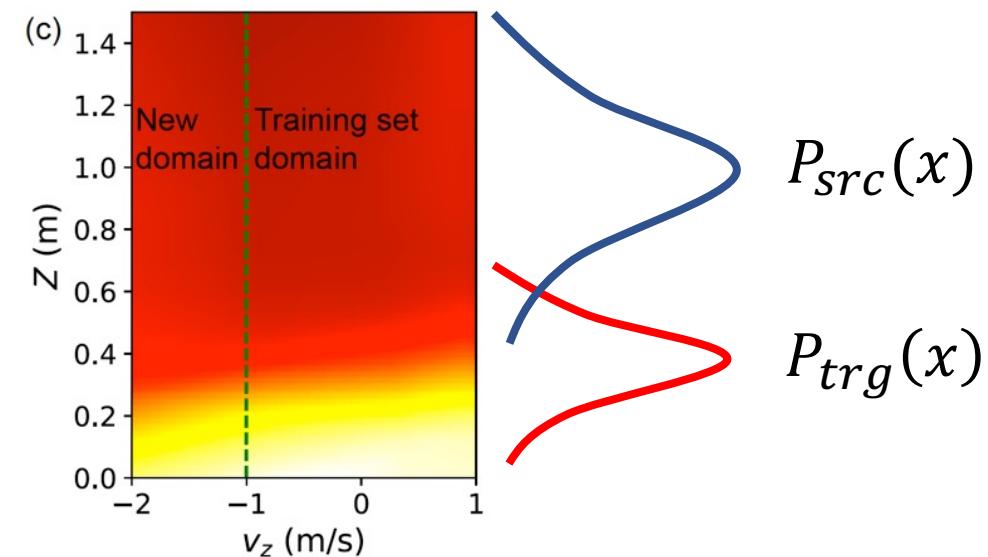
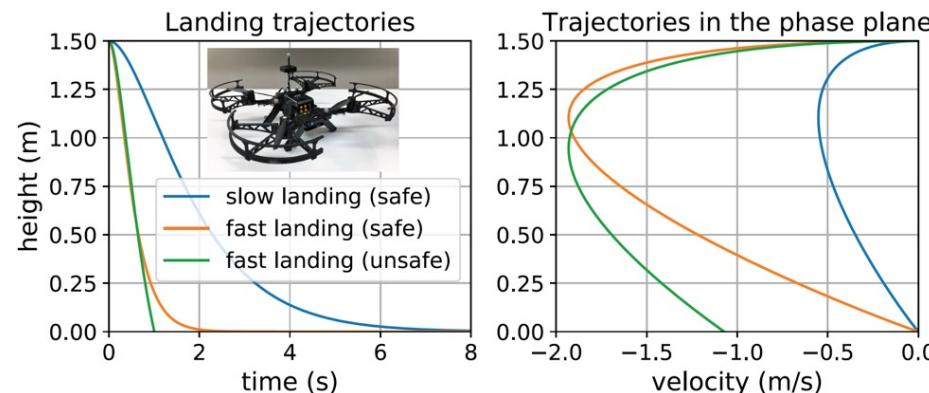
Animation  
time: 0.01 s



# Motivations and Challenges

- How to collect data without experts?
- Exploration v.s. exploitation: how to quantify uncertainty under domain shift?

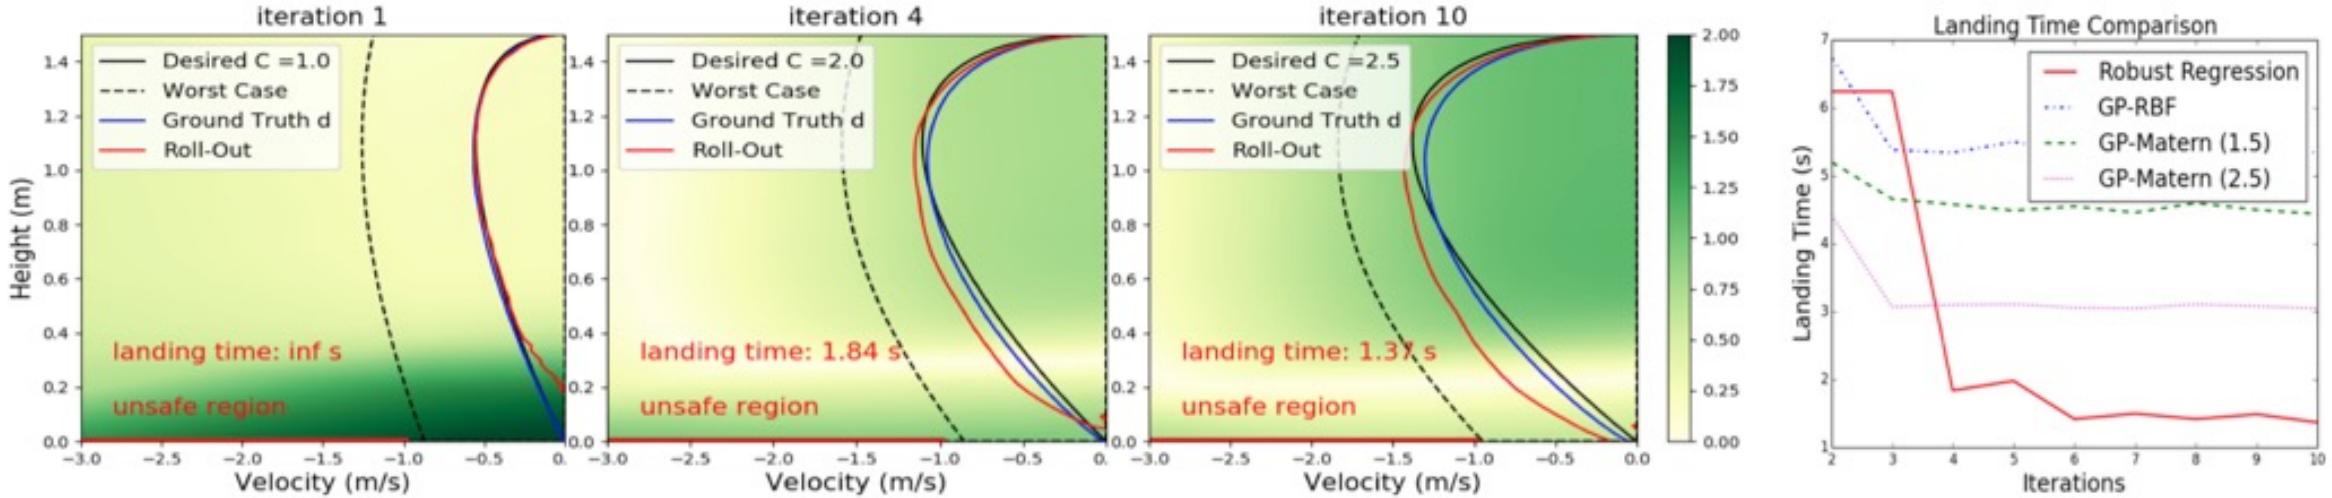
covariate shift:  $P(y|x)$  is fixed, but  $P_{trg}(x) \neq P_{src}(x)$



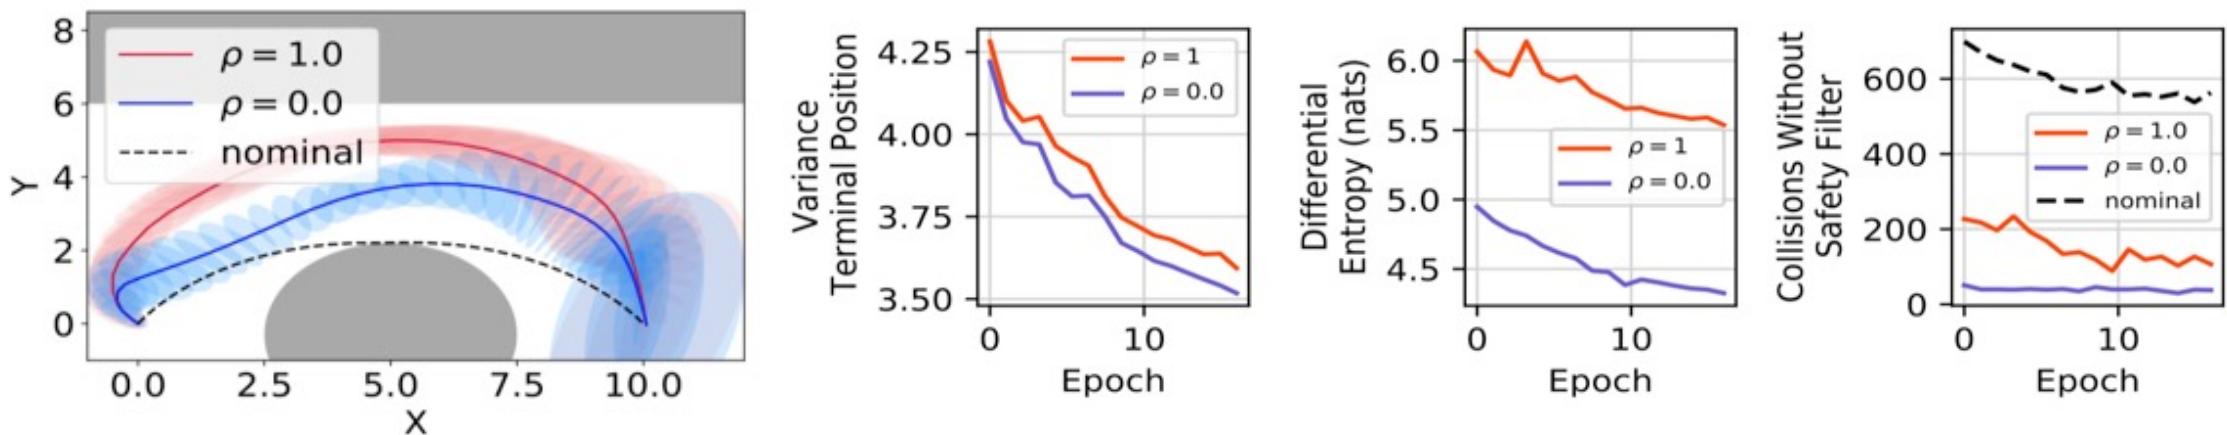
- Key idea: using **robust regression** to quantify uncertainty under covariate shift

# Results

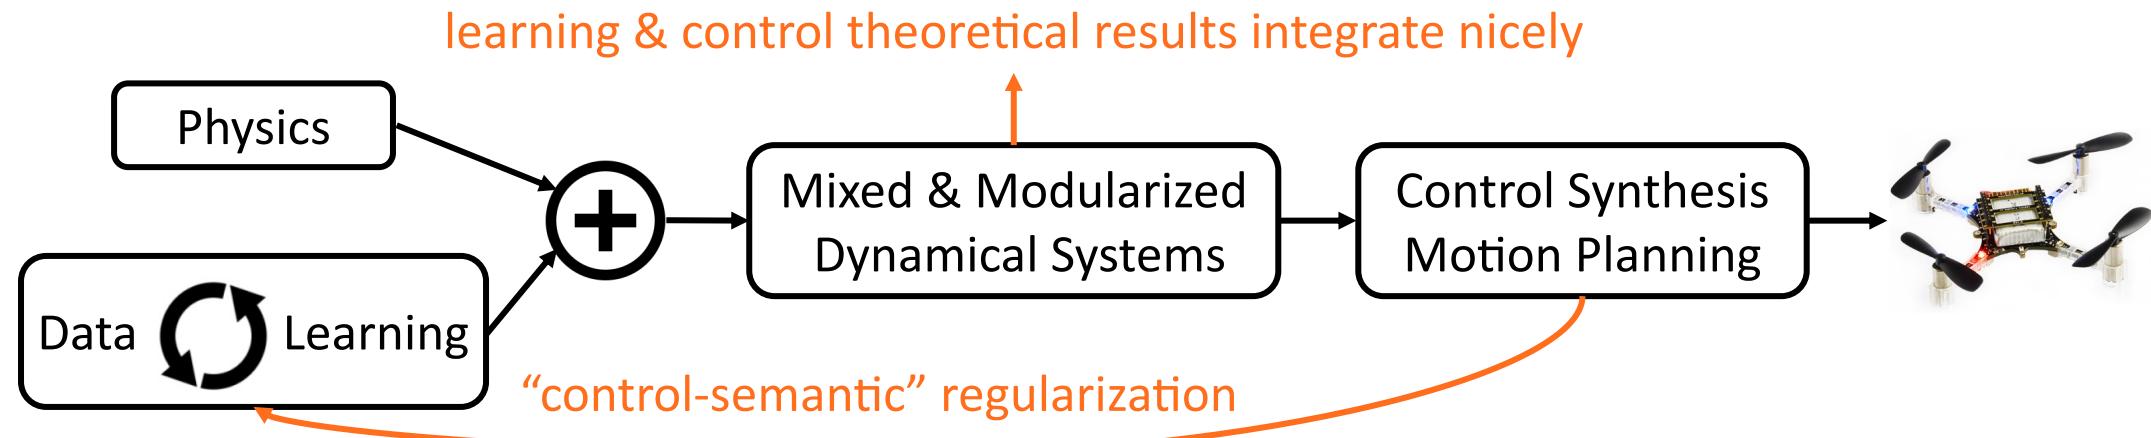
- Deterministic safety constraints and planning in a trajectory pool [Liu et al., L4DC 2020]



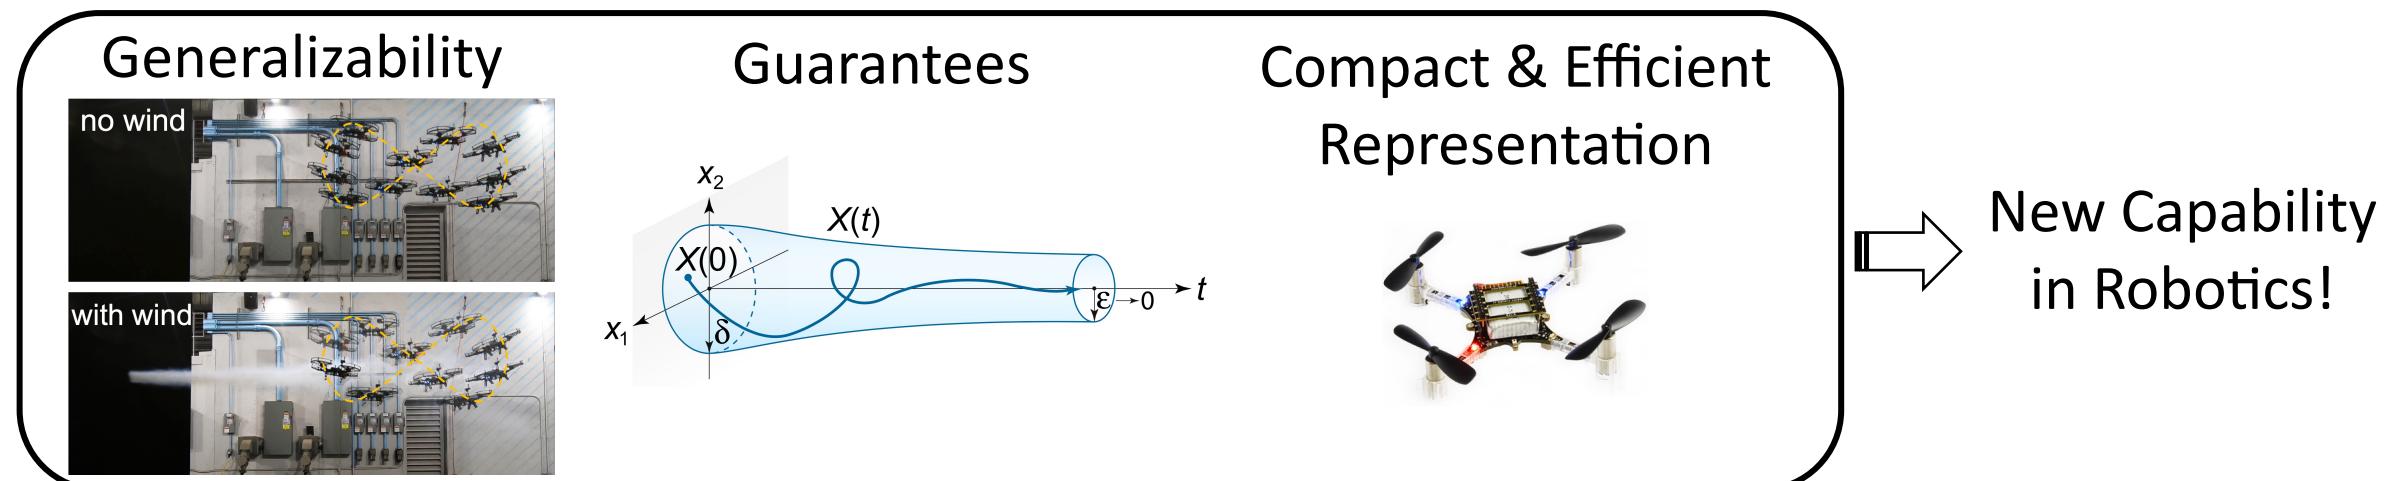
- Information-cost stochastic optimal control with chance-constraints [Nakka et al., RA-L 2020]



# Recap for Topic I: Dynamics-Level Residual Learning



- Key features we pursue:



# Some Other “Control-semantic” Regularization

- Learning stabilizable dynamics

Learning Stabilizable Nonlinear Dynamics  
with Contraction-Based Regularization

Sumeet Singh<sup>1</sup>, Spencer M. Richards<sup>1</sup>, Vikas Sindhwani<sup>2</sup>, Jean-Jacques E. Slotine<sup>3</sup>, and  
Marco Pavone<sup>1</sup>

$$\begin{aligned} \min_{\hat{f} \in \mathcal{H}} \quad & \sum_{i=1}^N \left\| \hat{f}(x_i, u_i) - \dot{x}_i \right\|_2^2 + \mu \|\hat{f}\|_{\mathcal{H}}^2 \\ \text{s.t.} \quad & \hat{f} \text{ is stabilizable,} \end{aligned}$$

- Learning Lagrangian systems

DEEP LAGRANGIAN NETWORKS:  
USING PHYSICS AS MODEL PRIOR FOR DEEP LEARNING

Michael Lutter, Christian Ritter & Jan Peters \*

covering all (rigid) robotic systems

$$\underbrace{\mathbf{H}(\mathbf{q})\ddot{\mathbf{q}} + \dot{\mathbf{H}}(\mathbf{q})\dot{\mathbf{q}} - \frac{1}{2} \left( \frac{\partial}{\partial \mathbf{q}} \left( \dot{\mathbf{q}}^T \mathbf{H}(\mathbf{q}) \dot{\mathbf{q}} \right) \right)^T}_{:= \mathbf{c}(\mathbf{q}, \dot{\mathbf{q}})} + \mathbf{g}(\mathbf{q}) = \boldsymbol{\tau}$$

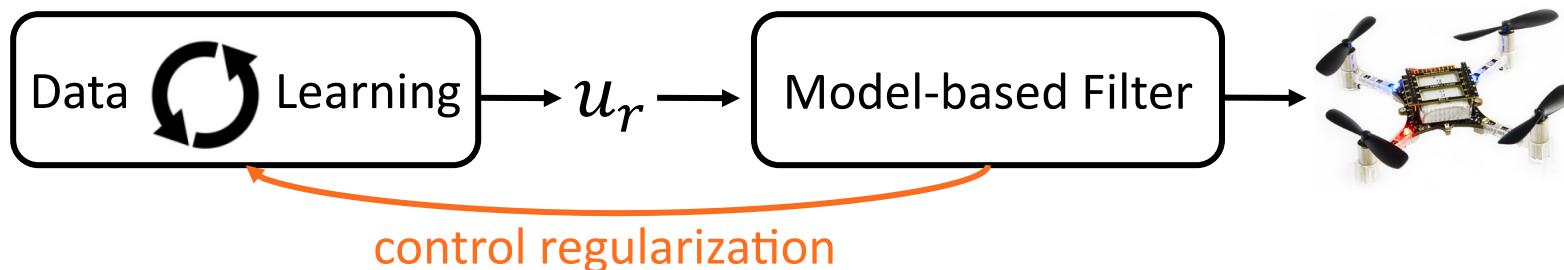
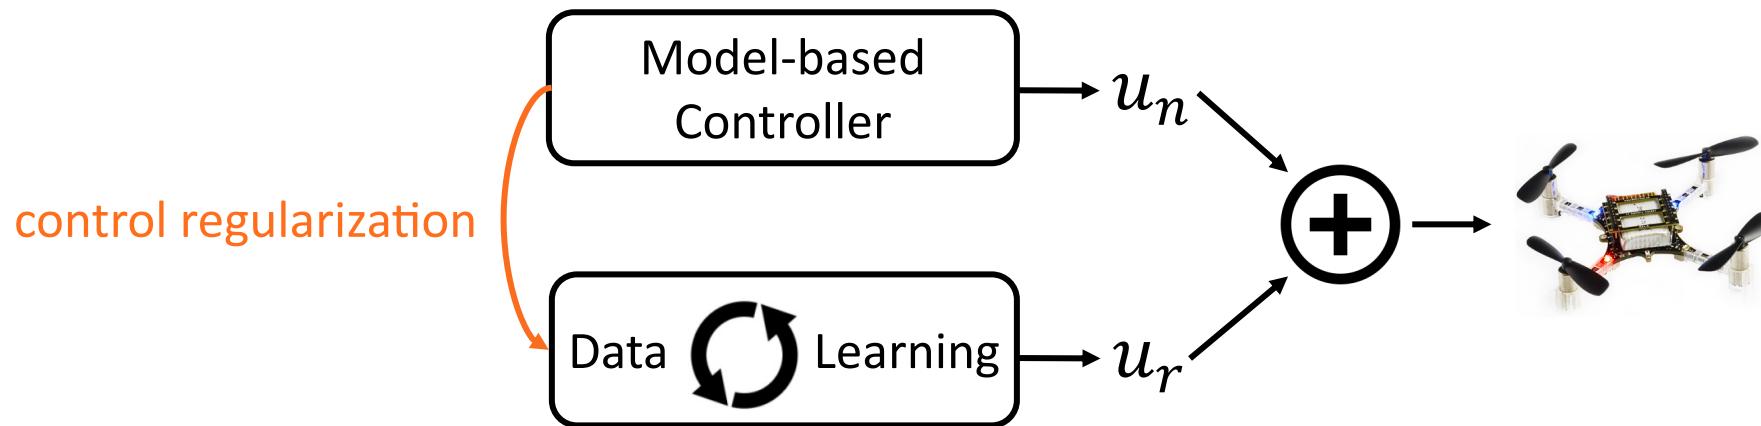
- Could be either hard constraints or regularizations

# References in Topic I: Dynamics-Level Residual Learning

- Zeng, A., Song, S., Lee, J., Rodriguez, A., & Funkhouser, T. (2020). Tossingbot: Learning to throw arbitrary objects with residual physics. *IEEE Transactions on Robotics*
- G. Shi\*, X. Shi\*, M. O'Connell\*, R. Yu, K. Azizzadenesheli, A. Anandkumar, Y. Yue, S.-J. Chung, "Neural Lander: Stable Drone Landing Control Using Learned Dynamics", *ICRA 2019*
- G. Shi, W. Hoenig, Y. Yue, S.-J. Chung, "Neural-Swarm: Decentralized Close-Proximity Multirotor Control Using Learned Interactions", *ICRA 2020*
- G. Shi, W. Hoenig, X. Shi, Y. Yue, S.-J. Chung, "Neural-Swarm2: Planning and Control of Heterogeneous Multirotor Swarms Using Learned Interactions", accepted by *T-RO*, 2021
- M. O'Connell\*, G. Shi, X. Shi, S.-J. Chung, "Meta-Learning-Based Robust Adaptive Flight Control Under Uncertain Wind Conditions", *arXiv preprint*
- A. Liu, G. Shi, S.-J. Chung, A. Anandkumar, Y. Yue, "Robust Regression for Safe Exploration in Control", *L4DC 2020*
- Y. K. Nakka, A. Liu, G. Shi, A. Anandkumar, Y. Yue, S.-J. Chung, "Chance-Constrained Trajectory Optimization for Safe Exploration and Learning of Nonlinear Systems", *RA-L 2020*
- P. Bartlett, D. Foster, D. Telgarsky, "Spectrally-normalized Margin Bounds for Neural Networks", *NeurIPS 2017*
- M. Zaheer, S. Kottur, S. Ravanbakhsh, B. Póczos, R. Salakhutdinov, A. J. Smola, "Deep Sets", *NeurIPS 2017*
- C. Finn, P. Abbeel, S. Levine, "Model-Agnostic Meta-Learning for Fast Adaptation of Deep Networks", *ICML 2017*
- H. Li, Z. Xu, G. Taylor, C. Studer, T. Goldstein, "Visualizing the Loss Landscape of Neural Nets", *NeurIPS 2018*

# Topic II: Action-level Residual Learning

- Two popular pipelines: 1) superposition and 2) filtering
- Some materials from [http://www.yisongyue.com/talks/safety\\_critical\\_learning.pdf](http://www.yisongyue.com/talks/safety_critical_learning.pdf)



# Example 1: Deep RL + Hand-engineered Controller

- In the training process,  $\pi_H$  is “encoded” in the dynamics

## Residual Reinforcement Learning for Robot Control

Tobias Johannink<sup>\*1,3</sup>, Shikhar Bahl<sup>\*2</sup>, Ashvin Nair<sup>\*2</sup>, Jianlan Luo<sup>1,2</sup>, Avinash Kumar<sup>1</sup>, Matthias Loskyll<sup>1</sup>, Juan Aparicio Ojea<sup>1</sup>, Eugen Solowjow<sup>1</sup>, Sergey Levine<sup>2</sup>

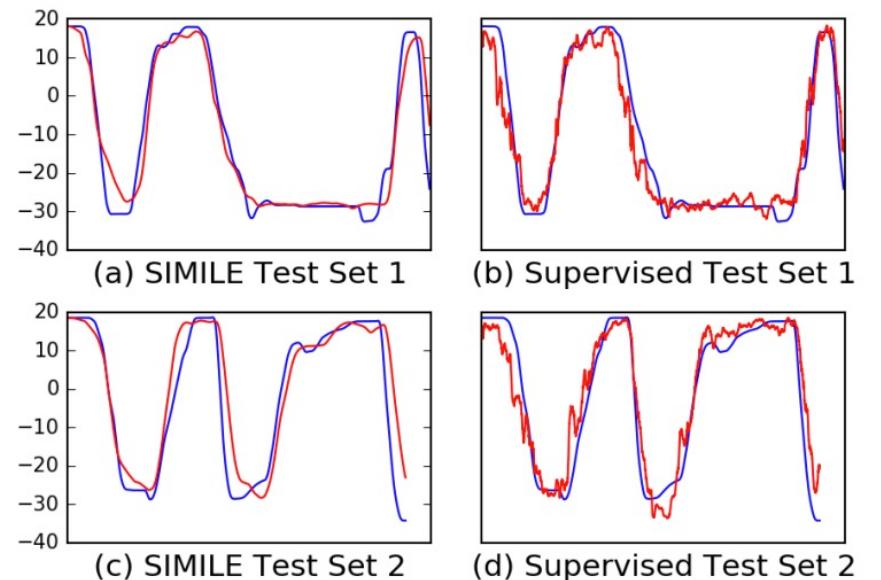
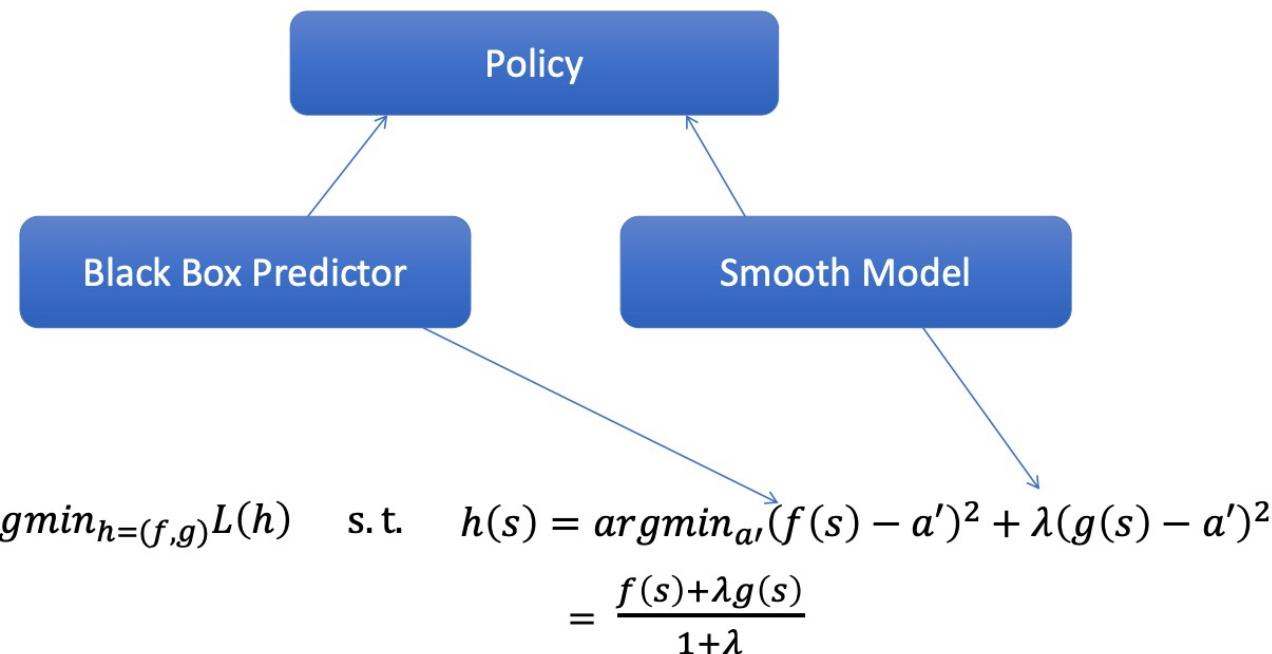
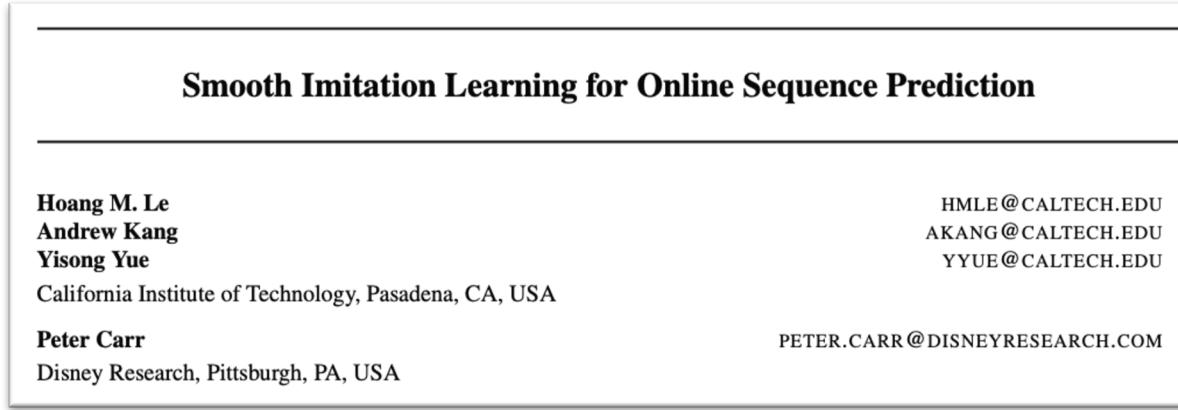
### Algorithm 1 Residual reinforcement learning

**Require:** policy  $\pi_\theta$ , hand-engineered controller  $\pi_H$ .

```
1: for  $n = 0, \dots, N - 1$  episodes do
2:   Initialize random process  $\mathcal{N}$  for exploration
3:   Sample initial state  $s_0 \sim E$ .
4:   for  $t = 0, \dots, H - 1$  steps do hand-engineered controller
5:     Get policy action  $u_t = \pi_\theta(s_t) + \mathcal{N}_t$ .
6:     Get action to execute  $u'_t = u_t + \pi_H(s_t)$ .
7:     Get next state  $s_{t+1} \sim p(\cdot | s_t, u'_t)$ .
8:     Store  $(s_t, u_t, s_{t+1})$  into replay buffer  $\mathcal{R}$ .
9:     Sample set of transitions  $(s, u, s') \sim \mathcal{R}$ .
10:    Optimize  $\theta$  using RL with sampled transitions.
11:   end for
12: end for
```



# Example 2: Smooth Imitation Learning



# Example 3: Control Regularization Reduces Variance in RL

## Control Regularization for Reduced Variance Reinforcement Learning

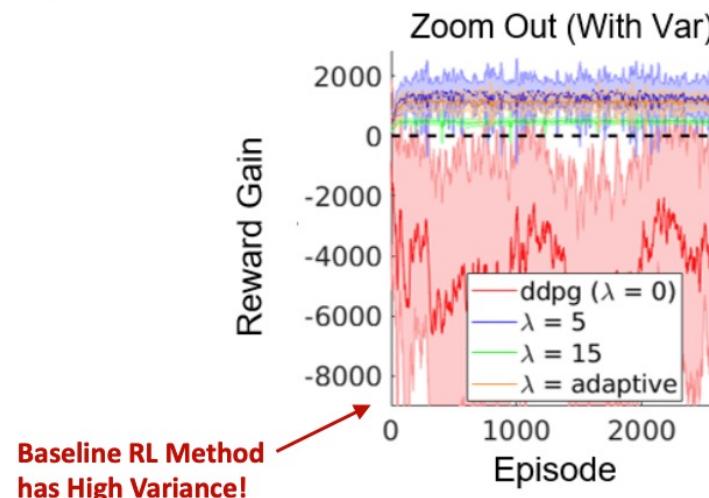
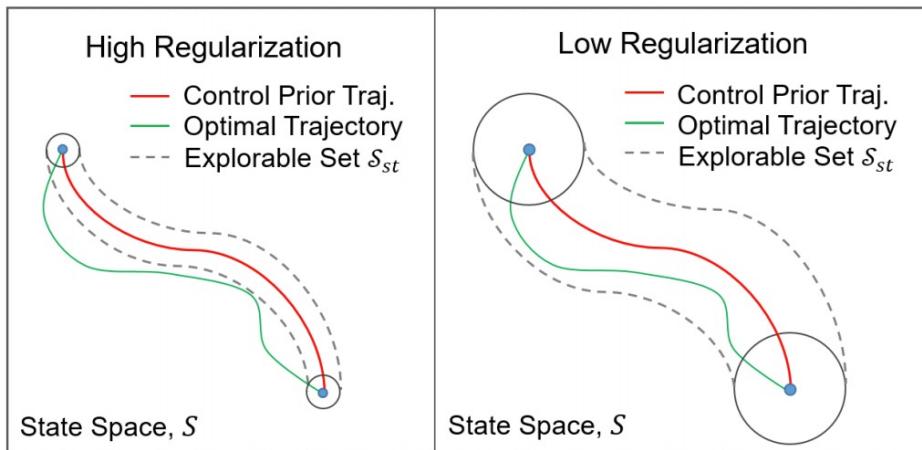
Richard Cheng<sup>1</sup> Abhinav Verma<sup>2</sup> Gábor Orosz<sup>3</sup> Swarat Chaudhuri<sup>2</sup> Yisong Yue<sup>1</sup> Joel W. Burdick<sup>1</sup>

- Theorem (informal):

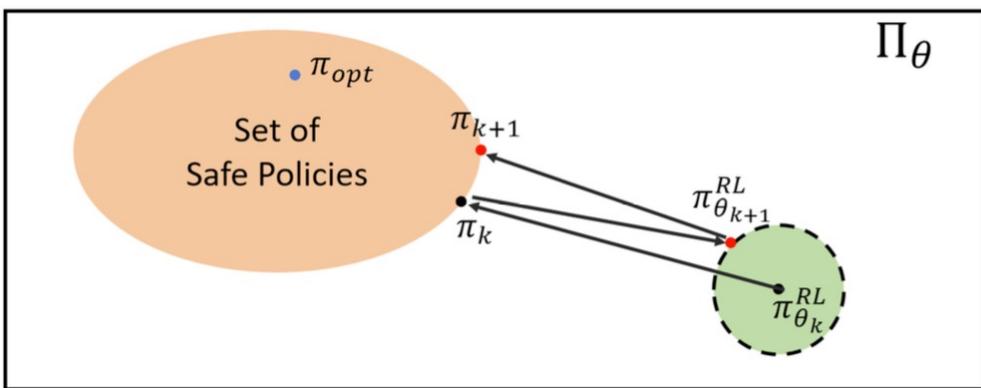
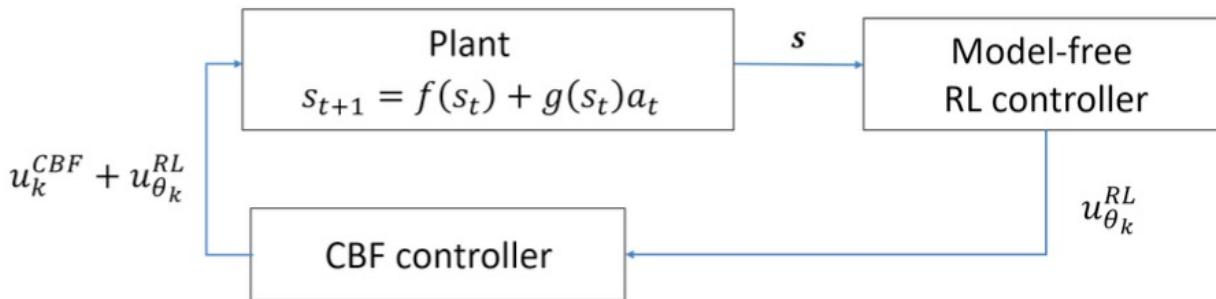
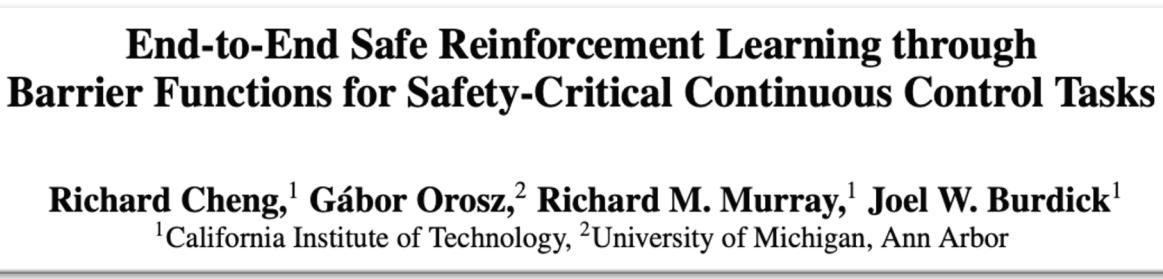
$$u_k(s) = \frac{1}{1+\lambda} u_{\theta_k}(s) + \frac{\lambda}{1+\lambda} u_{prior}(s)$$

- Variance of policy gradient decreases by factor of:  $\left(\frac{1}{1+\lambda}\right)^2$
- Bias converges to:  $\left(\frac{\lambda}{1+\lambda}\right) D_{TV}(h^*, g)$

Implies much faster learning!



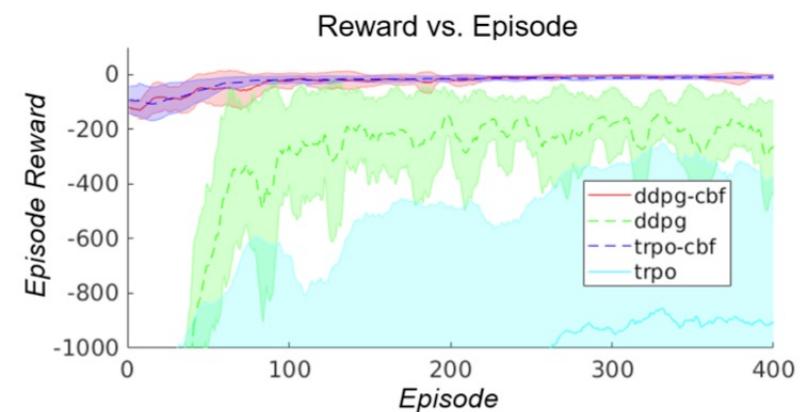
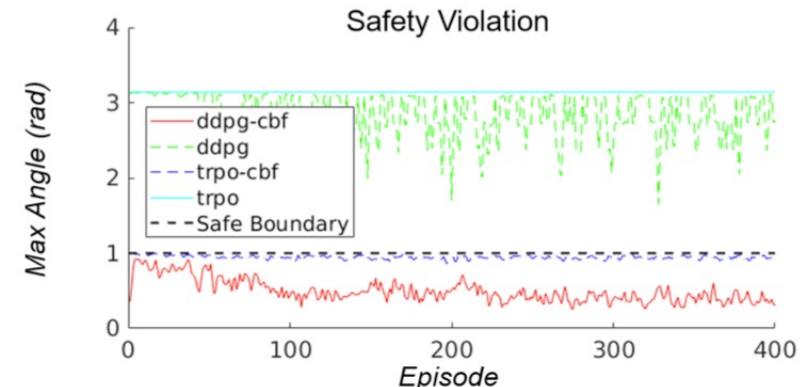
# Example 4: Model-based Controller as a “Filter”



$$(a_t, \epsilon) = \underset{a_t, \epsilon}{\operatorname{argmin}} \|a_t\|_2 + K_\epsilon \epsilon$$

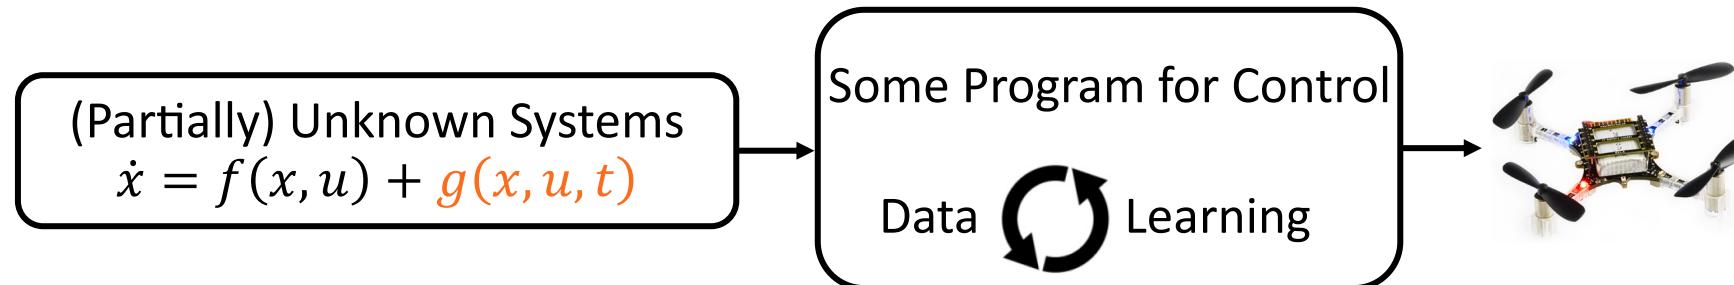
$$\begin{aligned} \text{s.t. } & p^T f(s_t) + p^T g(s_t) \left( u_{\theta_k}^{RL}(s_t) + a_t \right) + p^T \mu_d(s_t) \\ & - k_\delta |p|^T \sigma_d(s_t) + q \geq (1 - \eta) h(s_t) - \epsilon \\ & a_{low}^i \leq a_t^i + u_{\theta_k}^{RL(i)}(s_t) \leq a_{high}^i \text{ for } i = 1, \dots, M \end{aligned}$$

CBF  
safety  
constraints



# Topic III: Program-level Residual Learning

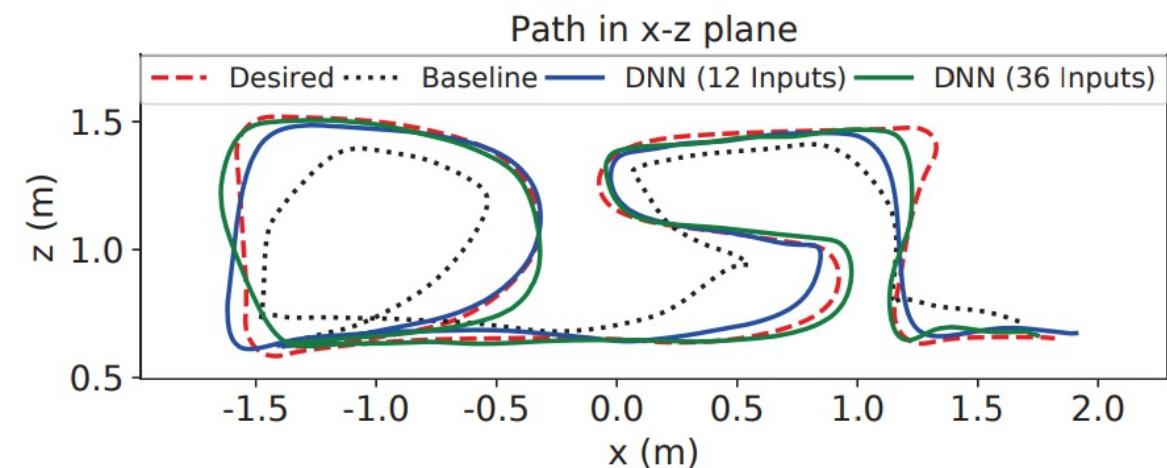
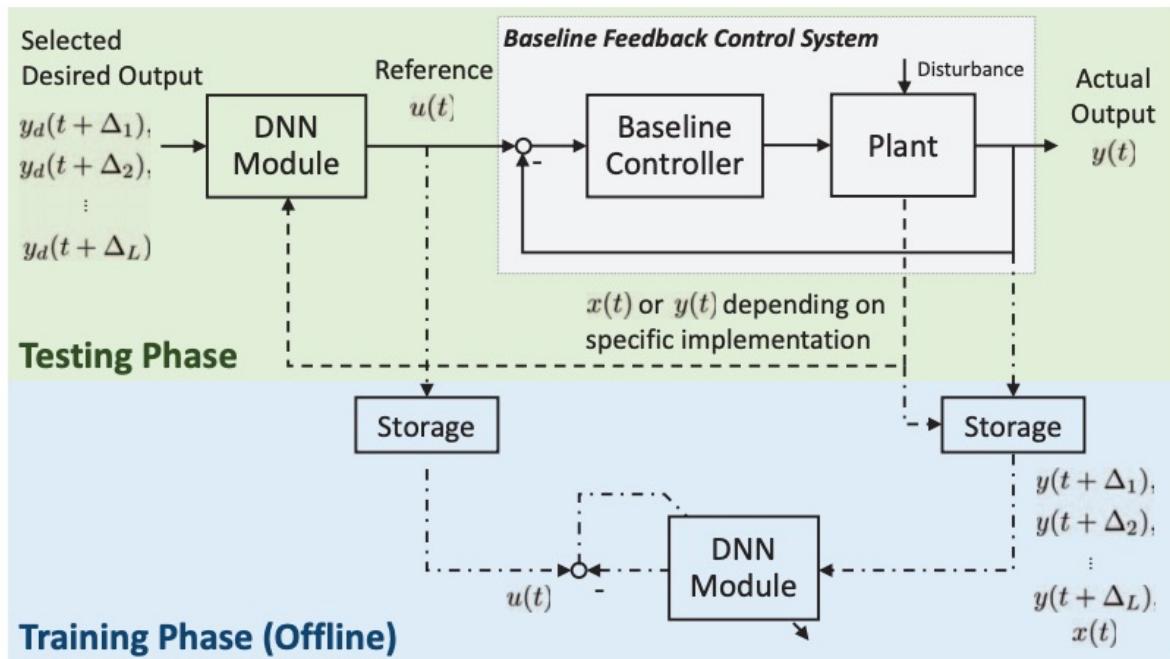
- The “program” design is on a case-by-case basis
- Algorithm design and analysis are not as clear as the dynamics-level and action-level residual learning
- A (very) general framework:



# Example 1: DNN to “Adapt” the Reference Signal

## Design of Deep Neural Networks as Add-on Blocks for Improving Impromptu Trajectory Tracking

Siqi Zhou, Mohamed K. Helwa, and Angela P. Schoellig



# Example 2: Learning Control Lyapunov Function Residual

## A Control Lyapunov Perspective on Episodic Learning via Projection to State Stability

Andrew J. Taylor<sup>1</sup>, Victor D. Dorobantu<sup>1</sup>, Meera Krishnamoorthy,  
Hoang M. Le, Yisong Yue, and Aaron D. Ames

- In control we only need to make sure

$$\mathcal{U}(\mathbf{x}) = \{\mathbf{u} \in \mathcal{U} : \dot{\hat{V}}(\mathbf{x}, \mathbf{u}) \leq -\alpha(\|\mathbf{x}\|)\},$$

$$\dot{\mathbf{x}} = \hat{\mathbf{f}}(\mathbf{x}) + \hat{\mathbf{g}}(\mathbf{x})\mathbf{u} + \underbrace{(\mathbf{g}(\mathbf{x}) - \hat{\mathbf{g}}(\mathbf{x}))\mathbf{u}}_{\mathbf{A}(\mathbf{x})} + \underbrace{\mathbf{f}(\mathbf{x}) - \hat{\mathbf{f}}(\mathbf{x})}_{\mathbf{b}(\mathbf{x})}, \quad (23) \quad \text{Residual dynamics}$$

$$\begin{aligned} \dot{\hat{V}}(\mathbf{x}, \mathbf{u}, \mathbf{d}) &= \overbrace{(\hat{\mathbf{f}}(\mathbf{x}) + \hat{\mathbf{g}}(\mathbf{x})\mathbf{u})^\top \nabla V(\mathbf{x})}^{\hat{V}(\mathbf{x}, \mathbf{u})} \\ &\quad + \underbrace{(\mathbf{A}(\mathbf{x})^\top \nabla V(\mathbf{x}))^\top \mathbf{u}}_{\mathbf{a}(\mathbf{x})} + \underbrace{\mathbf{b}(\mathbf{x})^\top \nabla V(\mathbf{x})}_{\mathbf{b}(\mathbf{x})}, \quad (24) \end{aligned}$$

If we have a control Lyapunov function  $V$ , we can just learn the projection

# Example 3: Differentiable MPC

## Differentiable MPC for End-to-end Planning and Control

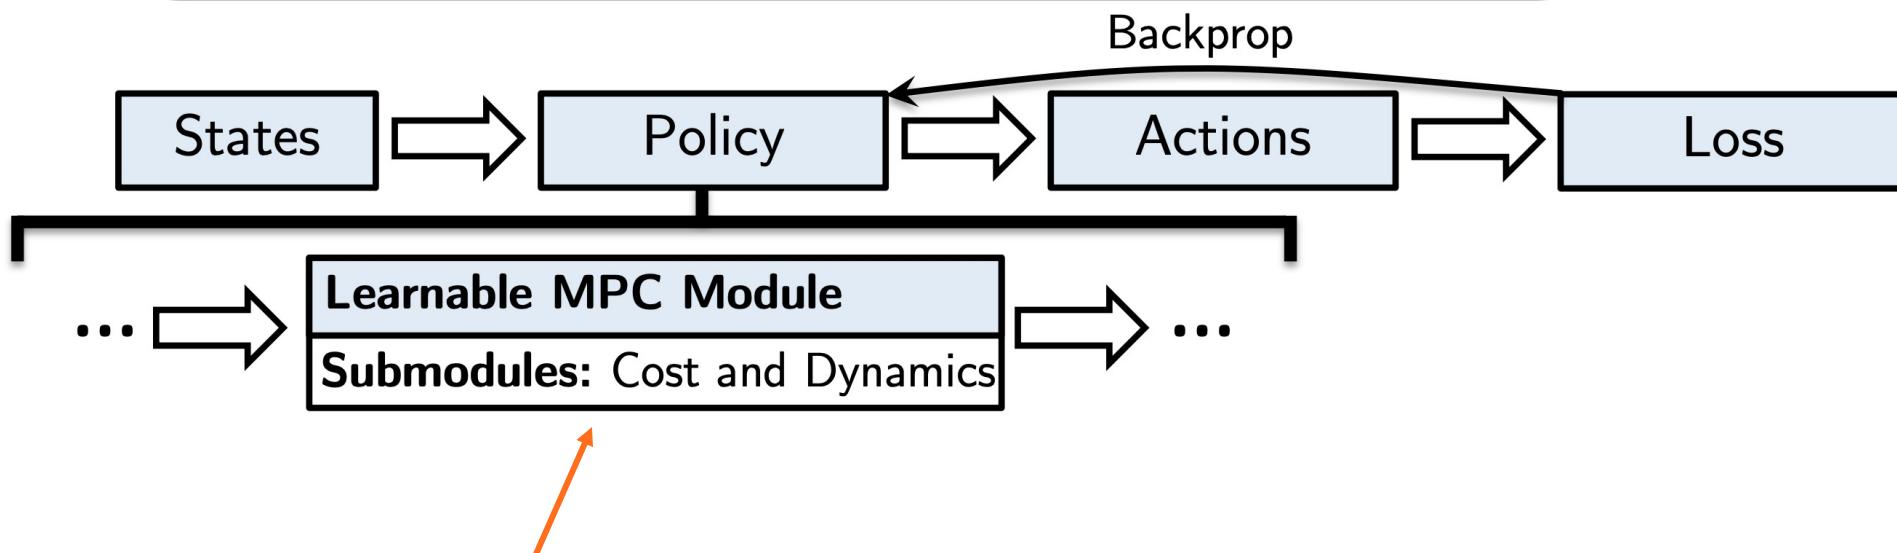
Brandon Amos<sup>1</sup> Ivan Dario Jimenez Rodriguez<sup>2</sup> Jacob Sacks<sup>2</sup>

Byron Boots<sup>2</sup> J. Zico Kolter<sup>13</sup>

<sup>1</sup>Carnegie Mellon University

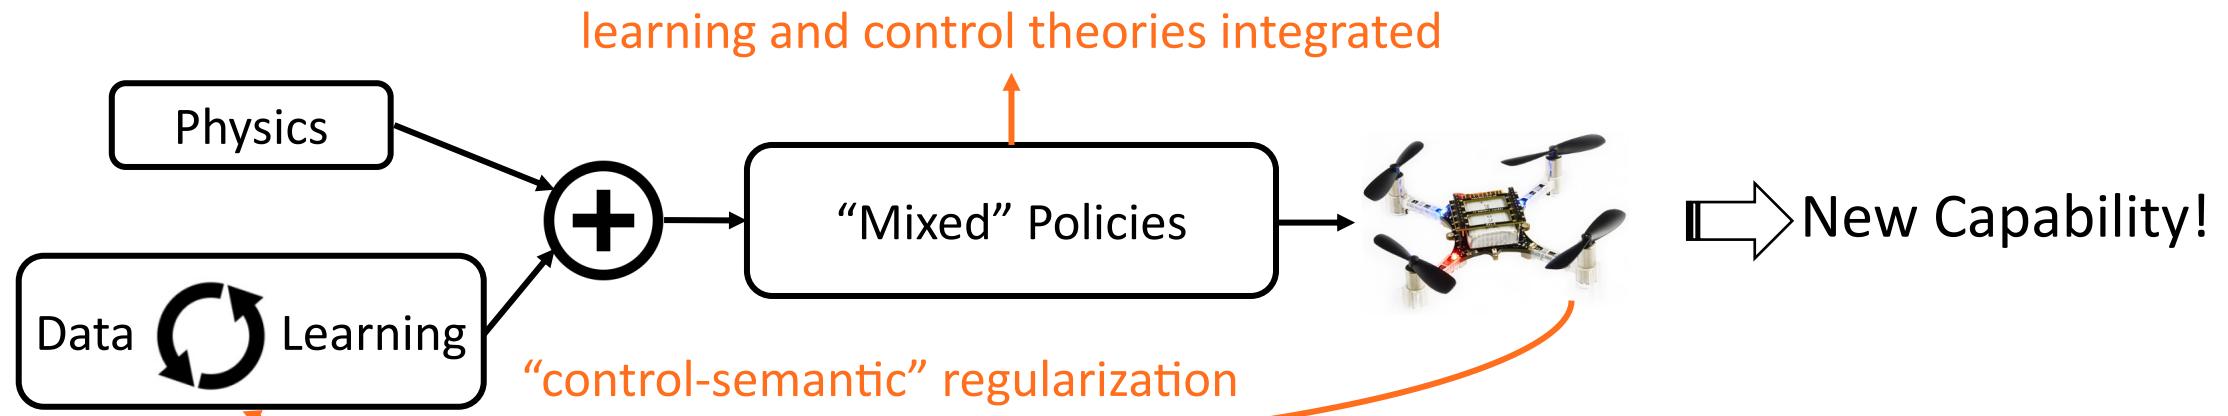
<sup>2</sup>Georgia Tech

<sup>3</sup>Bosch Center for AI



$$\tau_{1:T}^* = \underset{\tau_{1:T}}{\operatorname{argmin}} \sum_t C_{\theta,t}(\tau_t) \text{ subject to } x_1 = x_{\text{init}}, x_{t+1} = f_\theta(\tau_t), \underline{u} \leq u \leq \bar{u},$$

# Summary



- Topic I: dynamics-level residual learning
- Topic II: action-level residual learning
- Topic III: program-level residual learning

Some directions:

- Trade-offs (e.g., sampling complexity)?
- Combine control and learning theory (e.g., generalization)

**“1+1>2”**

approximation error of DNN ( $\|\hat{f}_a - f_a\|$ )

$$\frac{\epsilon_m}{\lambda - L_a \rho}$$