

# Neural Architecture Design



Jeremy Bernstein  
[bernstein@caltech.edu](mailto:bernstein@caltech.edu)

# Structure of this topic

Six lectures, covering:

1. Tools for understanding neural nets
2. Application: optimisation
3. Application: generalisation

Plus two homeworks.

# Agenda for today

1. Class philosophy
2. Neural network basics
3. Motivating questions
4. Architecture design
5. Perturbation theory

# Why theorise?

# Why theorise?

Some reasons people do machine learning theory:

- They like math (aesthetes).
- "*You couldn't possibly use an algorithm without a theoretical guarantee!*"

# Why theorise?

In this class, the main motivation will be:

- Build a better understanding of what works,  
so as to both improve and build upon it.

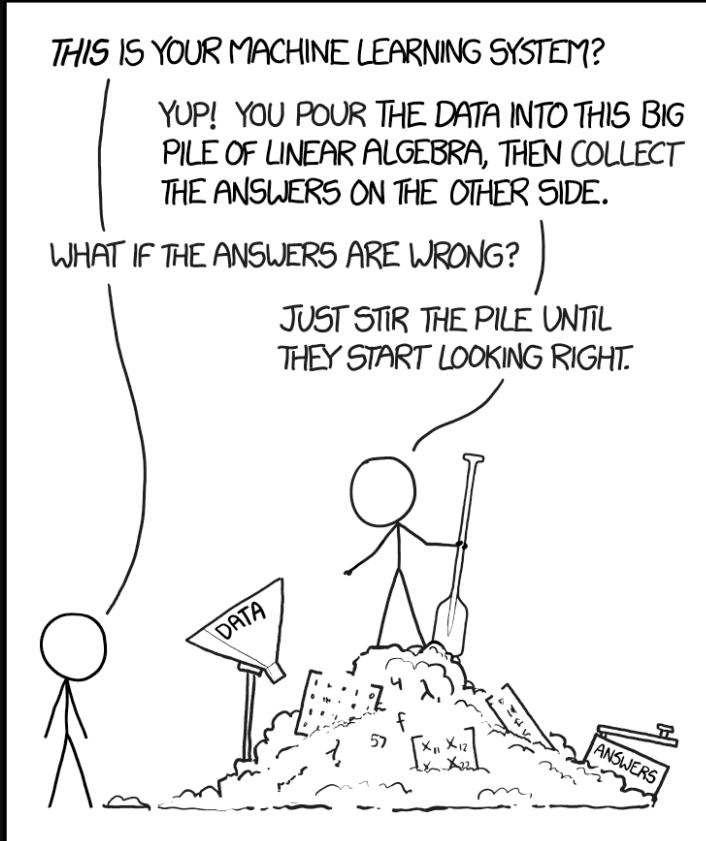


e.g. combine w/ control

# The stages of theory

1. Empirical exploration
2. Modelling  "all models are wrong  
... Some are useful"
3. Derivation
4. Empirical validation

# Pure exploration



[xkcd.com/1838](https://xkcd.com/1838)

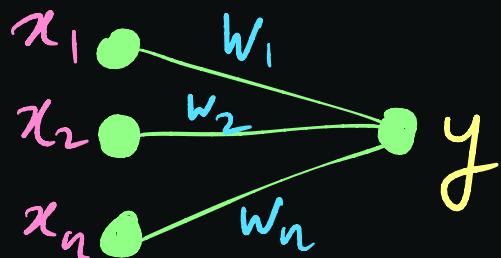
# Pure derivation



[kurzgesagt.org](http://kurzgesagt.org)

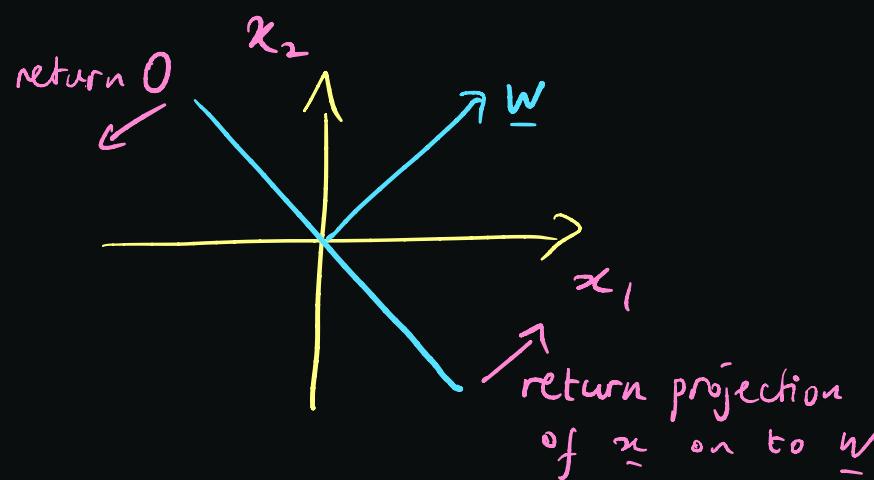
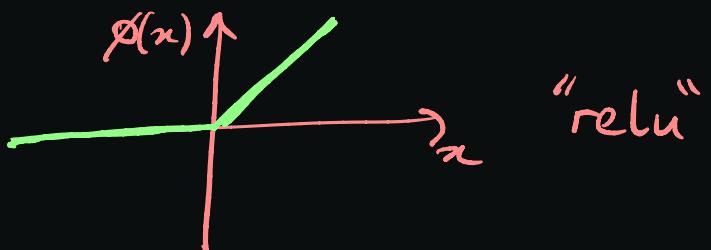
# Neural network basics

# The artificial neuron



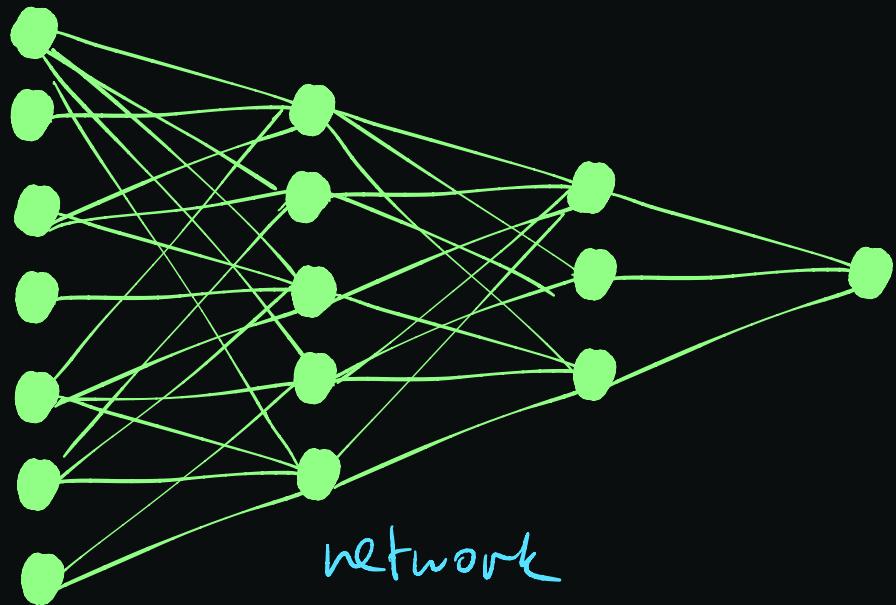
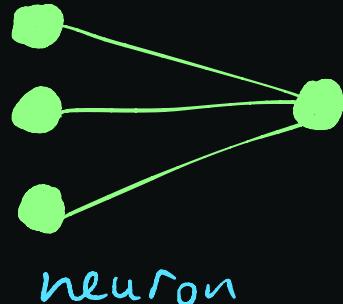
$$y = \phi\left(\sum_{i=1}^n w_i x_i\right) = \phi(\underline{w}^\top \underline{x})$$

$\phi$  is the nonlinearity,  
e.g.  $\phi(x) = \max(0, x)$



} geometric interpretation  
of relu neuron

# Composing neurons

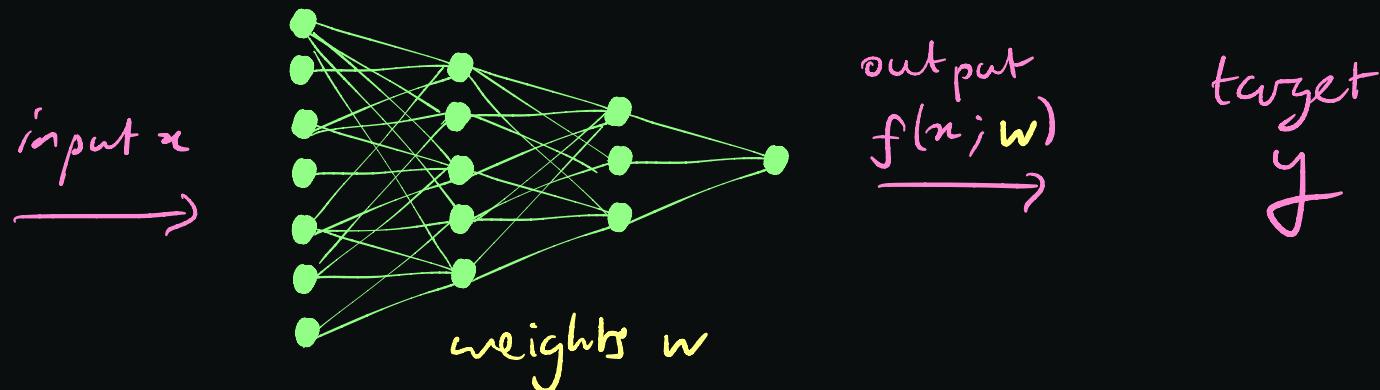


directed acyclic graph  
composed of neurons

General question: how do local properties of neurons  
translate to global properties of the network?

# Backprop: a global view

Wish to train network to fit some targets.



Supervised learning: dataset  $\{x^{(i)}, y^{(i)}\}_{i=1}^N$

Construct loss function  $\mathcal{L}(w) = \sum_{i=1}^N (f(x^{(i)}; w) - y^{(i)})^2$

Run gradient descent:  $w \rightarrow w - \eta \nabla_w \mathcal{L}(w)$

$\eta$  is the "learning rate" — how small should it be?

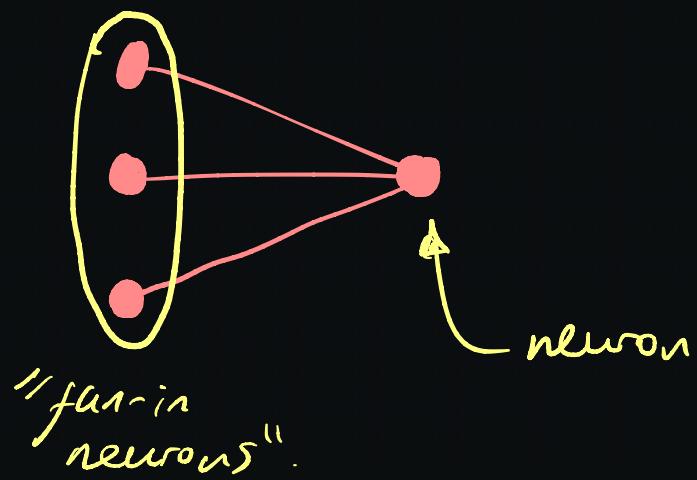
# Backprop: a local view

ith activation  
at layer  $l+1$

$$h_{l+1}^i = \text{relu}\left(\sum_j W_{l+1}^{ij} h_l^j\right)$$

(ij)th weight  
at layer  $l+1$

A neuron aggregates inputs over its "fan-in"

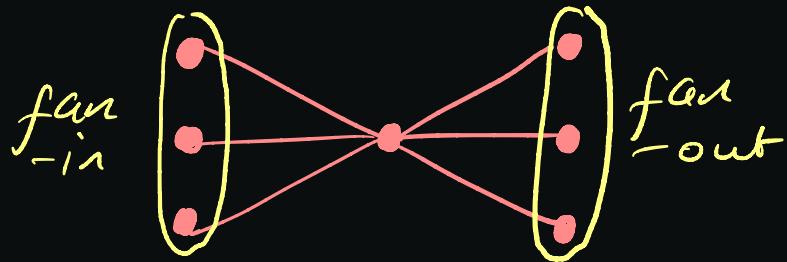


# Backprop: a local view

Backward pass (2 steps)

- $\frac{\partial \mathcal{L}}{\partial h_l^j} = \sum_i \frac{\partial \mathcal{L}}{\partial h_{l+1}^i} \times \mathbb{I}[h_{l+1}^i > 0] \times W_{l+1}^{ij}; \quad \textcircled{1}$
- $\frac{\partial \mathcal{L}}{\partial W_l^{ij}} = \frac{\partial \mathcal{L}}{\partial h_l^i} \times \mathbb{I}[h_l^i > 0] \times h_{l-1}^j. \quad \textcircled{2}$

① a neuron aggregates gradients over its "fan-out"



② the neuron uses this gradient to update its fan-in weights. 15

# Motivating questions

# Optimisation

- How do we design principled training algorithms for neural nets?
- How do we move beyond classic optimisation theory?

e.g. convex opt.

currently, practitioners tune not only the optimiser hyperparameters, but also the optimiser itself (e.g. Adam vs SGD)

# Generalisation

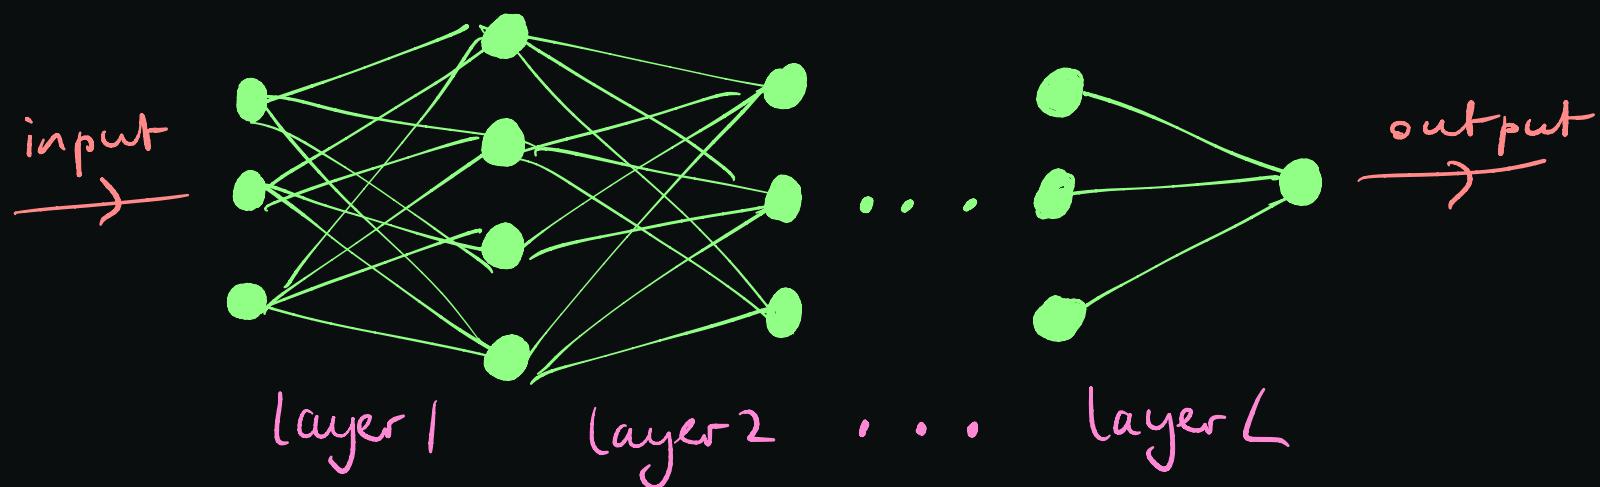
- Why do neural nets generalise when  
 $\#\text{parameters} \gg \#\text{data}$ ?
- Why do neural nets generalise when they have the capacity to fit any labelling of the training data?

( this violates the central premise of Vapnik - Chervonenkis Learning Theory .

# Neural architecture design

# Multilayer perceptron (MLP)

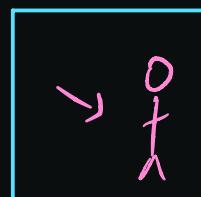
- layered structure
- each layer is a matrix followed by nonlinearity
- assumes little about the structure of the input.



$$f(x; w) = (\phi \circ w_L) \circ \dots \circ (\phi \circ w_1)(x)$$

# Convolutional neural network (CNN)

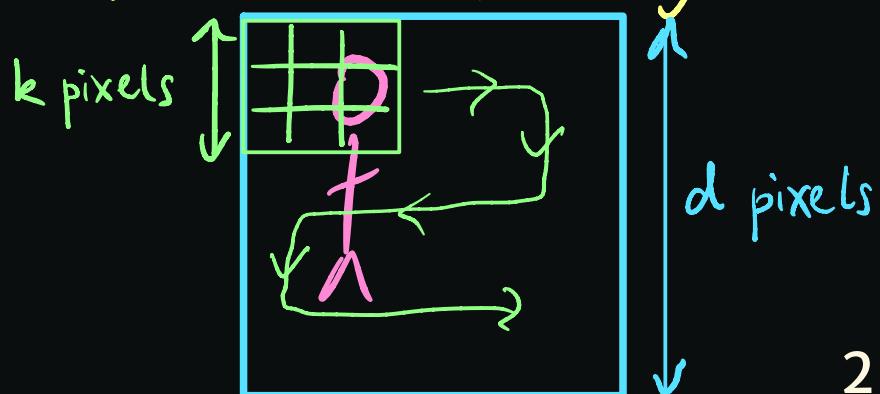
- assumes input is a 2D image
- input has translation invariance



A network layer exploits this structure by convolving a small filter with the image instead of doing a full matrix multiply.

Compare # parameters

CNN filter	$k \times k$
MLP neuron	$d \times d$



# Architecture zoo

Different architectures account for data with different structure. For example:

Structure

vector

image

sequence

Architecture

MLP

CNN

Transformer

# Neural architecture search (NAS)

NAS is a computational approach to discovering new architectures.

It comes in two main flavours:

- ① train lots of networks with slightly different architecture, e.g. NAS via reinforcement learning  
[ — can be viewed as an "evolutionary" outer loop ]  
where network training is the inner loop.
- ② try to learn the network weights and architecture at the same time, e.g. "DARTS"

# What's missing?

## Neural Architecture Search

- ① computationally expensive
- ② results of search biased by how the search space is defined — would NAS discover transformers?

## Intuitive approach "use CNNs for image data"

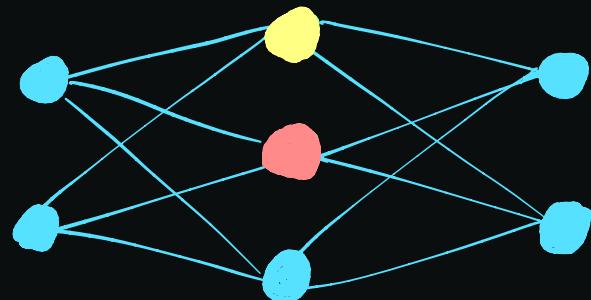
- ① not explicit about what the role of architecture is
- ② doesn't answer concrete questions like :  
"using architecture X to learn dataset Y will take Z datapoints"

We just looked at a global property of the architecture — network topology.

For now, let's turn to some local properties of neurons and the nonlinearity.

## Local properties of architecture

A good rule-of-thumb in architecture design is to ensure that the activations are all on the same scale.



We don't want the activity of neuron A to dominate the activity of neuron B

# Wiring constraints

Consider a "linear neuron" ( $y = \sum_{i=1}^n w_i x_i$ ) and impose two constraints on the weights:

$$(1) \sum_i w_i = 0 \quad \text{---} \quad \text{"balanced excitation & inhibition"}$$

$$(2) \sum_i w_i^2 = 1 \quad \text{---} \quad \text{hyperspherical constraint}$$

Assuming that the inputs  $x_i$  are uncorrelated random variables with the same mean and variance 1, then:

$$Ey = 0 \text{ by (1)} \quad \text{and} \quad Ey^2 = 1 \text{ by (2).}$$

So the output  $y$  has the same scale as the inputs  $x_i$ .

# Nonlinearity design

Consider the "scaled relu" nonlinearity  $\phi(n) = \alpha \cdot \max(0, n)$ .  
— What's the best  $\alpha$ ?

Suppose  $n \sim N(0, 1)$ . Then  $\phi(x)$  is "rectified Gaussian" with variance  $\frac{\alpha^2}{2} \left(1 - \frac{1}{\pi}\right) \approx 0.34\alpha^2$ .

For  $\alpha=1$ , the standard relu nonlinearity tends to "squash" its input by a factor of 0.34.

But by letting  $\alpha = \sqrt{\frac{2}{1 - \frac{1}{\pi}}}$  we avoid this, and obtain  $\text{Var}[\phi(n)] = \text{Var}[x] = 1$ .

We care about this  
because we train  
networks by  
perturbation!

# Perturbation theory

General question: for a network output  $f(x; w)$ ,  
how does  $\Delta f = f(x; w + \Delta w) - f(x; w)$  depend on  
the size of the perturbation  $\Delta w$ ?

# Matrix perturbation theory

There are a lot of results about how a matrix  $A$  behaves under perturbation  $A \mapsto A + \Delta A$ . For example:

## ① perturbation expansions

e.g.  $\lambda_i(A + \Delta A) \approx \lambda_i(A) + h(\Delta A) + \mathcal{O}(\|\Delta A\|^2)$

$\overset{i^{\text{th}} \text{ eigenvalue}}{\nearrow}$   $\overset{\text{some linear function}}{\nearrow}$

## ② perturbation bounds

e.g.  $\|A + \Delta A\|_F \leq \|A\|_F + \|\Delta A\|_F$  Frobenius norm

$\curvearrowleft$  triangle inequality

# Deep perturbation theory

A neural network is just a product of matrices (and nonlinearities).

Consider a toy example for a network with weight vector  $\underline{a} \in \mathbb{R}^d$ .

$$f(x; \underline{a}) = \left( \prod_{i=1}^d a_i \right) x \quad \text{"deep, linear, scalar network"}$$

Perturbation result:

$$\frac{f(x; \underline{a} + \Delta \underline{a}) - f(x; \underline{a})}{f(x; \underline{a})} = \prod_{i=1}^d \left( 1 + \frac{\Delta a_i}{a_i} \right) - 1.$$

Will generalise this to "deep, linear, matrix network" in HW 3.30

# Summary

- we looked at network topology and said things like "CNNs seem to be well-suited to images". We will return to this issue in **lecture 11** when we look at PAC-Bayesian generalisation theory.
- we looked at properties of neurons and saw how they effect the balance of network activity.
- we looked at perturbation theory of compositional functions. This will help in **lecture 9** when we look at optimisation theory of neural nets.

# Next lecture

We will develop a major tool of NN theory:

The Neural network — Gaussian process  
Correspondence

This will let us move from parameter space to function space so that we can study the typical kinds of function that an NN implements.

